



[Home](#) | [Book](#) | [Help](#) | [Contact](#) | [News](#)  | [Follow](#) 

Please see the post [Do not buy the print version of the Ruby on Rails Tutorial \(yet\)](#)

[skip to content](#) | [view as single page](#)

Ruby on Rails Tutorial

Learn Web Development with Rails

Michael Hartl

Contents

Chapter 1 From zero to deploy

1.1 Introduction

1.1.1 Comments for various readers

- 1.1.2 “Scaling” Rails
 - 1.1.3 Conventions in this book
- 1.2 Up and running
 - 1.2.1 Development environments
 - IDEs
 - Text editors and command lines
 - Browsers
 - A note about tools
 - 1.2.2 Ruby, RubyGems, Rails, and Git
 - Rails Installer (Windows)
 - Install Git
 - Install Ruby
 - Install RubyGems
 - Install Rails
 - 1.2.3 The first application
 - 1.2.4 Bundler
 - 1.2.5 rails server
 - 1.2.6 Model-view-controller (MVC)
- 1.3 Version control with Git
 - 1.3.1 Installation and setup
 - First-time system setup
 - First-time repository setup
 - 1.3.2 Adding and committing
 - 1.3.3 What good does Git do you?
 - 1.3.4 GitHub
 - 1.3.5 Branch, edit, commit, merge
 - Branch
 - Edit
 - Commit
 - Merge
 - Push
- 1.4 Deploying
 - 1.4.1 Heroku setup
 - 1.4.2 Heroku deployment, step one
 - 1.4.3 Heroku deployment, step two
 - 1.4.4 Heroku commands
- 1.5 Conclusion

Chapter 2 A demo app

- 2.1 Planning the application
 - 2.1.1 Modeling demo users
 - 2.1.2 Modeling demo microposts
- 2.2 The Users resource
 - 2.2.1 A user tour
 - 2.2.2 MVC in action
 - 2.2.3 Weaknesses of this Users resource
- 2.3 The Microposts resource
 - 2.3.1 A micropost microtour
 - 2.3.2 Putting the *micro* in microposts
 - 2.3.3 A user has_many microposts
 - 2.3.4 Inheritance hierarchies
 - 2.3.5 Deploying the demo app
- 2.4 Conclusion

Chapter 3 Mostly static pages

- 3.1 Static pages
 - 3.1.1 Truly static pages
 - 3.1.2 Static pages with Rails
- 3.2 Our first tests
 - 3.2.1 Test-driven development
 - 3.2.2 Adding a page
 - Red
 - Green
 - Refactor
- 3.3 Slightly dynamic pages
 - 3.3.1 Testing a title change
 - 3.3.2 Passing title tests
 - 3.3.3 Embedded Ruby
 - 3.3.4 Eliminating duplication with layouts
- 3.4 Conclusion
- 3.5 Exercises
- 3.6 Advanced setup
 - 3.6.1 Eliminating bundle exec
 - RVM Bundler integration
 - binstubs

- 3.6.2 Automated tests with Guard
- 3.6.3 Speeding up tests with Spork
Guard with Spork
- 3.6.4 Tests inside Sublime Text

Chapter 4 Rails-flavored Ruby

- 4.1 Motivation
- 4.2 Strings and methods
 - 4.2.1 Comments
 - 4.2.2 Strings
Printing
Single-quoted strings
 - 4.2.3 Objects and message passing
 - 4.2.4 Method definitions
 - 4.2.5 Back to the title helper
- 4.3 Other data structures
 - 4.3.1 Arrays and ranges
 - 4.3.2 Blocks
 - 4.3.3 Hashes and symbols
 - 4.3.4 CSS revisited
- 4.4 Ruby classes
 - 4.4.1 Constructors
 - 4.4.2 Class inheritance
 - 4.4.3 Modifying built-in classes
 - 4.4.4 A controller class
 - 4.4.5 A user class
- 4.5 Conclusion
- 4.6 Exercises

Chapter 5 Filling in the layout

- 5.1 Adding some structure
 - 5.1.1 Site navigation
 - 5.1.2 Bootstrap and custom CSS
 - 5.1.3 Partials
- 5.2 Sass and the asset pipeline
 - 5.2.1 The asset pipeline
Asset directories

- Manifest files
 - Preprocessor engines
 - Efficiency in production
- 5.2.2 Syntactically awesome stylesheets
 - Nesting
 - Variables
- 5.3 Layout links
 - 5.3.1 Route tests
 - 5.3.2 Rails routes
 - 5.3.3 Named routes
 - 5.3.4 Pretty RSpec
- 5.4 User signup: A first step
 - 5.4.1 Users controller
 - 5.4.2 Signup URI
- 5.5 Conclusion
- 5.6 Exercises

Chapter 6 Modeling users

- 6.1 User model
 - 6.1.1 Database migrations
 - 6.1.2 The model file
 - Model annotation
 - Accessible attributes
 - 6.1.3 Creating user objects
 - 6.1.4 Finding user objects
 - 6.1.5 Updating user objects
- 6.2 User validations
 - 6.2.1 Initial user tests
 - 6.2.2 Validating presence
 - 6.2.3 Length validation
 - 6.2.4 Format validation
 - 6.2.5 Uniqueness validation
 - The uniqueness caveat
- 6.3 Adding a secure password
 - 6.3.1 An encrypted password
 - 6.3.2 Password and confirmation
 - 6.3.3 User authentication

- 6.3.4 User has secure password
 - 6.3.5 Creating a user
- 6.4 Conclusion
- 6.5 Exercises

Chapter 7 Sign up

- 7.1 Showing users
 - 7.1.1 Debug and Rails environments
 - 7.1.2 A Users resource
 - 7.1.3 Testing the user show page (with factories)
 - 7.1.4 A Gravatar image and a sidebar
- 7.2 Signup form
 - 7.2.1 Tests for user signup
 - 7.2.2 Using `form_for`
 - 7.2.3 The form HTML
- 7.3 Signup failure
 - 7.3.1 A working form
 - 7.3.2 Signup error messages
- 7.4 Signup success
 - 7.4.1 The finished signup form
 - 7.4.2 The flash
 - 7.4.3 The first signup
 - 7.4.4 Deploying to production with SSL
- 7.5 Conclusion
- 7.6 Exercises

Chapter 8 Sign in, sign out

- 8.1 Sessions and signin failure
 - 8.1.1 Sessions controller
 - 8.1.2 Signin tests
 - 8.1.3 Signin form
 - 8.1.4 Reviewing form submission
 - 8.1.5 Rendering with a flash message
- 8.2 Signin success
 - 8.2.1 Remember me
 - 8.2.2 A working `sign_in` method
 - 8.2.3 Current user

- 8.2.4 Changing the layout links
 - 8.2.5 Signin upon signup
 - 8.2.6 Signing out
- 8.3 Introduction to Cucumber (optional)
 - 8.3.1 Installation and setup
 - 8.3.2 Features and steps
 - 8.3.3 Counterpoint: RSpec custom matchers
- 8.4 Conclusion
- 8.5 Exercises

Chapter 9 Updating, showing, and deleting users

- 9.1 Updating users
 - 9.1.1 Edit form
 - 9.1.2 Unsuccessful edits
 - 9.1.3 Successful edits
- 9.2 Authorization
 - 9.2.1 Requiring signed-in users
 - 9.2.2 Requiring the right user
 - 9.2.3 Friendly forwarding
- 9.3 Showing all users
 - 9.3.1 User index
 - 9.3.2 Sample users
 - 9.3.3 Pagination
 - 9.3.4 Partial refactoring
- 9.4 Deleting users
 - 9.4.1 Administrative users
 - Revisiting `attr_accessible`
 - 9.4.2 The destroy action
- 9.5 Conclusion
- 9.6 Exercises

Chapter 10 User microposts

- 10.1 A Micropost model
 - 10.1.1 The basic model
 - 10.1.2 Accessible attributes and the first validation
 - 10.1.3 User/Micropost associations
 - 10.1.4 Micropost refinements

- Default scope
 - Dependent: destroy
 - 10.1.5 Content validations
- 10.2 Showing microposts
 - 10.2.1 Augmenting the user show page
 - 10.2.2 Sample microposts
- 10.3 Manipulating microposts
 - 10.3.1 Access control
 - 10.3.2 Creating microposts
 - 10.3.3 A proto-feed
 - 10.3.4 Destroying microposts
- 10.4 Conclusion
- 10.5 Exercises

Chapter 11 Following users

- 11.1 The Relationship model
 - 11.1.1 A problem with the data model (and a solution)
 - 11.1.2 User/relationship associations
 - 11.1.3 Validations
 - 11.1.4 Followed users
 - 11.1.5 Followers
- 11.2 A web interface for following users
 - 11.2.1 Sample following data
 - 11.2.2 Stats and a follow form
 - 11.2.3 Following and followers pages
 - 11.2.4 A working follow button the standard way
 - 11.2.5 A working follow button with Ajax
- 11.3 The status feed
 - 11.3.1 Motivation and strategy
 - 11.3.2 A first feed implementation
 - 11.3.3 Subselects
 - 11.3.4 The new status feed
- 11.4 Conclusion
 - 11.4.1 Extensions to the sample application
 - Replies
 - Messaging
 - Follower notifications

Foreword

My former company (CD Baby) was one of the first to loudly switch to Ruby on Rails, and then even more loudly switch back to PHP (Google me to read about the drama). This book by Michael Hartl came so highly recommended that I had to try it, and the *Ruby on Rails Tutorial* is what I used to switch back to Rails again.

Though I've worked my way through many Rails books, this is the one that finally made me “get” it. Everything is done very much “the Rails way”—a way that felt very unnatural to me before, but now after doing this book finally feels natural. This is also the only Rails book that does test-driven development the entire time, an approach highly recommended by the experts but which has never been so clearly demonstrated before. Finally, by including Git, GitHub, and Heroku in the demo examples, the author really gives you a feel for what it's like to do a real-world project. The tutorial's code examples are not in isolation.

The linear narrative is such a great format. Personally, I powered through the *Rails Tutorial* in three long days, doing all the examples and challenges at the end of each chapter. Do it from start to finish, without jumping around, and you'll get the ultimate benefit.

Enjoy!

[Derek Sivers \(sivers.org\)](http://sivers.org)

Formerly: Founder, [CD Baby](#)

Currently: Founder, [Thoughts Ltd.](#)

Acknowledgments

The *Ruby on Rails Tutorial* owes a lot to my previous Rails book, *RailsSpace*, and hence to my coauthor [Aurelius Prochazka](#). I'd like to thank Aure both for the work he did on that book and for his support of this one. I'd also like to thank Debra Williams Cauley, my editor on both *RailsSpace* and the *Ruby on Rails Tutorial*; as long as she keeps taking me to baseball games, I'll keep writing books for her.

I'd like to acknowledge a long list of Rubyists who have taught and inspired me over the years: David Heinemeier Hansson, Yehuda Katz, Carl Lerche, Jeremy Kemper, Xavier Noria, Ryan Bates, Geoffrey Grosenbach, Peter Cooper, Matt Aimonetti, Gregg Pollack, Wayne E. Seguin, Amy Hoy, Dave Chelimsky, Pat Maddox, Tom Preston-Werner, Chris Wanstrath, Chad Fowler, Josh Susser, Obie Fernandez, Ian McFarland, Steven Bristol, Wolfram Arnold, Alex Chaffee, Giles Bowkett, Evan Dorn, Long Nguyen, James Lindenbaum, Adam Wiggins, Tikhon Bernstam, Ron Evans, Wyatt Greene, Miles Forrest, the good people at Pivotal Labs, the Heroku gang, the thoughtbot guys, and the GitHub crew. Finally, many, many readers—far too many to list—have contributed a huge number of bug reports and suggestions during the writing of this book, and I gratefully acknowledge their help in making it as good as it can be.

About the author

[Michael Hartl](#) is the author of the [Ruby on Rails Tutorial](#), the leading introduction to web development with [Ruby on Rails](#). His prior experience includes writing and developing *RailsSpace*, an extremely obsolete Rails tutorial book, and developing Insoshi, a once-popular and now-obsolete social networking platform in Ruby on Rails. In 2011, Michael received a [Ruby Hero Award](#) for his contributions to the Ruby community. He is a graduate of [Harvard College](#), has a [Ph.D. in Physics](#) from [Caltech](#), and is an alumnus of the [Y Combinator](#) entrepreneur program.

Copyright and license

Ruby on Rails Tutorial: Learn Web Development with Rails. Copyright © 2012 by Michael Hartl.

All source code in the *Ruby on Rails Tutorial* is available jointly under the [MIT License](#) and the [Beerware License](#).

The MIT License

Copyright (c) 2012 Michael Hartl

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

```
/*
 * -----
 * "THE BEER-WARE LICENSE" (Revision 42):
 * Michael Hartl wrote this code. As long as you retain this notice you
 * can do whatever you want with this stuff. If we meet some day, and you think
 * this stuff is worth it, you can buy me a beer in return.
 * -----
 */
```

Chapter 4

Rails-flavored Ruby

Grounded in examples from [Chapter 3](#), this chapter explores some elements of Ruby important for Rails. Ruby is a big language, but fortunately the subset needed to be productive as a Rails developer is relatively small. Moreover, this subset is *different* from the usual approaches to learning Ruby, which is why, if your goal is making dynamic web applications, I recommend learning Rails first, picking up bits of Ruby along the way. To become a Rails *expert*, you need to understand Ruby more deeply, and this book gives you a good foundation for developing that expertise. As noted in [Section 1.1.1](#), after finishing the *Rails Tutorial* I suggest reading a pure Ruby book such as [Beginning Ruby](#), [The Well-Grounded Rubyist](#), or [The Ruby Way](#).

This chapter covers a lot of material, and it's OK not to get it all on the first pass. I'll refer back to it frequently in future chapters.

4.1 Motivation

As we saw in the last chapter, it's possible to develop the skeleton of a Rails application, and even start testing it, with essentially no knowledge of the underlying Ruby language. We did this by relying on the test code provided by the tutorial and addressing each error message until the test suite was passing. This situation can't last forever, though, and we'll open this chapter with an addition to the site that brings us face-to-face with our Ruby limitations.

When we last saw our new application, we had just updated our mostly static pages to use Rails layouts to eliminate duplication in our views ([Listing 4.1](#)).



PDF/Screencasts



📡 ✉️ 📘 🐦

Rails 3.2 • **Rails 3.0**

👍 Like 2k

Listing 4.1. The sample application site layout.

`app/views/layouts/application.html.erb`

```
<!DOCTYPE html>
<html>
  <head>
    <title>Ruby on Rails Tutorial Sample App | <%= yield(:title) %></title>
    <%= stylesheet_link_tag "application", :media => "all" %>
    <%= javascript_include_tag "application" %>
    <%= csrf_meta_tags %>
  </head>
  <body>
    <%= yield %>
  </body>
</html>
```

Let's focus on one particular line in [Listing 4.1](#):

```
<%= stylesheet_link_tag "application", :media => "all" %>
```

This uses the built-in Rails function `stylesheet_link_tag` (which you can read more about at the [Rails API](#)) to include `application.css` for all [media types](#) (including computer screens and printers). To an experienced Rails developer, this line looks simple, but there are at least four potentially confusing Ruby ideas: built-in Rails methods, method invocation with missing parentheses, symbols, and hashes. We'll cover all of these ideas in this chapter.

In addition to coming equipped with a large number of built-in functions for use in the views, Rails also allows the creation of new ones. Such functions are called *helpers*; to see how to make a custom helper, let's start by examining the title line from [Listing 4.1](#):

This relies on the definition of a page title (using **provide**) in each view, as in

```
<% provide(:title, 'Home') %>
<h1>Sample App</h1>
<p>
  This is the home page for the
  <a href="http://railstutorial.org/">Ruby on Rails Tutorial</a>
  sample application.
</p>
```

But what if we don't provide a title? It's a good convention to have a *base title* we use on every page, with an optional page title if we want to be more specific. We've *almost* achieved that with our current layout, with one wrinkle: as you can see if you delete the **provide** call in one of the views, in the absence of a page-specific title the full title appears as follows:

Ruby on Rails Tutorial Sample App |

In other words, there's a suitable base title, but there's also a trailing vertical bar character | at the end.

To solve the problem of a missing page title, we'll define a custom helper called **full_title**. The **full_title** helper returns a base title, "Ruby on Rails Tutorial Sample App", if no page title is defined, and adds a vertical bar followed by the page title if one is defined ([Listing 4.2](#)).¹

Listing 4.2. Defining a **full_title** helper.

app/helpers/application_helper.rb

```

module ApplicationHelper

  # Returns the full title on a per-page basis.
  def full_title(page_title)
    base_title = "Ruby on Rails Tutorial Sample App"
    if page_title.empty?
      base_title
    else
      "#{base_title} | #{page_title}"
    end
  end
end

```

Now that we have a helper, we can use it to simplify our layout by replacing

```

<title>Ruby on Rails Tutorial Sample App | <%= yield(:title) %></title>

```

with

```

<title><%= full_title(yield(:title)) %></title>

```

as seen in [Listing 4.3](#).

Listing 4.3. The sample application site layout.

`app/views/layouts/application.html.erb`

```

<!DOCTYPE html>
<html>
  <head>
    <title><%= full_title(yield(:title)) %></title>
    <%= stylesheet_link_tag "application", :media => "all" %>
    <%= javascript_include_tag "application" %>
    <%= csrf_meta_tags %>

```

```
</head>
<body>
  <%= yield %>
</body>
</html>
```

To put our helper to work, we can eliminate the unnecessary word “Home” from the Home page, allowing it to revert to the base title. We do this by first updating our test with the code in [Listing 4.4](#), which updates the previous title test and adds one to test for the absence of the custom ‘Home’ string in the title.

Listing 4.4. Updated tests for the Home page’s title.

`spec/requests/static_pages_spec.rb`

```
require 'spec_helper'

describe "Static pages" do
  describe "Home page" do
    it "should have the h1 'Sample App'" do
      visit '/static_pages/home'
      page.should have_selector('h1', :text => 'Sample App')
    end

    it "should have the base title" do
      visit '/static_pages/home'
      page.should have_selector('title',
                               :text => "Ruby on Rails Tutorial Sample App")
    end

    it "should not have a custom page title" do
      visit '/static_pages/home'
      page.should_not have_selector('title', :text => '| Home')
    end
  end
end
```



```
end
```

See if you can figure out why we've added a new test instead of just altering the current one. (*Hint:* The answer is in [Section 3.3.1.](#))

Let's run the test suite to verify that one test fails:

```
$ bundle exec rspec spec/requests/static_pages_spec.rb
```

To get the test suite to pass, we'll remove the `provide` line from the Home page's view, as seen in [Listing 4.5.](#)

Listing 4.5. The Home page with no custom page title.

`app/views/static_pages/home.html.erb`

```
<h1>Sample App</h1>
<p>
  This is the home page for the
  <a href="http://railstutorial.org/">Ruby on Rails Tutorial</a>
  sample application.
</p>
```

At this point the tests should pass:

```
$ bundle exec rspec spec/requests/static_pages_spec.rb
```

As with the line to include the application stylesheet, the code in [Listing 4.2](#) may look simple to the eyes of an experienced Rails developer, but it's *full* of potentially confusing Ruby ideas: modules,

comments, local variable assignment, booleans, control flow, string interpolation, and return values. This chapter will cover all of these ideas as well.

4.2 Strings and methods

Our principal tool for learning Ruby will be the *Rails console*, a command-line tool for interacting with Rails applications first seen in [Section 2.3.3](#). The console itself is built on top of interactive Ruby (**irb**), and thus has access to the full power of the Ruby language. (As we'll see in [Section 4.4.4](#), the console also has access to the Rails environment.) Start the console at the command line as follows:

```
$ rails console
Loading development environment
>>
```

By default, the console starts in a *development environment*, which is one of three separate environments defined by Rails (the others are *test* and *production*). This distinction won't be important in this chapter, but we'll learn more about environments in [Section 7.1.1](#).

The console is a great learning tool, and you should feel free to explore—don't worry, you (probably) won't break anything. When using the console, type Ctrl-C if you get stuck, or Ctrl-D to exit the console altogether. Throughout the rest of this chapter, you might find it helpful to consult the [Ruby API](#). It's packed (perhaps even *too* packed) with information; for example, to learn more about Ruby strings you can look at the Ruby API entry for the **String** class.

4.2.1 Comments

Ruby *comments* start with the pound sign **#** (also called the “hash mark” or, more poetically, the “octothorpe”) and extend to the end of the line. Ruby ignores comments, but they are useful for

human readers (including, often, the original author!). In the code

```
# Returns the full title on a per-page basis.
def full_title(page_title)
  .
  .
  .
end
```

the first line is a comment indicating the purpose of the subsequent function definition.

You don't ordinarily include comments in console sessions, but for instructional purposes I'll include some comments in what follows, like this:

```
$ rails console
>> 17 + 42   # Integer addition
=> 59
```

If you follow along in this section typing or copying-and-pasting commands into your own console, you can of course omit the comments if you like; the console will ignore them in any case.

4.2.2 Strings

Strings are probably the most important data structure for web applications, since web pages ultimately consist of strings of characters sent from the server to the browser. Let's start exploring strings with the console, this time started with **rails c**, which is a shortcut for **rails console**:

```
$ rails c
>> ""        # An empty string
=> ""
```

```
>> "foo"      # A nonempty string
=> "foo"
```

These are *string literals* (also, amusingly, called *literal strings*), created using the double quote character `"`. The console prints the result of evaluating each line, which in the case of a string literal is just the string itself.

We can also concatenate strings with the `+` operator:

```
>> "foo" + "bar"    # String concatenation
=> "foobar"
```

Here the result of evaluating `"foo"` plus `"bar"` is the string `"foobar"`.²

Another way to build up strings is via *interpolation* using the special syntax `#{}:`³

```
>> first_name = "Michael"    # Variable assignment
=> "Michael"
>> "#{first_name} Hartl"    # String interpolation
=> "Michael Hartl"
```

Here we've *assigned* the value `"Michael"` to the variable `first_name` and then interpolated it into the string `"#{first_name} Hartl"`. We could also assign both strings a variable name:

```
>> first_name = "Michael"
=> "Michael"
>> last_name = "Hartl"
=> "Hartl"
>> first_name + " " + last_name    # Concatenation, with a space in between
=> "Michael Hartl"
```

```
>> "#{first_name} #{last_name}" # The equivalent interpolation
=> "Michael Hartl"
```

Note that the final two expressions are equivalent, but I prefer the interpolated version; having to add the single space " " seems a bit awkward.

Printing

To *print* a string, the most commonly used Ruby function is **puts** (pronounced “put ess”, for “put string”):

```
>> puts "foo" # put string
foo
=> nil
```

The **puts** method operates as a *side-effect*: the expression **puts "foo"** prints the string to the screen and then returns literally nothing: **nil** is a special Ruby value for “nothing at all”. (In what follows, I’ll sometimes suppress the **=> nil** part for simplicity.)

Using **puts** automatically appends a newline character `\n` to the output; the related **print** method does not:

```
>> print "foo" # print string (same as puts, but without the newline)
foo=> nil
>> print "foo\n" # Same as puts "foo"
foo
=> nil
```

Single-quoted strings

All the examples so far have used *double-quoted strings*, but Ruby also supports *single-quoted strings*. For many uses, the two types of strings are effectively identical:

```
>> 'foo'           # A single-quoted string
=> "foo"
>> 'foo' + 'bar'
=> "foobar"
```

There's an important difference, though; Ruby won't interpolate into single-quoted strings:

```
>> '#{foo} bar'    # Single-quoted strings don't allow interpolation
=> "\#{foo} bar"
```

Note how the console returns values using double-quoted strings, which requires a backslash to *escape* special characters such as `#`.

If double-quoted strings can do everything that single-quoted strings can do, and interpolate to boot, what's the point of single-quoted strings? They are often useful because they are truly literal, and contain exactly the characters you type. For example, the “backslash” character is special on most systems, as in the literal newline `\n`. If you want a variable to contain a literal backslash, single quotes make it easier:

```
>> '\n'           # A literal 'backslash n' combination
=> "\\n"
```

As with the `#` character in our previous example, Ruby needs to escape the backslash with an additional backslash; inside double-quoted strings, a literal backslash is represented with *two* backslashes. For a small example like this, there's not much savings, but if there are lots of things to

escape it can be a real help:

```
>> 'Newlines (\n) and tabs (\t) both use the backslash character \.'
```

```
=> "Newlines (\n) and tabs (\t) both use the backslash character \\'."
```

4.2.3 Objects and message passing

Everything in Ruby, including strings and even `nil`, is an *object*. We'll see the technical meaning of this in [Section 4.4.2](#), but I don't think anyone ever understood objects by reading the definition in a book; you have to build up your intuition for objects by seeing lots of examples.

It's easier to describe what objects *do*, which is respond to messages. An object like a string, for example, can respond to the message `length`, which returns the number of characters in the string:

```
>> "foobar".length      # Passing the "length" message to a string
```

```
=> 6
```

Typically, the messages that get passed to objects are *methods*, which are functions defined on those objects.⁴ Strings also respond to the `empty?` method:

```
>> "foobar".empty?
```

```
=> false
```

```
>> "".empty?
```

```
=> true
```

Note the question mark at the end of the `empty?` method. This is a Ruby convention indicating that the return value is *boolean*: `true` or `false`. Booleans are especially useful for *control flow*:

```
>> s = "foobar"
>> if s.empty?
>>   "The string is empty"
>> else
>>   "The string is nonempty"
>> end
=> "The string is nonempty"
```

Booleans can also be combined using the `&&` (“and”), `||` (“or”), and `!` (“not”) operators:

```
>> x = "foo"
=> "foo"
>> y = ""
=> ""
>> puts "Both strings are empty" if x.empty? && y.empty?
=> nil
>> puts "One of the strings is empty" if x.empty? || y.empty?
"One of the strings is empty"
=> nil
>> puts "x is not empty" if !x.empty?
"x is not empty"
=> nil
```

Since everything in Ruby is an object, it follows that `nil` is an object, so it too can respond to methods. One example is the `to_s` method that can convert virtually any object to a string:

```
>> nil.to_s
=> ""
```

This certainly appears to be an empty string, as we can verify by *chaining* the messages we pass to `nil`:


```
>> nil.empty?  
NoMethodError: You have a nil object when you didn't expect it!  
You might have expected an instance of Array.  
The error occurred while evaluating nil.empty?  
>> nil.to_s.empty?      # Message chaining  
=> true
```

We see here that the **nil** object doesn't itself respond to the **empty?** method, but **nil.to_s** does.

There's a special method for testing for **nil**-ness, which you might be able to guess:

```
>> "foo".nil?  
=> false  
>> "".nil?  
=> false  
>> nil.nil?  
=> true
```

The code

```
puts "x is not empty" if !x.empty?
```

also shows an alternate use of the **if** keyword: Ruby allows you to write a statement that is evaluated only if the statement following **if** is true. There's a complementary **unless** keyword that works the same way:

```
>> string = "foobar"  
>> puts "The string '#{string}' is nonempty." unless string.empty?
```

```
The string 'foobar' is nonempty.  
=> nil
```

It's worth noting that the `nil` object is special, in that it is the *only* Ruby object that is false in a boolean context, apart from `false` itself:

```
>> if nil  
>>   true  
>> else  
>>   false      # nil is false  
>> end  
=> false
```

In particular, all other Ruby objects are *true*, even 0:

```
>> if 0  
>>   true      # 0 (and everything other than nil and false itself) is true  
>> else  
>>   false  
>> end  
=> true
```

4.2.4 Method definitions

The console allows us to define methods the same way we did with the `home` action from [Listing 3.6](#) or the `full_title` helper from [Listing 4.2](#). (Defining methods in the console is a bit cumbersome, and ordinarily you would use a file, but it's convenient for demonstration purposes.) For example, let's define a function `string_message` that takes a single *argument* and returns a message based on whether the argument is empty or not:

```
>> def string_message(string)
>>   if string.empty?
>>     "It's an empty string!"
>>   else
>>     "The string is nonempty."
>>   end
>> end
=> nil
>> puts string_message("")
It's an empty string!
>> puts string_message("foobar")
The string is nonempty.
```

Note that Ruby functions have an *implicit return*, meaning they return the last statement evaluated—in this case, one of the two message strings, depending on whether the method's argument **string** is empty or not. Ruby also has an explicit return option; the following function is equivalent to the one above:

```
>> def string_message(string)
>>   return "It's an empty string!" if string.empty?
>>   return "The string is nonempty."
>> end
```

The alert reader might notice at this point that the second **return** here is actually unnecessary—being the last expression in the function, the string **"The string is nonempty."** will be returned regardless of the **return** keyword, but using **return** in both places has a pleasing symmetry to it.

4.2.5 Back to the title helper

We are now in a position to understand the **full_title** helper from [Listing 4.2](#):⁵

```

module ApplicationHelper

  # Returns the full title on a per-page basis.
  def full_title(page_title)
    base_title = "Ruby on Rails Tutorial Sample App"
    if page_title.empty?
      base_title
    else
      "#{base_title} | #{page_title}"
    end
  end
end

```

Documentation comment
 # Method definition
 # Variable assignment
 # Boolean test
 # Implicit return
 # String interpolation

These elements—function definition, variable assignment, boolean tests, control flow, and string interpolation—come together to make a compact helper method for use in our site layout. The final element is `module ApplicationHelper`: modules give us a way to package together related methods, which can then be *mixed in* to Ruby classes using `include`. When writing ordinary Ruby, you often write modules and include them explicitly yourself, but in the case of a helper module Rails handles the inclusion for us. The result is that the `full_title` method is automagically available in all our views.

4.3 Other data structures

Although web apps are ultimately about strings, actually *making* those strings requires using other data structures as well. In this section, we'll learn about some Ruby data structures important for writing Rails applications.

4.3.1 Arrays and ranges

An array is just a list of elements in a particular order. We haven't discussed arrays yet in the *Rails Tutorial*, but understanding them gives a good foundation for understanding hashes ([Section 4.3.3](#)) and for aspects of Rails data modeling (such as the `has_many` association seen in [Section 2.3.3](#) and covered more in [Section 10.1.3](#)).

So far we've spent a lot of time understanding strings, and there's a natural way to get from strings to arrays using the `split` method:

```
>> "foo bar baz".split      # Split a string into a three-element array
=> ["foo", "bar", "baz"]
```

The result of this operation is an array of three strings. By default, `split` divides a string into an array by splitting on whitespace, but you can split on nearly anything else as well:

```
>> "fooxbarxbazx".split('x')
=> ["foo", "bar", "baz"]
```

As is conventional in most computer languages, Ruby arrays are *zero-offset*, which means that the first element in the array has index 0, the second has index 1, and so on:

```
>> a = [42, 8, 17]
=> [42, 8, 17]
>> a[0]                # Ruby uses square brackets for array access.
=> 42
>> a[1]
=> 8
>> a[2]
=> 17
>> a[-1]                # Indices can even be negative!
=> 17
```

We see here that Ruby uses square brackets to access array elements. In addition to this bracket notation, Ruby offers synonyms for some commonly accessed elements:⁶

```
>> a                # Just a reminder of what 'a' is
=> [42, 8, 17]
>> a.first
=> 42
>> a.second
=> 8
>> a.last
=> 17
>> a.last == a[-1]   # Comparison using ==
=> true
```

This last line introduces the equality comparison operator `==`, which Ruby shares with many other languages, along with the associated `!=` (“not equal”), etc.:

```
>> x = a.length      # Like strings, arrays respond to the 'length' method.
=> 3
>> x == 3
=> true
>> x == 1
=> false
>> x != 1
=> true
>> x >= 1
=> true
>> x < 1
=> false
```

In addition to **length** (seen in the first line above), arrays respond to a wealth of other methods:

```
>> a
=> [42, 8, 17]
>> a.sort
=> [8, 17, 42]
>> a.reverse
=> [17, 8, 42]
>> a.shuffle
```

```
=> [17, 42, 8]
>> a
=> [42, 8, 17]
```

Note that none of the methods above changes **a** itself. To *mutate* the array, use the corresponding “bang” methods (so-called because the exclamation point is usually pronounced “bang” in this context):

```
>> a
=> [42, 8, 17]
>> a.sort!
=> [8, 17, 42]
>> a
=> [8, 17, 42]
```

You can also add to arrays with the **push** method or its equivalent operator, **<<**:

```
>> a.push(6)           # Pushing 6 onto an array
=> [42, 8, 17, 6]
>> a << 7              # Pushing 7 onto an array
=> [42, 8, 17, 6, 7]
>> a << "foo" << "bar"  # Chaining array pushes
=> [42, 8, 17, 6, 7, "foo", "bar"]
```

This last example shows that you can chain pushes together, and also that, unlike arrays in many other languages, Ruby arrays can contain a mixture of different types (in this case, integers and strings).

Before we saw **split** convert a string to an array. We can also go the other way with the **join** method:

```
>> a
=> [42, 8, 17, 7, "foo", "bar"]
>> a.join           # Join on nothing
=> "428177foobar"
>> a.join(', ')     # Join on comma-space
=> "42, 8, 17, 7, foo, bar"
```

Closely related to arrays are *ranges*, which can probably most easily be understood by converting them to arrays using the `to_a` method:

```
>> 0..9
=> 0..9
>> 0..9.to_a        # Oops, call to_a on 9
NoMethodError: undefined method `to_a' for 9:Fixnum
>> (0..9).to_a      # Use parentheses to call to_a on the range
=> [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

Though `0..9` is a valid range, the second expression above shows that we need to add parentheses to call a method on it.

Ranges are useful for pulling out array elements:

```
>> a = %w[foo bar baz quux]      # Use %w to make a string array.
=> ["foo", "bar", "baz", "quux"]
>> a[0..2]
=> ["foo", "bar", "baz"]
```

Ranges also work with characters:

```
>> ('a'..'e').to_a
=> ["a", "b", "c", "d", "e"]
```


4.3.2 Blocks

Both arrays and ranges respond to a host of methods that accept *blocks*, which are simultaneously one of Ruby's most powerful and most confusing features:

```
>> (1..5).each { |i| puts 2 * i }  
2  
4  
6  
8  
10  
=> 1..5
```

This code calls the **each** method on the range **(1..5)** and passes it the block **{ |i| puts 2 * i }**. The vertical bars around the variable name in **|i|** are Ruby syntax for a block variable, and it's up to the method to know what to do with the block. In this case, the range's **each** method can handle a block with a single local variable, which we've called **i**, and it just executes the block for each value in the range.

Curly braces are one way to indicate a block, but there is a second way as well:

```
>> (1..5).each do |i|  
?>   puts 2 * i  
>> end  
2  
4  
6  
8  
10  
=> 1..5
```

Blocks can be more than one line, and often are. In the *Rails Tutorial* we'll follow the common convention of using curly braces only for short one-line blocks and the **do . end** syntax for longer one-liners and for multi-line blocks:

```
>> (1..5).each do |number|
?>   puts 2 * number
>>   puts '---'
>> end
2
--
4
--
6
--
8
--
10
--
=> 1..5
```

Here I've used **number** in place of **i** just to emphasize that any variable name will do.

Unless you already have a substantial programming background, there is no shortcut to understanding blocks; you just have to see them a lot, and eventually you'll get used to them.² Luckily, humans are quite good at making generalizations from concrete examples; here are a few more, including a couple using the **map** method:

```
>> 3.times { puts "Betelgeuse!" } # 3.times takes a block with no variables.
"Betelgeuse!"
"Betelgeuse!"
"Betelgeuse!"
=> 3
>> (1..5).map { |i| i**2 } # The ** notation is for 'power'.
=> [1, 4, 9, 16, 25]
>> %w[a b c] # Recall that %w makes string arrays.
```

```
=> ["a", "b", "c"]
>> %w[a b c].map { |char| char.upcase }
=> ["A", "B", "C"]
>> %w[A B C].map { |char| char.downcase }
=> ["a", "b", "c"]
```

As you can see, the **map** method returns the result of applying the given block to each element in the array or range.

By the way, we're now in a position to understand the line of Ruby I threw into [Section 1.4.4](#) to generate random subdomains:

```
('a'..'z').to_a.shuffle[0..7].join
```

Let's build it up step-by-step:

```
>> ('a'..'z').to_a                # An alphabet array
=> ["a", "b", "c", "d", "e", "f", "g", "h", "i", "j", "k", "l", "m", "n", "o",
    "p", "q", "r", "s", "t", "u", "v", "w", "x", "y", "z"]
>> ('a'..'z').to_a.shuffle        # Shuffle it.
=> ["c", "g", "l", "k", "h", "z", "s", "i", "n", "d", "y", "u", "t", "j", "q",
    "b", "r", "o", "f", "e", "w", "v", "m", "a", "x", "p"]
>> ('a'..'z').to_a.shuffle[0..7]  # Pull out the first eight elements.
=> ["f", "w", "i", "a", "h", "p", "c", "x"]
>> ('a'..'z').to_a.shuffle[0..7].join # Join them together to make one string.
=> "mznpybuj"
```

4.3.3 Hashes and symbols

Hashes are essentially a generalization of arrays: you can think of hashes as basically like arrays, but not limited to integer indices. (In fact, some languages, especially Perl, sometimes call hashes

associative arrays for this reason.) Instead, hash indices, or *keys*, can be almost any object. For example, we can use strings as keys:

```
>> user = {}                                # {} is an empty hash.
=> {}
>> user["first_name"] = "Michael"           # Key "first_name", value "Michael"
=> "Michael"
>> user["last_name"] = "Hartl"              # Key "last_name", value "Hartl"
=> "Hartl"
>> user["first_name"]                       # Element access is like arrays.
=> "Michael"
>> user                                     # A literal representation of the hash
=> {"last_name"=>"Hartl", "first_name"=>"Michael"}
```

Hashes are indicated with curly braces containing key-value pairs; a pair of braces with no key-value pairs—i.e., `{}`—is an empty hash. It's important to note that the curly braces for hashes have nothing to do with the curly braces for blocks. (Yes, this can be confusing.) Although hashes resemble arrays, one important difference is that hashes don't generally guarantee keeping their elements in a particular order.⁸ If order matters, use an array.

Instead of defining hashes one item at a time using square brackets, it's easy to use a literal representation with keys and values separated by `=>`, called a “hashrocket”:

```
>> user = { "first_name" => "Michael", "last_name" => "Hartl" }
=> {"last_name"=>"Hartl", "first_name"=>"Michael"}
```

Here I've used the usual Ruby convention of putting an extra space at the two ends of the hash—a convention ignored by the console output. (Don't ask me why the spaces are conventional; probably some early influential Ruby programmer liked the look of the extra spaces, and the convention stuck.)

So far we've used strings as hash keys, but in Rails it is much more common to use *symbols* instead. Symbols look kind of like strings, but prefixed with a colon instead of surrounded by quotes. For example, `:name` is a symbol. You can think of symbols as basically strings without all the extra baggage:⁹

```
>> "name".split('')
=> ["n", "a", "m", "e"]
>> :name.split('')
NoMethodError: undefined method `split' for :name:Symbol
>> "foobar".reverse
=> "raboof"
>> :foobar.reverse
NoMethodError: undefined method `reverse' for :foobar:Symbol
```

Symbols are a special Ruby data type shared with very few other languages, so they may seem weird at first, but Rails uses them a lot, so you'll get used to them fast.

In terms of symbols as hash keys, we can define a `user` hash as follows:

```
>> user = { :name => "Michael Hartl", :email => "michael@example.com" }
=> {:name=>"Michael Hartl", :email=>"michael@example.com"}
>> user[:name]           # Access the value corresponding to :name.
=> "Michael Hartl"
>> user[:password]      # Access the value of an undefined key.
=> nil
```

We see here from the last example that the hash value for an undefined key is simply `nil`.

Since it's so common for hashes to use symbols as keys, Ruby 1.9 supports a new syntax just for this special case:

```
>> h1 = { :name => "Michael Hartl", :email => "michael@example.com" }
=> {:name=>"Michael Hartl", :email=>"michael@example.com"}
>> h2 = { name: "Michael Hartl", email: "michael@example.com" }
=> {:name=>"Michael Hartl", :email=>"michael@example.com"}
>> h1 == h2
=> true
```

The second syntax replaces the symbol/hashrocket combination with the name of the key followed by a colon and a value:

```
{ name: "Michael Hartl", email: "michael@example.com" }
```

This construction more closely follows the hash notation in other languages (such as JavaScript) and enjoys growing popularity in the Rails community. Both syntaxes are still in common use, so it's essential to be able to recognize them. Most hashes in the rest of this book use the new notation, which won't work with Ruby 1.8.7 or earlier; if you are using an earlier version of Ruby, you will either have to upgrade to Ruby 1.9 (recommended) or use the old hash notation.

Hash values can be virtually anything, even other hashes, as seen in [Listing 4.6](#).

Listing 4.6. Nested hashes.

```
>> params = {}           # Define a hash called 'params' (short for 'parameters').
=> {}
>> params[:user] = { name: "Michael Hartl", email: "mhartl@example.com" }
=> {:name=>"Michael Hartl", :email=>"mhartl@example.com"}
>> params
=> {:user=>{:name=>"Michael Hartl", :email=>"mhartl@example.com"}}
>> params[:user][:email]
=> "mhartl@example.com"
```

These sorts of hashes-of-hashes, or *nested hashes*, are heavily used by Rails, as we'll see starting in

Section 7.3.

As with arrays and ranges, hashes respond to the **each** method. For example, consider a hash named **flash** with keys for two conditions, **:success** and **:error**:

```
>> flash = { success: "It worked!", error: "It failed." }
=> {:success=>"It worked!", :error=>"It failed."}
>> flash.each do |key, value|
?>   puts "Key #{key.inspect} has value #{value.inspect}"
>> end
Key :success has value "It worked!"
Key :error has value "It failed."
```

Note that, while the **each** method for arrays takes a block with only one variable, **each** for hashes takes two, a *key* and a *value*. Thus, the **each** method for a hash iterates through the hash one key-value *pair* at a time.

The last example uses the useful **inspect** method, which returns a string with a literal representation of the object it's called on:

```
>> puts (1..5).to_a           # Put an array as a string.
1
2
3
4
5
>> puts (1..5).to_a.inspect   # Put a literal array.
[1, 2, 3, 4, 5]
>> puts :name, :name.inspect
name
:name
>> puts "It worked!", "It worked!".inspect
It worked!
"It worked!"
```

By the way, using `inspect` to print an object is common enough that there's a shortcut for it, the `p` function:

```
>> p :name           # Same as 'puts :name.inspect'
:name
```

4.3.4 CSS revisited

It's time now to revisit the line from [Listing 4.1](#) used in the layout to include the cascading style sheets:

```
<%= stylesheet_link_tag "application", :media => "all" %>
```

We are now nearly in a position to understand this. As mentioned briefly in [Section 4.1](#), Rails defines a special function to include stylesheets, and

```
stylesheet_link_tag "application", :media => "all"
```

is a call to this function. But there are two mysteries. First, where are the parentheses? In Ruby, they are optional; these two lines are equivalent:

```
# Parentheses on function calls are optional.
stylesheet_link_tag("application", :media => "all")
stylesheet_link_tag "application", :media => "all"
```


Second, the `:media` argument sure looks like a hash, but where are the curly braces? When hashes are the *last* argument in a function call, the curly braces are optional; these two lines are equivalent:

```
# Curly braces on final hash arguments are optional.
stylesheet_link_tag "application", { :media => "all" }
stylesheet_link_tag "application", :media => "all"
```

So, we see now that the line

```
stylesheet_link_tag "application", :media => "all"
```

calls the `stylesheet_link_tag` function with two arguments: a string, indicating the path to the stylesheet, and a hash, indicating the media type. Because of the `<%= %>` brackets, the results are inserted into the template by ERb, and if you view the source of the page in your browser you should see the HTML needed to include a stylesheet ([Listing 4.7](#)). (You may see some extra things, like `?body=1`, after the CSS filenames. These are inserted by Rails to ensure that browsers reload the CSS when it changes on the server.)

Listing 4.7. The HTML source produced by the CSS includes.

```
<link href="/assets/application.css" media="all" rel="stylesheet"
type="text/css" />
```

If you actually view the CSS file by navigating to <http://localhost:3000/assets/application.css>, you'll see that (apart from some comments) it is empty. We'll set about changing this in [Chapter 5](#).

4.4 Ruby classes

We've said before that everything in Ruby is an object, and in this section we'll finally get to define some of our own. Ruby, like many object-oriented languages, uses *classes* to organize methods; these classes are then *instantiated* to create objects. If you're new to object-oriented programming, this may sound like gibberish, so let's look at some concrete examples.

4.4.1 Constructors

We've seen lots of examples of using classes to instantiate objects, but we have yet to do so explicitly. For example, we instantiated a string using the double quote characters, which is a *literal constructor* for strings:

```
>> s = "foobar"      # A literal constructor for strings using double quotes
=> "foobar"
>> s.class
=> String
```

We see here that strings respond to the method `class`, and simply return the class they belong to.

Instead of using a literal constructor, we can use the equivalent *named constructor*, which involves calling the `new` method on the class name:^{[10](#)}

```
>> s = String.new("foobar")  # A named constructor for a string
=> "foobar"
>> s.class
=> String
>> s == "foobar"
=> true
```

This is equivalent to the literal constructor, but it's more explicit about what we're doing.

Arrays work the same way as strings:

```
>> a = Array.new([1, 3, 2])
=> [1, 3, 2]
```

Hashes, in contrast, are different. While the array constructor **Array.new** takes an initial value for the array, **Hash.new** takes a *default* value for the hash, which is the value of the hash for a nonexistent key:

```
>> h = Hash.new
=> {}
>> h[:foo]           # Try to access the value for the nonexistent key :foo.
=> nil
>> h = Hash.new(0)   # Arrange for nonexistent keys to return 0 instead of nil.
=> {}
>> h[:foo]
=> 0
```

When a method gets called on the class itself, as in the case of **new**, it's called a *class method*. The result of calling **new** on a class is an object of that class, also called an *instance* of the class. A method called on an instance, such as **length**, is called an *instance method*.

4.4.2 Class inheritance

When learning about classes, it's useful to find out the *class hierarchy* using the **superclass** method:

```
>> s = String.new("foobar")
=> "foobar"
>> s.class           # Find the class of s.
=> String
```

```
>> s.class.superclass           # Find the superclass of String.  
=> Object  
>> s.class.superclass.superclass # Ruby 1.9 uses a new BasicObject base class  
=> BasicObject  
>> s.class.superclass.superclass.superclass  
=> nil
```

A diagram of this inheritance hierarchy appears in [Figure 4.1](#). We see here that the superclass of **String** is **Object** and the superclass of **Object** is **BasicObject**, but **BasicObject** has no superclass. This pattern is true of every Ruby object: trace back the class hierarchy far enough and every class in Ruby ultimately inherits from **BasicObject**, which has no superclass itself. This is the technical meaning of “everything in Ruby is an object”.




Figure 4.1: The inheritance hierarchy for the **String** class.

To understand classes a little more deeply, there’s no substitute for making one of our own. Let’s make a **Word** class with a **palindrome?** method that returns **true** if the word is the same spelled forward and backward:

```
>> class Word
>>   def palindrome?(string)
>>     string == string.reverse
>>   end
>> end
=> nil
```

We can use it as follows:

```
>> w = Word.new           # Make a new Word object.
=> #<Word:0x22d0b20>
>> w.palindrome?("foobar")
=> false
>> w.palindrome?("level")
=> true
```

If this example strikes you as a bit contrived, good; this is by design. It's odd to create a new class just to create a method that takes a string as an argument. Since a word is *a* string, it's more natural to have our **Word** class *inherit* from **String**, as seen in [Listing 4.8](#). (You should exit the console and re-enter it to clear out the old definition of **Word**.)

Listing 4.8. Defining a **Word** class in the console.

```
>> class Word < String      # Word inherits from String.
>>   # Returns true if the string is its own reverse.
>>   def palindrome?
>>     self == self.reverse  # self is the string itself.
>>   end
>> end
=> nil
```

Here **Word < String** is the Ruby syntax for inheritance (discussed briefly in [Section 3.1.2](#)), which ensures that, in addition to the new **palindrome?** method, words also have all the same

methods as strings:

```
>> s = Word.new("level")    # Make a new Word, initialized with "level".
=> "level"
>> s.palindrome?           # Words have the palindrome? method.
=> true
>> s.length                 # Words also inherit all the normal string methods.
=> 5
```

Since the **Word** class inherits from **String**, we can use the console to see the class hierarchy explicitly:

```
>> s.class
=> Word
>> s.class.superclass
=> String
>> s.class.superclass.superclass
=> Object
```

This hierarchy is illustrated in [Figure 4.2](#).




Figure 4.2: The inheritance hierarchy for the (non-built-in) **Word** class from [Listing 4.8](#).

In [Listing 4.8](#), note that checking that the word is its own reverse involves accessing the word inside the **Word** class. Ruby allows us to do this using the **self** keyword: inside the **Word** class, **self** is the object itself, which means we can use

```
self == self.reverse
```

to check if the word is a palindrome.^{[11](#)}

4.4.3 Modifying built-in classes

While inheritance is a powerful idea, in the case of palindromes it might be even more natural to add the **palindrome?** method to the **String** class itself, so that (among other things) we can call

`palindrome?` on a string literal, which we currently can't do:

```
>> "level".palindrome?  
NoMethodError: undefined method `palindrome?' for "level":String
```

Somewhat amazingly, Ruby lets you do just this; Ruby classes can be *opened* and modified, allowing ordinary mortals such as ourselves to add methods to them:^{[12](#)}

```
>> class String  
>>   # Returns true if the string is its own reverse.  
>>   def palindrome?  
>>     self == self.reverse  
>>   end  
>> end  
=> nil  
>> "deified".palindrome?  
=> true
```

(I don't know which is cooler: that Ruby lets you add methods to built-in classes, or that `"deified"` is a palindrome.)

Modifying built-in classes is a powerful technique, but with great power comes great responsibility, and it's considered bad form to add methods to built-in classes without having a *really* good reason for doing so. Rails does have some good reasons; for example, in web applications we often want to prevent variables from being *blank*—e.g., a user's name should be something other than spaces and other [whitespace](#)—so Rails adds a `blank?` method to Ruby. Since the Rails console automatically includes the Rails extensions, we can see an example here (this won't work in plain `irb`):

```
>> "".blank?  
=> true  
>> " ".empty?
```



```
=> false
>> "   ".blank?
=> true
>> nil.blank?
=> true
```

We see that a string of spaces is not *empty*, but it is *blank*. Note also that `nil` is blank; since `nil` isn't a string, this is a hint that Rails actually adds `blank?` to `String`'s base class, which (as we saw at the beginning of this section) is `Object` itself. We'll see some other examples of Rails additions to Ruby classes in [Section 8.2.1](#).

4.4.4 A controller class

All this talk about classes and inheritance may have triggered a flash of recognition, because we have seen both before, in the `StaticPages` controller ([Listing 3.15](#)):

```
class StaticPagesController < ApplicationController

  def home
  end

  def help
  end

  def about
  end
end
```

You're now in a position to appreciate, at least vaguely, what this code means:

`StaticPagesController` is a class that inherits from `ApplicationController`, and comes equipped with `home`, `help`, and `about` methods. Since each Rails console session loads the local Rails environment, we can even create a controller explicitly and examine its class hierarchy: ¹³

```
>> controller = StaticPagesController.new
=> #<StaticPagesController:0x22855d0>
>> controller.class
=> StaticPagesController
>> controller.class.superclass
=> ApplicationController
>> controller.class.superclass.superclass
=> ActionController::Base
>> controller.class.superclass.superclass.superclass
=> ActionController::Metal
>> controller.class.superclass.superclass.superclass.superclass
=> AbstractController::Base
>> controller.class.superclass.superclass.superclass.superclass.superclass
=> Object
```

A diagram of this hierarchy appears in [Figure 4.3](#).

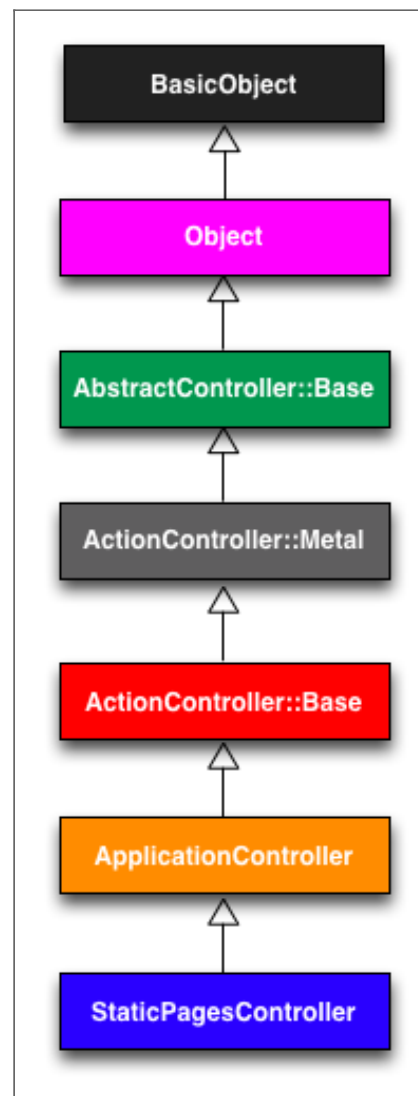


Figure 4.3: The inheritance hierarchy for the StaticPages controller.

We can even call the controller actions inside the console, which are just methods:

```
>> controller.home  
=> nil
```

Here the return value is `nil` because the `home` action is blank.

But wait—actions don't have return values, at least not ones that matter. The point of the `home` action, as we saw in [Chapter 3](#), is to render a web page, not to return a value. And I sure don't remember ever calling `StaticPagesController.new` anywhere. What's going on?

What's going on is that Rails is *written in* Ruby, but Rails isn't Ruby. Some Rails classes are used like ordinary Ruby objects, but some are just [grist](#) for Rails' magic mill. Rails is [sui generis](#), and should be studied and understood separately from Ruby. This is why, if your principal programming interest is writing web applications, I recommend learning Rails first, then learning Ruby, then looping back to Rails.

4.4.5 A user class

We end our tour of Ruby with a complete class of our own, a `User` class that anticipates the User model coming up in [Chapter 6](#).

So far we've entered class definitions at the console, but this quickly becomes tiresome; instead, create the file `example_user.rb` in your application root directory and fill it with the contents of [Listing 4.9](#).

Listing 4.9. Code for an example user.

`example_user.rb`

```
class User
  attr_accessor :name, :email

  def initialize(attributes = {})
    @name = attributes[:name]
    @email = attributes[:email]
  end
end
```

```
end

def formatted_email
  "#{@name} <#{@email}>"
end
end
```

There's quite a bit going on here, so let's take it step by step. The first line,

```
attr_accessor :name, :email
```

creates *attribute accessors* corresponding to a user's name and email address. This creates “getter” and “setter” methods that allow us to retrieve (get) and assign (set) `@name` and `@email` instance variables, which were mentioned briefly in [Section 2.2.2](#). In Rails, the principal importance of instance variables is that they are automatically available in the views, but in general they are used for variables that need to be available throughout a Ruby class. (We'll have more to say about this in a moment.) Instance variables always begin with an `@` sign, and are `nil` when undefined.

The first method, `initialize`, is special in Ruby: it's the method called when we execute `User.new`. This particular `initialize` takes one argument, `attributes`:

```
def initialize(attributes = {})
  @name = attributes[:name]
  @email = attributes[:email]
end
```

Here the `attributes` variable has a *default value* equal to the empty hash, so that we can define a user with no name or email address (recall from [Section 4.3.3](#) that hashes return `nil` for nonexistent keys, so `attributes[:name]` will be `nil` if there is no `:name` key, and similarly for

```
attributes[:email]).
```

Finally, our class defines a method called `formatted_email` that uses the values of the assigned `@name` and `@email` variables to build up a nicely formatted version of the user's email address using string interpolation ([Section 4.2.2](#)):

```
def formatted_email
  "#{@name} <#{@email}>"
end
```

Because `@name` and `@email` are both instance variables (as indicated with the `@` sign), they are automatically available in the `formatted_email` method.

Let's fire up the console, `require` the example user code, and take our User class out for a spin:

```
>> require './example_user'      # This is how you load the example_user code.
=> ["User"]
>> example = User.new
=> #<User:0x224ceec @email=nil, @name=nil>
>> example.name                   # nil since attributes[:name] is nil
=> nil
>> example.name = "Example User"   # Assign a non-nil name
=> "Example User"
>> example.email = "user@example.com" # and a non-nil email address
=> "user@example.com"
>> example.formatted_email
=> "Example User <user@example.com>"
```

Here the `'.'` is Unix for “current directory”, and `'./example_user'` tells Ruby to look for an example user file relative to that location. The subsequent code creates an empty example user and then fills in the name and email address by assigning directly to the corresponding attributes (assignments made possible by the `attr_accessor` line in [Listing 4.9](#)). When we write

```
example.name = "Example User"
```

Ruby is setting the `@name` variable to `"Example User"` (and similarly for the `email` attribute), which we then use in the `formatted_email` method.

Recalling from [Section 4.3.4](#) we can omit the curly braces for final hash arguments, we can create another user by passing a hash to the `initialize` method to create a user with pre-defined attributes:

```
>> user = User.new(name: "Michael Hartl", email: "mhartl@example.com")  
=> #<User:0x225167c @email="mhartl@example.com", @name="Michael Hartl">  
>> user.formatted_email  
=> "Michael Hartl <mhartl@example.com>"
```

We will see starting in [Chapter 7](#) that initializing objects using a hash argument, a technique known as *mass assignment*, is common in Rails applications.

4.5 Conclusion

This concludes our overview of the Ruby language. In [Chapter 5](#), we'll start putting it to good use in developing the sample application.

We won't be using the `example_user.rb` file from [Section 4.4.5](#), so I suggest removing it:

```
$ rm example_user.rb
```

Then commit the other changes to the main source code repository:

```
$ git add .  
$ git commit -m "Add a full_title helper"
```

4.6 Exercises

1. By replacing the question marks in [Listing 4.10](#) with the appropriate methods, combine `split`, `shuffle`, and `join` to write a function that shuffles the letters in a given string.
2. Using [Listing 4.11](#) as a guide, add a `shuffle` method to the `String` class.
3. Create three hashes called `person1`, `person2`, and `person3`, with first and last names under the keys `:first` and `:last`. Then create a `params` hash so that `params[:father]` is `person1`, `params[:mother]` is `person2`, and `params[:child]` is `person3`. Verify that, for example, `params[:father][:first]` has the right value.
4. Find an online version of the Ruby API and read about the `Hash` method `merge`.
5. Find and follow the Ruby Koans to reach Ruby enlightenment.

Listing 4.10. Skeleton for a string shuffle function.

```
>> def string_shuffle(s)  
>>   s.split('').??.?  
>> end  
=> nil  
>> string_shuffle("foobar")
```

Listing 4.11. Skeleton for a `shuffle` method attached to the `String` class.


```
>> class String
>>   def shuffle
>>     self.split('').??.?
>>   end
>> end
=> nil
>> "foobar".shuffle
```

[« Chapter 3 Mostly static pages](#)

[Chapter 5 Filling in the layout »](#)

1. If a helper is specific to a particular controller, you should put it in the corresponding helper file; for example, helpers for the StaticPages controller generally go in `app/helpers/static_pages_helper.rb`. In our case, we expect the `full_title` helper to be used on all the site's pages, and Rails has a special helper file for this case: `app/helpers/application_helper.rb`. ↑
2. For more on the origins of “foo” and “bar”—and, in particular, the possible *non*-relation of “foobar” to “FUBAR”—see the [Jargon File entry on “foo”](#). ↑
3. Programmers familiar with Perl or PHP should compare this to the automatic interpolation of dollar sign variables in expressions like `"foo $bar"`. ↑
4. Apologies in advance for switching haphazardly between *function* and *method* throughout this chapter; in Ruby, they're the same thing: all methods are functions, and all functions are methods, because everything is an object. ↑
5. Well, there will still be *one* thing left that we don't understand, which is how Rails ties this all together: mapping URIs to actions, making the `full_title` helper available in views, etc. This is an interesting subject, and I encourage you to investigate it further, but knowing exactly *how* Rails works is not necessary when *using* Rails. (For a deeper understanding, I recommend [The Rails 3 Way](#) by Obie Fernandez.) ↑

6. The **second** method used here isn't currently part of Ruby itself, but rather is added by Rails. It works in this case because the Rails console automatically includes the Rails extensions to Ruby. ↑
7. Programming experts, on the other hand, might benefit from knowing that blocks are *closures*, which are one-shot anonymous functions with data attached. ↑
8. Ruby 1.9 actually guarantees that hashes keep their elements in the same order entered, but it would be unwise ever to count on a particular ordering. ↑
9. As a result of having less baggage, symbols are easier to compare to each other; strings need to be compared character by character, while symbols can be compared all in one go. This makes them ideal for use as hash keys. ↑
10. These results will vary based on the version of Ruby you are using. This example assumes you are using Ruby 1.9.3. ↑
11. For more on Ruby classes and the **self** keyword, see the [RailsTips](#) post "[Class and Instance Variables in Ruby](#)". ↑
12. For those familiar with JavaScript, this functionality is comparable to using a built-in class prototype object to augment the class. (Thanks to reader [Erik Eldridge](#) for pointing this out.) ↑
13. You don't have to know what each class in this hierarchy does. *I* don't know what they all do, and I've been programming in Ruby on Rails since 2005. This means either that (a) I'm grossly incompetent or (b) you can be a skilled Rails developer without knowing all its innards. I hope for both our sakes that it's the latter. ↑