

CG2271

Real-Time Operating Systems

Lecture 2

I/O Systems

colintan@nus.edu.sg



NUS
National University
of Singapore

School *of* Computing

Learning Objectives

- **To understand and know what input/output (I/O) systems are, and how to program them.**
- **WHY:**
 - Accessing I/O is fundamental to real-time systems.
 - ✓ **Reading sensors to get information about the environment.**
 - ✓ **Writing to actuators to manipulate the environment.**
 - Much of microcontroller programming is about reading sensors, performing computations and writing the results to actuators.

Introduction

- **In the last lecture we looked briefly at what is in a real-time system.**
 - Sensors on the input end.
 - Controllers in the middle.
 - Actuators on the output end.
- **In this lecture we will look in more detail at these ideas:**
 - Input/Output Systems.
 - Input/Output Programming Techniques.
 - Basic Concepts.

I/O Systems

SOME PRELIMINARIES

Assembly Language

- It's easier to explain some concepts in assembly language, and assembly language is essential for some tasks.

- Assembly language: This is a very low-level processor dependent language that actually manipulates the hardware on the processor.

- ✓ Example instructions include ADD, SUB, MUL, DIV, CMP, BLT, BGT, BEQ, JMP etc. E.g. to do “for(i=0; i<20; i++)”, we have:

- LI R0, 0*

- LI R1, 20*

- Loop: ADDI R0, R0, 1*

- CMP R0, R2*

- BLT Loop*

Assembly Language

- **Registers:** A microprocessor has special locations called “registers” to store data to be operated on, and to store results. A typical CPU has between 4 and 32 registers, usually labeled R0, R1, .. R31. We have just seen an example on the previous page. 😊
- **Interrupts:** Sometimes hardware needs to get the CPU’s attention. It issues an “interrupt” causing the CPU to suspend whatever it’s doing and attend to the hardware.

I/O Systems

COMMUNICATION WITH EXTERNAL SYSTEMS

Communicating with External Devices

- **For a processor to do anything useful, it must be able to communicate with the outside world.**

- Some examples of external devices include:

- ✓ **Input devices:**

- Temperature sensors, light sensors, skid sensors, pitot tubes + static ports, etc etc*

- ✓ **Output devices:**

- piezo-electric alarms, LCD/LED displays, actuators, servos, etc etc.*

Communicating with External Devices


Desktops/Servers vs. Microcontrollers

- **I/O design in desktop and server class machines are different from microcontroller class machines.**
 - **Desktop and Servers (including Notebooks)**
 - ✓ **General purpose, meant to “interface” with a large and (at design time) unknown set of devices.**
 - ✓ **Design is much more open-ended, not much built into the processor.**
 - Consists just of “address”, “data” and “control” lines.*
 - ✓ **A “bridge” circuit is instead designed around the processor to provide I/O options.**
 - Generally limited to PCI, PCMCIA (notebooks), AGP/PCIe (specialized forms of PCI for video cards) and USB.*
 - ✓ **This open-endedness can lead to much more complicated designs in the devices to be attached.**

Communicating with External Devices

Desktops/Servers vs. Microcontrollers

- **Desktops and Servers:**



Communicating with External Devices

Desktops/Servers vs. Microcontrollers

- **I/O design in desktop and server class machines are different from microcontroller class machines.**
 - **Microcontrollers**
 - ✓ **Generally a larger, but more specialized set of input/output options.**

See list on next slide.
 - ✓ **These are BUILT INTO the processor, rather than on a circuit surrounding the processor.**
 - ✓ **This results in simpler interfaces with the external devices. 😊**

Communicating with External Devices

Desktops/Servers vs. Microcontrollers

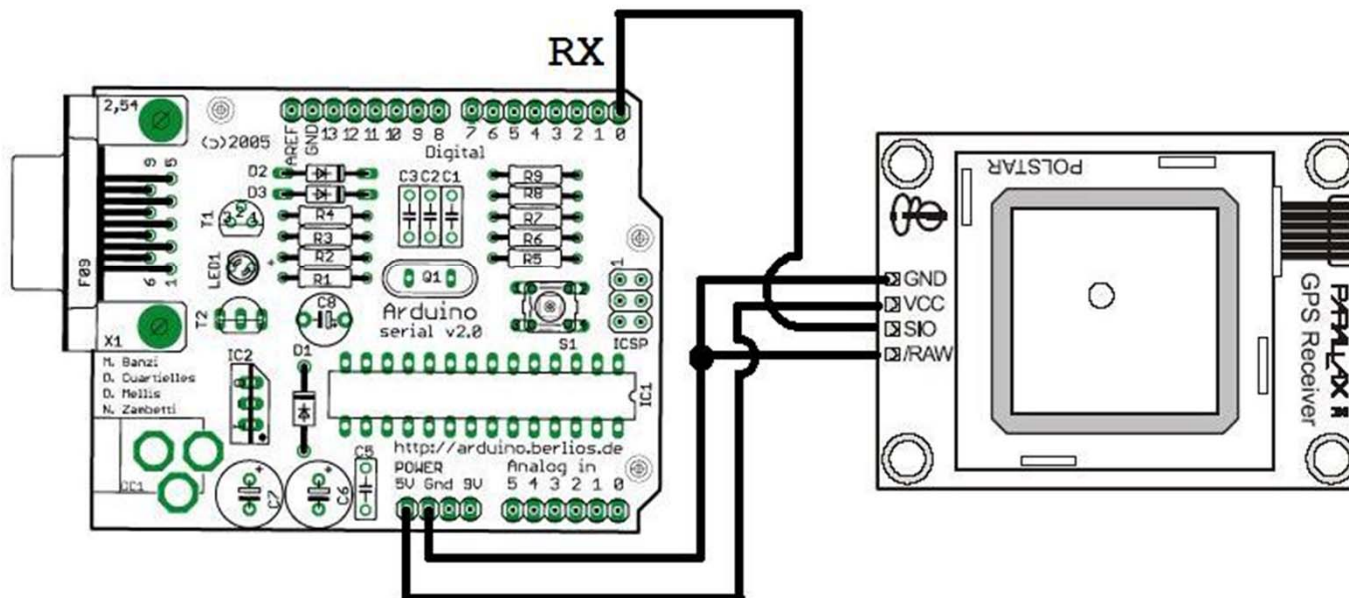
- Typical I/O “ports” in a microcontroller.
 - ✓ **I2C/SPI** – For communicating with other integrated circuits.
SPI is used for communicating with SD cards.
 - ✓ **CAN** – For communicating across many different devices in a vehicle.
 - ✓ **Analog-to-digital converters** – For converting between voltages and digital values.
 - ✓ **PWM ports** – For converting from digital to analog values.
 - ✓ **Ethernet ports** – For communicating across a local area network.
 - ✓ **USB ports** – For communicating with PCs.
 - ✓ ...

Communicating with External Devices

Desktops/Servers vs. Microcontrollers

- Microcontrollers:**

(PCINT14/RESET) PC6	1	28	PC5 (ADC5/SCL/PCINT13)
(PCINT16/RXD) PD0	2	27	PC4 (ADC4/SDA/PCINT12)
(PCINT17/TXD) PD1	3	26	PC3 (ADC3/PCINT11)
(PCINT18/INT0) PD2	4	25	PC2 (ADC2/PCINT10)
(PCINT19/OC2B/INT1) PD3	5	24	PC1 (ADC1/PCINT9)
(PCINT20/XCK/T0) PD4	6	23	PC0 (ADC0/PCINT8)
VCC	7	22	GND
GND	8	21	AREF
(PCINT6/XTAL1/TOSC1) PB6	9	20	AVCC
(PCINT7/XTAL2/TOSC2) PB7	10	19	PB5 (SCK/PCINT5)
(PCINT21/OC0B/T1) PD5	11	18	PB4 (MISO/PCINT4)
(PCINT22/OC0A/AIN0) PD6	12	17	PB3 (MOSI/OC2A/PCINT3)
(PCINT23/AIN1) PD7	13	16	PB2 (SS/OC1B/PCINT2)
(PCINT0/CLKO/ICP1) PB0	14	15	PB1 (OC1A/PCINT1)



Communicating with External Devices

- **To communicate with these devices, there must be:**
 - **A channel for communication**
 - ✓ **Typically the address, data and control buses.**
 - Address bus: Electrical lines that contain the addresses of memory or devices to be accessed.*
 - Data bus: Electrical lines that carry data to and from memory or other devices.*
 - Control bus: Electrical lines that tell the memory or device to either READ the data on the bus or WRITE data to the bus.*
 - **An agreed method for communication**
 - ✓ **A “protocol” that specifies which signal lines are to be de/asserted, how many clock cycles should go by before data appears etc.**
 - ✓ **This is usually specified in a “timing diagram”.**

I/O Systems

TYPES OF I/O PORTS

Types of I/O Ports

- **Two main types of I/O ports in DESKTOP and SERVER class systems.**
 - More common in large desktop and server systems:
 - ✓ **Memory mapped I/O**
 - ✓ **Direct mapped I/O**
- **Micro-controllers are designed a little differently.**
- **We will look at DESKTOP and SERVER class systems first.**

Communicating with External Devices

- **Typically every device connected to the CPU is given a unique identifier:**
 - This serves as the “address” (memory mapped) or “I/O port” (direct mapped) of the device.
 - A CPU will write to this address or I/O port to send data to the device, and read from it to get data from the device.
 - ✓ **Similar in idea to reading/writing memory.**

Types of I/O Ports

Desktops and Servers

- **Two main types:**
 - **Memory mapped I/O:**
 - ✓ **I/O identification numbers come from the normal memory addressing range of the CPU**

E.g. memory locations 8192, 8193, 8194 and 8195 are to identify devices that read from temperature sensors 0 and 1, activate an alarm and call the fire department respectively.
 - ✓ **The CPU sees the devices as part of its memory**

Desktops and Servers

Memory Mapped I/O

- **Memory Mapped I/O example:**
 - Read from temperature sensors 0 and 1 and write to memory locations 400 and 401

```
LI R0, 8192 ; Address of sensor 0
LW R0, (R0) ; Read from temp sensor 0
LI R1, 400 ; Address to write to
SW R0, (R1) ; Write temperature
LI R0, 8193 ; Address of sensor 1
LW R0, (R0) ; Read from temp sensor 1
ADDI R1, R1, 1 ; Increment R1 to 401
SW R0, (R1) ; Write temperature to 401
```

Desktops and Servers

Memory Mapped I/O

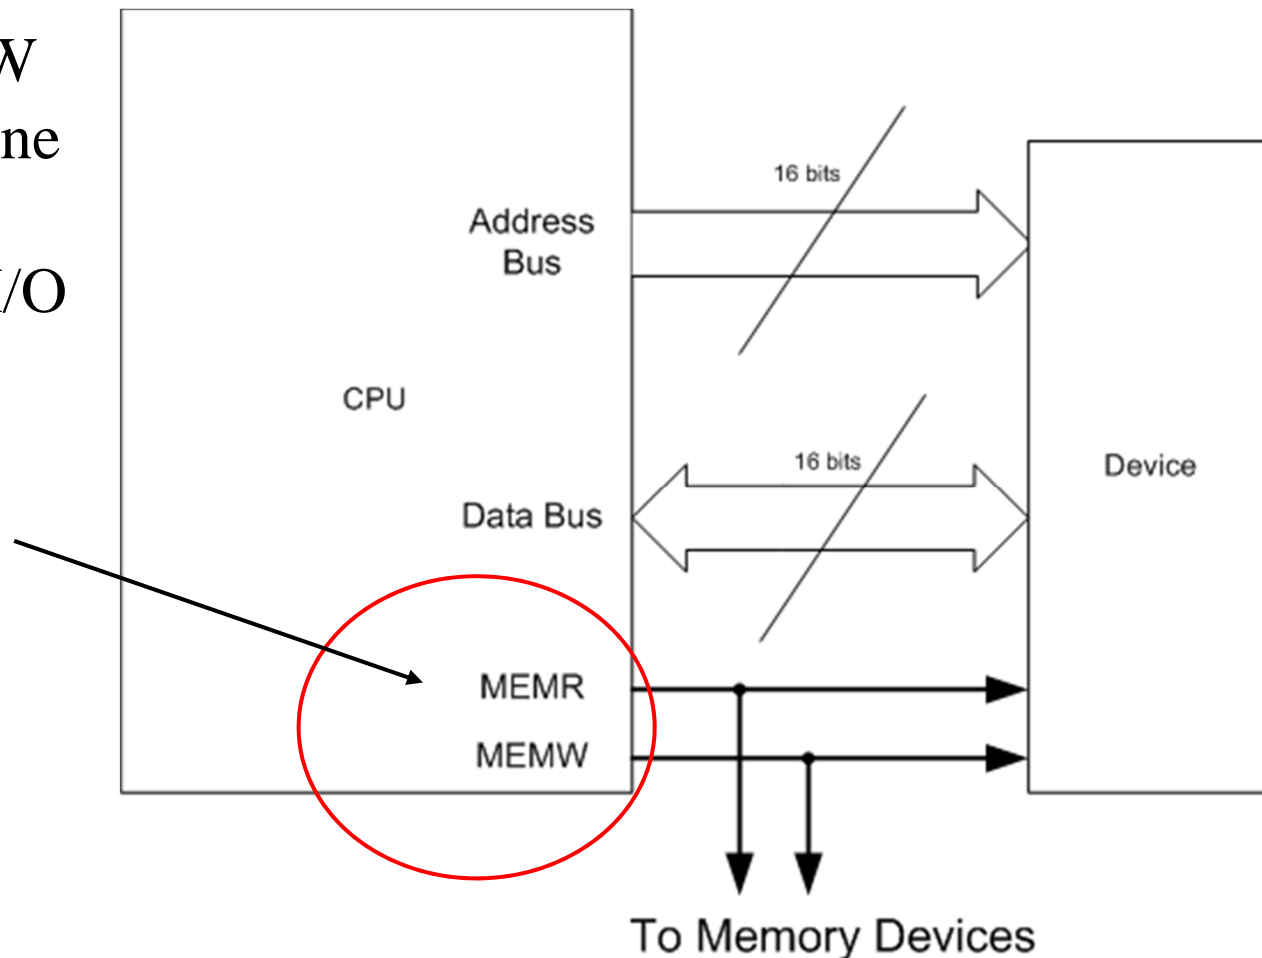
- **Important points:**

- The port numbers for reading sensors 1 and 2 (addresses 8192 and 8193) are treated exactly the same way as memory addresses (e.g. 400 and 401).
- Obviously memory ranges reserved for I/O ports *cannot* be used as normal memory.
 - ✓ **The CPU actually loses some memory addressing range.**
- It is however a simple and attractive scheme.

Desktops and Servers

Memory Mapped I/O

Single MEMW
and MEMR line
shared with
memory and I/O
device



Desktops and Servers

Direct Mapped I/O

- **Second type is Direct Mapped (or Isolated) I/O:**
 - The addressing range of the I/O devices is completely separate from memory.
 - This means that the CPU does not lose any memory addressing range at all
 - AND I/O devices potentially get an equally large addressing range!

Desktops and Servers

Direct Mapped I/O

- **BUT:**
 - CPUs supporting direct mapped I/O must provide:
 - ✓ **Two separate sets of READ/WRITE signals:**
 - One set for memory*
MEMW and MEMR
 - One set of I/O*
IOW and IOR

Desktops and Servers

Direct Mapped I/O

- ✓ **New commands to read/write the I/O devices**

Can no longer use LW/SW or similar commands.

These would cause you to read or write memory instead of I/O devices!

- ✓ **Typically these instructions take the form of:**

*IN R0, portnum ; Read the device at portnum and
write to register R0*

OUT R0, potrnum ; Write R0 to the device at portnum

Desktops and Servers

Direct Mapped I/O

- **Direct Mapped I/O example:**

- Read from temperature sensors at port numbers 8192 and 8193
- Write to memory locations 8192 and 8193

```
IN R0, 8192      ; Read sensor 1
LI R1, 8192
SW R0, (R1)      ; Write to memory
                  location 8192
ADDI R1, R1, 1   ; Inc. to location 8193
IN R0, 8193      ; Read sensor 2
SW R0, (R1)      ; Write to memory location
                  ; 8193
```

Desktops and Servers


Direct Mapped I/O

- **Notice that:**
 - It does not matter that the sensors were at port numbers 8192 and 8193:
 - ✓ **We can still read/write memory locations 8192 and 8193 because their addressing ranges are separate!**
 - We read/write memory using standard LW and SW instructions
 - We read/write I/O devices using special IN and OUT instructions.

Desktops and Servers

Direct Mapped I/O

MEMW and MEMR
lines only for memory
devices



Separate IOW and
IOR lines for I/O
devices

Microcontrollers


Register Mapped I/O

- **Register-mapped I/O is a variation of memory-mapped I/O.**
 - In memory mapped I/O, any location that is not used to store data can be used for I/O.
 - ✓ The corresponding device is activated when the “decoder” circuit matches the address on the address bus with the device’s ID.
 - In register-mapped I/O, the memory locations that are used for I/O is fixed.
 - ✓ These fixed locations are typically called “registers”, which is different from CPU registers that you are used to.
 - ✓ Suitable for microcontrollers as the set of peripherals is usually fixed.

Timing Diagrams

- Typically a timing diagram specifies (in seconds) how long an event is supposed to occur after the occurrence of another event.
 - E.g. to write to a particular RAM device:
 - ✓ Place address and data onto address and data buses respectively.
 - ✓ After at least 25 ns, assert the WRITE signal (T_{WAIT1})
 - ✓ Hold for at least 100 ns (T_{HOLD})
 - ✓ Deassert WRITE signal, remove all other signals
 - ✓ Wait for at least 50 ns before accessing device again (T_{WAIT2})

Timing Diagrams



Timing Diagrams

- Actual bus signals, as seen on a digital bus analyzer:



I/O Systems

I/O PROGRAMMING

I/O Programming

- **The 3 types of I/O ports shown earlier correspond to how the hardware is designed.**
- **On the software side, there are 3 ways to access the I/O devices.**
 - These can work with any of the 3 types of I/O ports shown earlier.
 - The 3 ways to access I/O are:
 - ✓ **Programmed I/O**
 - ✓ **Interrupt-driven I/O**
 - ✓ **Direct Memory Access**

I/O Programming

Programmed I/O

- **This is the simplest form of I/O programming – The CPU does all the work.**
- **Consider a process that prints “ABCDEFGH” on the printer:**
 - User process acquires the printer by making a system call. This call returns an error or blocks if the printer is busy, until it is available.
 - The OS kernel then copies the string to be printed into its own buffer.


I/O Programming

Programmed I/O

- The OS then copies character by character onto the printer's latch, and the printer prints it out.
 - **Copy the first character and advance the buffer's pointer.**
 - **Check that the printer is ready for the next character. If not, wait.**
 - *This is called “busy-waiting” or “polling”.*
 - **Copy the next character. Repeat until buffer is empty.**

I/O Programming

Programmed I/O



```

copy_from_user(buffer, p, count);
for (i = 0; i < count; i++) {
    while (*printer_status_reg != READY) ;
    *printer_data_register = p[i];
}
return_to_user();

```

```

/* p is the kernel bufer */
/* loop on every character */
/* loop until ready */
/* output one character */

```

I/O Programming

Programmed I/O

- **Issue:**
 - It takes perhaps 10ms to print a character.
 - During this time, the CPU will be busy-waiting until the printer is done printing.
 - On a 96MHz processor like the ARM STR912, this is equivalent to wasting 960,000 instruction cycles!
 - ✓ **Assumption is that 1 instruction completes in 1 cycle. Realistic for a RISC processor.**

I/O Programming

Interrupt I/O

- **Since 960,000 instructions is a large number, why not let the CPU do other stuff while the printer is busy?**
 - After the string is copied to kernel space, the OS will send a character to the printer, then switch to a task.
 - When the printer is done, it will interrupt the CPU by asserting one of the “interrupt request” (IRQ) lines on the CPU.
 - This triggers a software routine called an “interrupt handler” , “interrupt service routine” or “interrupt service procedure” which will then load the next character. We will use the term “ISR”.

I/O Programming

Interrupt I/O

```
copy_from_user(buffer, p, count);  
enable_interrupts( );  
while (*printer_status_reg != READY) ;  
*printer_data_register = p[0];  
scheduler( );
```

(a)

```
if (count == 0) {  
    unblock_user( );  
} else {  
    *printer_data_register = p[i];  
    count = count - 1;  
    i = i + 1;  
}  
acknowledge_interrupt( );  
return_from_interrupt( );
```

(b)

I/O Programming

Interrupt I/O

- **Issue:**
 - The printer interrupts the CPU every time it prints a character!
 - Each interrupt has overheads:
 - ✓ CPU must stop what it is doing, look up a vector table for the address of the appropriate ISR, then hand control to the ISR.
 - ✓ ISR must save the registers, process the interrupt, and restore the registers.
 - ✓ CPU then hands control back to the interrupted process.

I/O Programming

Interrupt I/O

- **There are typically several hundred thousand to several million characters in a document!**
 - This means triggering millions of interrupts to print each document.
 - The overheads can be sizeable.
- **Would be ideal if we could just dump the document on the printer and worry only when it has finished printing everything!**
 - This is where Direct Memory Access comes in.

I/O Programming

Direct Memory Access.

- **In Direct Memory Access (DMA):**
 - A separate processor called a DMA Controller (DMAC) is used to transfer data between a device and memory.
- **In a DMA READ:**
 - The user process initiates the READ by making an OS system call, providing:
 - ✓ **Information on which device to read. In the case of a disk drive, the user process will provide information of which part of the disk (side number, track number, starting block number) to read.**
 - ✓ **Information on how many bytes (or blocks of bytes) to read.**

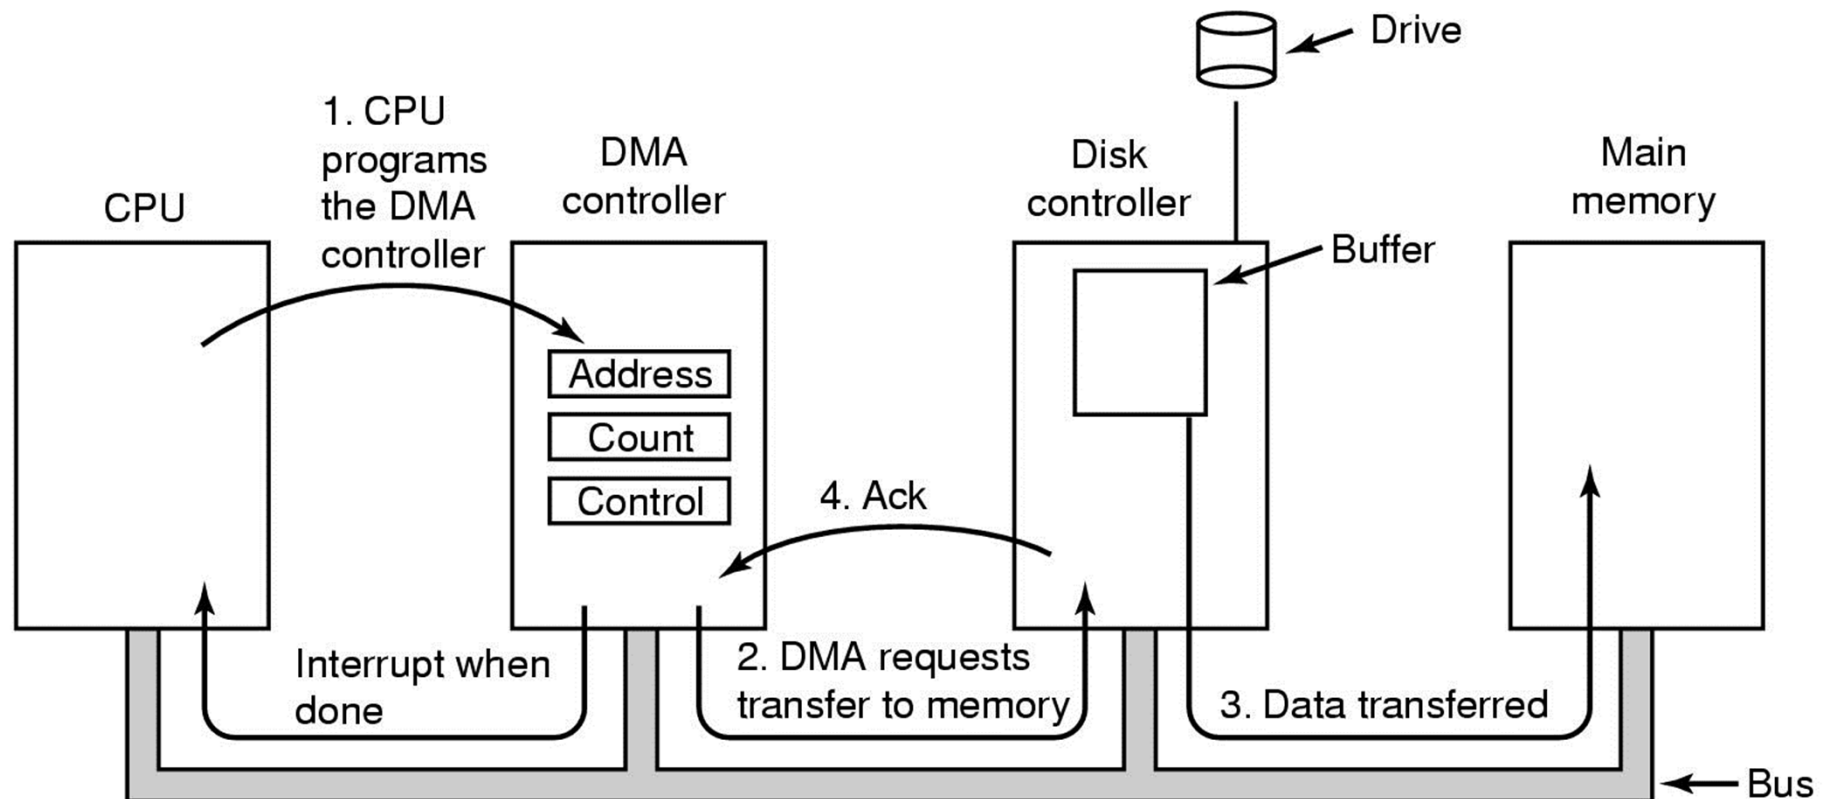
I/O Programming

Direct Memory Access.

- ✓ **Starting address of the buffer to write to.**
- ✓ **The process is blocked until the transfer is done.**
- The OS then (using appropriate I/O ports) passes this information to the DMAC, and tells the DMAC to start transferring.
- The OS then switches to another process, while the DMAC transfers the data independently of the CPU.
- When the DMAC is done, it interrupts the CPU.
 - ✓ **The ISR for the DMAC then calls routines in the OS to move the blocked process to the READY queue.**

I/O Programming

Direct Memory Access.



I/O Programming

Direct Memory Access.

- **In our printer example:**
 - The user process makes an OS call with a pointer to the text buffer, and the number of characters to print.
 - The OS copies the text into its own buffers, then sets up the DMA transfer and blocks the calling process.
 - The OS initiates a DMA transfer, then hands control to another process.
 - When the DMAC is done, it interrupts the OS. The OS moves the calling process into the READY state.

I/O Programming

Direct Memory Access.

```
copy_from_user(buffer, p, count);  
set_up_DMA_controller( );  
scheduler( );
```

(a)

```
acknowledge_interrupt( );  
unblock_user( );  
return_from_interrupt( );
```

(b)

Summary

- **In this lecture we looked at some basic ideas:**
 - What is assembly language?
 - I/O Ports
 - Designing programs to access I/O
 - ✓ Programmed I/O
 - ✓ Interrupt I/O
 - ✓ Direct Memory Access I/O