

CG2271

Real-Time Operating Systems

Lecture 4

Real-Time Software Architectures

colintan@nus.edu.sg



Learning Objectives

- **By the end of this lecture you will be able to:**
 - Understand the various options for designing software for real-time systems.
 - Understand what an RTOS is, and when they are useful.

Introduction

- **A “software architecture” refers to the structure of a piece of software designed to solve a problem.**
 - This structure is important as it determines whether a piece of software can meet the constraints of a problem, and whether it is actually correct.
- **Not all real-time systems require an operating system!!**
 - Some alternatives:
 - ✓ **Round-robin.**
 - ✓ **Round robin with interrupts.**
 - ✓ **Function queue scheduling.**
 - ✓ **Timed loops.**

Real-Time Software Architectures

ROUND ROBIN

Real-Time Software Architectures

Round Robin

- **This is the simplest of all architectures. E.g. to process 3 sensors:**
 - 1. Read sensor 1
 - 2. Do something with sensor 1 data
 - 3. Read sensor 2.
 - 4. Do something with sensor 2 data
 - 5. Read sensor 3.
 - 6. Do something with sensor 3 data
 - 7. Goto 1

Real-Time Software Architectures

Round Robin

- **Round Robin architecture works particularly well when:**
 - There are few devices to service.
 - There is no long complicated processing to be done with the data read in.
 - There are no tight time requirements.
 - ✓ E.g. in a multimeter, if a user turns a dial, it is unlikely that he will notice that the loop is currently reading the voltage probes, before it comes round and reads the new dial settings.

Real-Time Software Architectures

Round Robin

- **However it fails when:**

- Any device requires attention in less time than it takes for the CPU to go around the loop.

- If there is lengthy processing to do. In the multimeter example in the book, if it takes 3 seconds to process a voltage reading:

- ✓ It may take as long as 3 seconds for the multimeter to respond to the user changing the dial setting.

Real-Time Software Architectures

Round Robin

- **This architecture is also fragile:**
 - If we added as little as just one more device, the performance may no longer be acceptable.
 - ✓ **For example, if the new device takes a long time to get a reading.**

Real-Time Software Architectures

ROUND ROBIN WITH INTERRUPTS

Real-Time Software Architectures

Round Robin with Interrupts

- **In this architecture:**
 - When a sensor has data to send to the CPU, it will interrupt the CPU.
 - The interrupt handler reads the device and sets a flag.
 - The main loop polls these flags, and takes action if any of the flags is set.

Real-Time Software Architectures

Round Robin with Interrupts

- **In the example we are going to see, there are 3 devices that must be read, processed, and the output sent to 3 actuators.**
 - Data from sensors written to 3 global variables dev1_data, dev2_data and dev3_data.
 - We have flags dev1_flag, dev2_flag and dev3_flag to indicate newly arrived data.
 - Devices send interrupts when new data arrives.
 - ISRs are set up to deal with new data from devices.
 - Main loops infinitely checking the dev?_flag.
 - ✓ **If set, process the data in dev?_data.**

Real-Time Software Architectures

Round Robin with Interrupts

```
...  
int dev1_flag=0, dev2_flag=0, dev3_flag=0;  
int dev1_data=0, dev2_data=0, dev3_data=0;  
ISR(DEVICE1_vect)  
{  
    !! Read device 1 to dev1_data;  
    dev1_flag=1;  
}  
ISR(DEVICE2_vect)  
{  
    !! Read device 2 to dev2_data;  
    dev2_flag=1;  
}  
ISR(DEVICE3_vect)  
{  
    !! Read device 3 to dev3_data  
    dev3_flag=1;  
}
```

Real-Time Software Architectures

Round Robin with Interrupts

```
int main()
{
    !! Set up interrupts, devices, etc.
    dev1_flag=0; dev2_flag=0; dev3_flag=0;
    while(1)
    {
        if(dev1_flag)
        {
            !! Process dev1_data;
            !! Write to actuator1;
            dev1_flag=0;
        } // if
        if(dev2_flag)
        {
            !! Process dev2_data;
            !! Write to actuator2
            dev2_flag=0;
        } //if
    }
```

Real-Time Software Architectures

Round Robin with Interrupts

```
if(dev3_flag)
{
    !!Process dev3_data
    !! Write to actuator 3
    dev3_flag=0;
} // if
} // while
} // main
```

- **This architecture is an improvement over simple round robin:**
 - Devices get attended to immediately. Unlikely to suffer loss of data.
- **However, it may still take a while for this data to actually be processed!**

Real-Time Software Architectures

Round Robin with Interrupts

- **For example, suppose all 3 devices A, B and C interrupt the CPU:**
 - Data from all 3 devices take 200 ms each in the main loop to process.
 - If the processing sequence is A, B, C, A, B, C, ..., then C will have to wait as long as 400ms to be processed!
- **Even C is a high priority device, it still has to wait for A and B to be done first.**

Real-Time Software Architectures

Round Robin with Interrupts

- **Solutions:**

- Move the processing code for C to its interrupt handler.

- ✓ **Interrupt handler for C will now take 200 ms!**

- ✓ **Unacceptable if the handler is high priority and blocks all lower priority interrupts during this time.**

- Poll C more often:

- ✓ **Instead of A, B, C, A, B, C,..., we poll A, C, B, C, C, A, C, B, C, C, A, C, B, C, C, ...**

Real-Time Software Architectures

FUNCTION QUEUE SCHEDULING

Real-Time Software Architectures

Function Queue Scheduling

- **This is similar to Round Robin with Interrupts:**
 - Interrupt handlers read the data from the device.
 - BUT instead of setting a flag, it inserts a pointer to the function to process this data.
- **The main loop then takes a function from this queue and executes it.**

Real-Time Software Architectures

Function Queue Scheduling

```
void function_A()  
{  
    !! Process data from sensor A  
}  
void function_B()  
{  
    !! Process data from sensor B  
}  
void function_C()  
{  
    !! Process data from sensor C  
}
```

Real-Time Software Architectures

Function Queue Scheduling

```
ISR(DEVICEA_vect)
{
    !! read data from sensor A
    !! Insert pointer to function_A into queue
}
ISR(DEVICEB_vect)
{
    !! read data from sensor B
    !! Insert pointer to function_B into queue
}
ISR(DEVICEC_vect)
{
    !! read data from sensor C
    !! Insert pointer to function_C into queue
}
```

Real-Time Software Architectures

Function Queue Scheduling

```
void main()  
{  
    while(true)  
    {  
        !! Remove function from queue and execute  
        it.  
    } /* While */  
}
```

Real-Time Software Architectures

Function Queue Scheduling

- **It is easy to enforce priorities in this scheme:**
 - Just have priorities in the way the queue is managed!
 - For example, function_C is always placed at the front of the queue ahead of everyone else.
 - This gives function_C priority over everyone else.

Real-Time Software Architectures

TIMED LOOPS

Real-Time Software Architectures

Timed Loops

- **A timed-loop is similar to round-robin.**
 - A while loop repeatedly calls functions to handle processing.
 - Each function is called in a fixed order.
- **Difference:**
 - Functions are called at fixed intervals.
 - Before calling the function, the main loop checks if a sufficient number of clock cycles have passed by.
- **This is useful for routines that must be called at fixed times.**

Real-Time Software Architectures

Timed Loops

- **Example:**

- PID loops require fixed timings for accurate computation.

$$p(t_k) = K(r(t_k) - y(t_k))$$

$$d(t_k) = \frac{T_d}{T_d + Nh} (d(t_{k-1}) - kN(y(t_k) - y(t_{k-1})))$$

$$i(t_{k+1}) = i(t_k) + \frac{Kh}{T_i} (r(t_k) - y(t_k))$$

$$u(t_k) = p(t_k) + i(t_k) + d(t_k)$$

- The computations for $d(t_k)$ and $i(t_{k+1})$ compute integrations and differentiations.

- ✓ Accuracy is dependent on period h , which must be fixed.
- ✓ In our example, we assume a 50 Hz cycle, so $h=20$ ms.


Real-Time Software Architectures

Timed Loops

- **Example code from Ardupilot flight control software.**

```
void loop()
{
    // Execute in 50 Hz loops.
    while(1)
    if (millis()-fast_loopTimer > 19)
    {
        deltaMiliSeconds      = millis() - fast_loopTimer;
        G_Dt = (float)deltaMiliSeconds / 1000.f;
        fast_loopTimer= millis();
        mainLoop_count++;

        // Execute the fast loop
        // -----
        fast_loop();
        // Execute the medium loop
        medium_loop();
    }
}
```



Real-Time Software Architectures

Timed Loops

- **This is how the loop works:**
 - The “`millis()`” function returns the number of milliseconds since the program started execution.
 - The “`fast_loopTimer= millis();`” statement therefore keeps track of the current time.
 - The “`if (millis()-fast_loopTimer > 19)`” at the start therefore checks that 20ms have elapsed before calling “`fast_loop`” and “`medium_loop`”.
 - ✓ This works out to a 50Hz cycle.
 - ✓ Note that both `fast_loop` and `medium_loop` MUST complete under 20 ms or this scheme breaks!

Real-Time Software Architectures

Timed Loops

- **The “fast_loop” function does all the critical flight-control stuff including reading airspeed, computing control responses, etc.**
 - Everything in this loop is executed 50 times a second.

```
void fast_loop()  
{  
    // This is the fast loop - we want it to execute at 50Hz if possible  
    // -----  
    if (deltaMiliSeconds > G_Dt_max)  
        G_Dt_max = deltaMiliSeconds;  
  
    // Read radio  
    // -----  
    read_radio();  
}
```

Real-Time Software Architectures

Timed Loops

```
// check for throttle failsafe condition
// -----
set_failsafe(ch3_failsafe);

// read in the plane's attitude
// -----
GPS.update();
airspeed = (float)GPS.airspeed;
calc_airspeed_errors();

// Read in aircraft's attitude
read_AHRS();
roll_sensor      = -roll_sensor;
pitch_sensor     = -pitch_sensor;
yaw_sensor       = -yaw_sensor;

read_airspeed();
```

Real-Time Software Architectures

Timed Loops

```
read_AHRS();  
// altitude smoothing  
// -----  
if (control_mode != FLY_BY_WIRE_B)  
    calc_altitude_error();  
  
// custom code/exceptions for flight modes  
// -----  
update_current_flight_mode();  
  
// apply desired roll, pitch and yaw to the plane  
// -----  
  
    stabilize();  
  
// write out the servo PWM values  
// -----  
set_servos_4();  
  
}
```

Real-Time Software Architectures

Timed Loops

- **The `medium_loop` function operates at a 10 Hz cycle, but is called 50 times per second by `loop()`.**
 - A counter called “`medium_loopCounter`” is incremented each time the `medium_loop` function is called.
 - A switch-case statement is then used to divide the function into 5 chunks, based on the value of `medium_loopCounter`.
 - ✓ Effectively each chunk is executed only once in every 5 calls to `medium_loop`.
 - ✓ Since `medium_loop` is called 50 times a second, each chunk is called 10 times a second, giving a 10 Hz frequency!

Real-Time Software Architectures

Timed Loops

```
void medium_loop()
{
    // This is the start of the medium (10 Hz) loop pieces
    // -----
    switch(medium_loopCounter) {

        // This case deals with the GPS
        //-----
        case 0:
            medium_loopCounter++;
            ...
            break;

        // This case performs some navigation computations
        //-----
        case 1:
            medium_loopCounter++;
            ...
            break;
    }
}
```


Real-Time Software Architectures

Timed Loops

```
// unused
//-----
case 2:
    medium_loopCounter++;
    ...
    break;
// This case deals with sending high rate telemetry
//-----
case 3:
    medium_loopCounter++;
    ...
// This case controls the slow loop
//-----
case 4:
    medium_loopCounter=0;
    slow_loop();
    break;
}
}
```

Real-Time Software Architectures

Timed Loops

- **The last loop is the `slow_loop` which runs at 3 1/3 Hz.**
 - This is called 10 times per second by `medium_loop`.
 - We use the same trick as before to “step-down” the frequency.
 - ✓ We have a variable called `slow_loopCounter` that is incremented each time `slow_loop` is called.
 - ✓ A switch-case statement divides `slow_loop` into 3 chunks.
 - ✓ Each chunk is executed once every 3 calls, effectively giving a 3 1/3 Hz (10/3) frequency.

Real-Time Software Architectures

Timed Loops

```
void slow_loop()
{
    // This is the slow (3 1/3 Hz) loop pieces
    //-----
    switch (slow_loopCounter){
        case 0:
            slow_loopCounter++;
            !! Handle telemetry comms with ground station.
            break;

        case 1:
            slow_loopCounter++;
            !! Read switches off RC radio
            !! Read battery level.

        case 2:
            slow_loopCounter = 0;
            update_events();
            break;
    }
}
```

Real-Time Software Architectures

REAL-TIME OPERATING SYSTEMS

Real-Time Software Architectures

Real-Time Operating Systems

- **A Real-Time Operating System (RTOS, pronounced ‘arr-toss’), is a suite of software routines that provide:**
 - A (possibly uniform) interface to access hardware devices.
 - ✓ Provides abstraction to layers above the interface.
 - ✓ Hardware accesses below the interface can be customized to each platform.
 - Interrupt handling services.
 - ✓ Interrupts from the hardware are vectored to handlers within the RTOS.
 - ✓ The handlers can be designed so that they respond to different interrupt in *exactly* the same amount of time. This introduces predictability.

Real-Time Software Architectures

Real-Time Operating Systems

- Infrastructure to handle multiple tasks.
 - ✓ Most embedded platforms have single CPUs that can only handle one task at a time.
 - ✓ The RTOS provides services that automatically switch between tasks, to give the illusion that multiple tasks are executing at the same time.
- Infrastructure to coordinate tasks.
 - ✓ Message passing, protecting critical sections, etc.
- Other services.
 - ✓ Memory management.
 - ✓ Disk access.
 - ✓ Etc.

Real-Time Software Architectures

Real-Time Operating Systems

- **When are RTOS good?**
 - Complicated applications that involve many parts that have to interact.
 - ✓ **Reading sensors, reading keypads, reading operator buttons, sounding alarms, controlling actuators, sending/receiving data over the network, driving multiple displays, performing computations, etc.**
 - +RTOS provides a clean, convenient and predictable way to control complex applications.
 - +RTOS are often audited by independent organizations to prove that they are reliable.
 - ✓ **E.g. uC/OS is certified by EuroCAE to be safe for installing on commercial airliners.**
 - RTOS are relatively huge, must be customized, and can take some time to learn and understand.
- **The rest of this course is about RTOS, so we won't go into any further detail here.**

Summary

- **In this lecture we looked at various ways to design real-time software.**
 - **Round-robin.**
 - ✓ **Good for simple applications that read/write to a small set of devices with relatively slow cycle times.**
 - **Round-robin with Interrupts.**
 - ✓ **Good for simple applications that use somewhat faster devices, and each device has equal priority.**
 - **Function Queue Scheduling**
 - ✓ **Like round-robin with interrupts, but good when some devices have higher priority than others.**

Summary

- Timed Loops
 - ✓ Like Round Robin, but with device access at fixed intervals, e.g. every 20 ms.
- Real-Time Operating Systems.
 - ✓ Good for more complex applications with many parts that must be coordinated.
- In subsequent lectures we will look at how RTOS work, and how to design applications using an RTOS.