

CG2271

# Real-Time Operating Systems

## Lecture 5

### Task Management

[colintan@nus.edu.sg](mailto:colintan@nus.edu.sg)



**NUS**  
National University  
of Singapore

School *of* Computing

# Learning Objectives

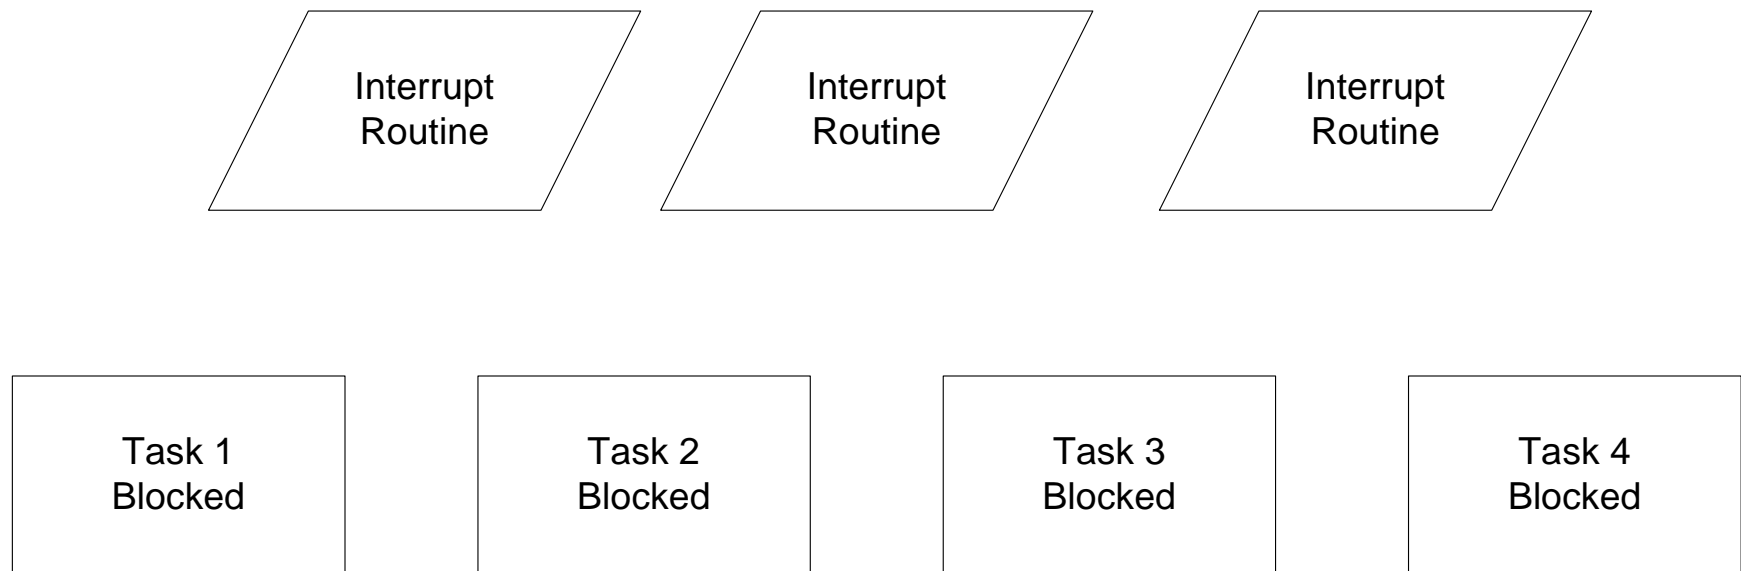
- **By the end of this lecture you will be able to:**
  - Understand the concept of a task, and how tasks behave in an RTOS environment.
  - Understand how tasks are created and managed in an RTOS environment.
  - Understand how an RTOS can simulate multiple tasks running even on a single core single processor system.

## Task Management

# GENERAL OPERATION OF EMBEDDED APPLICATIONS

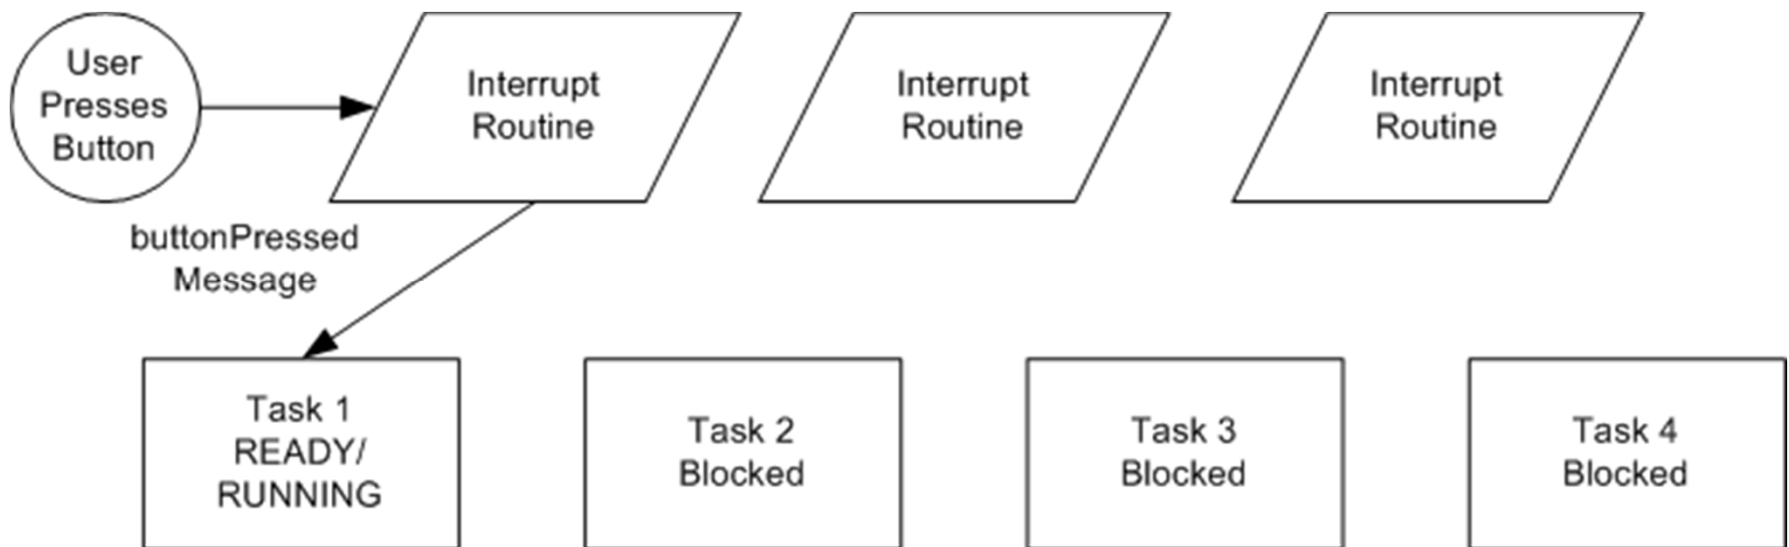
# General Operation of an Embedded Application

Most of the time, tasks are blocked, waiting for something to happen.



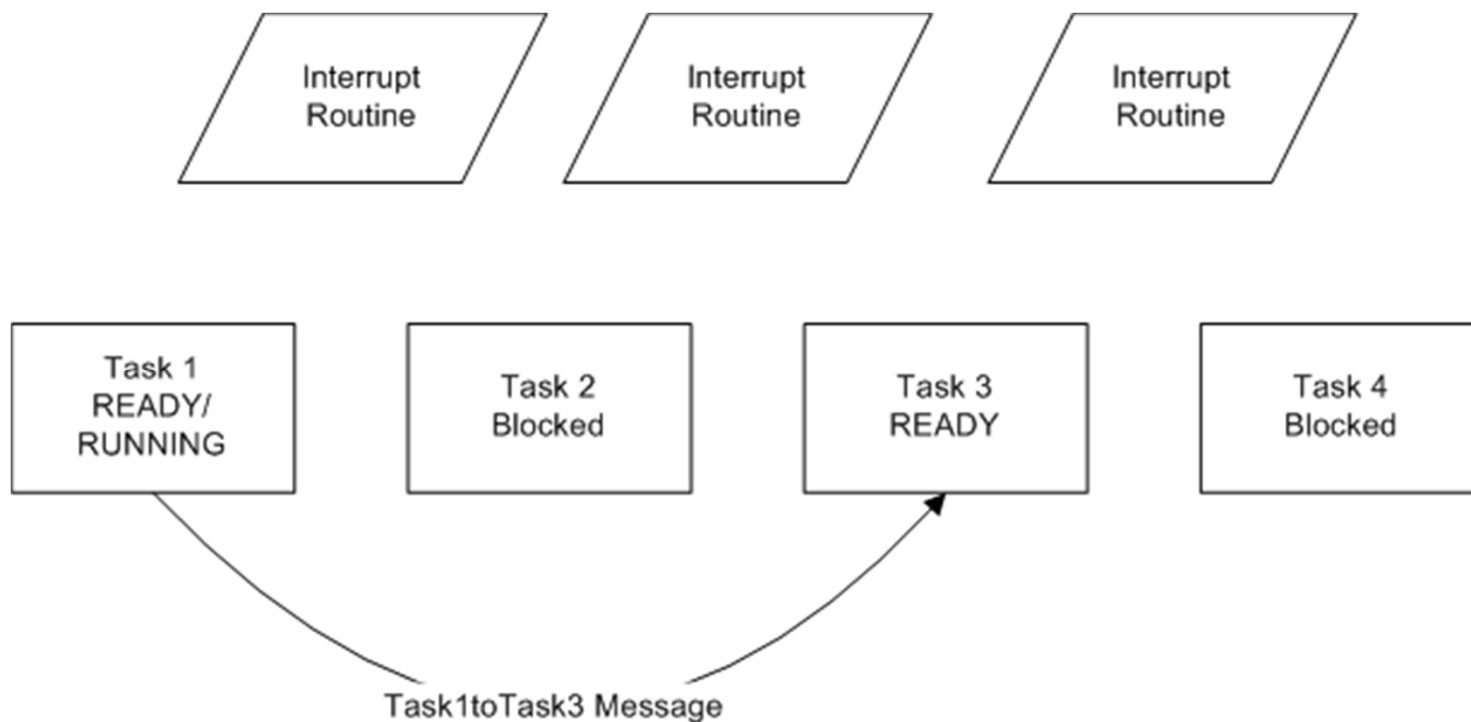
# General Operation of an Embedded Application

Something happens, triggering an interrupt. Interrupt handler sends a message, causing a task to unblock.



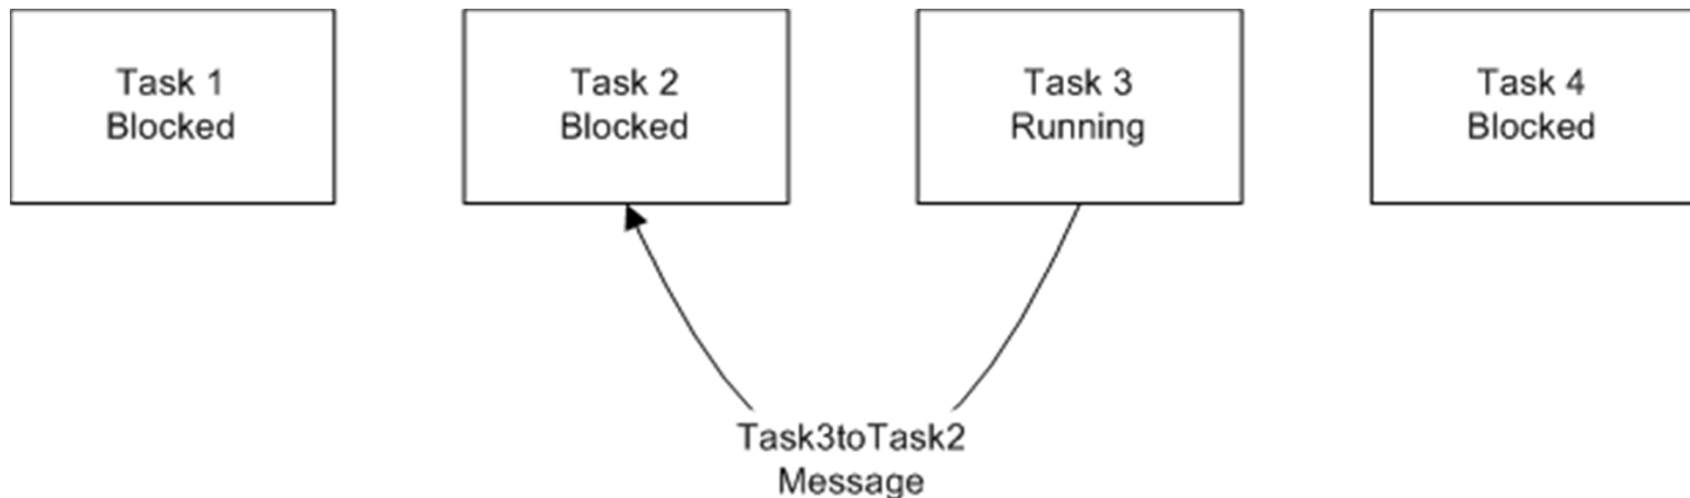
# General Operation of an Embedded Application

Task 1 computes something and passes it to Task 3, which moves from BLOCKED to READY.



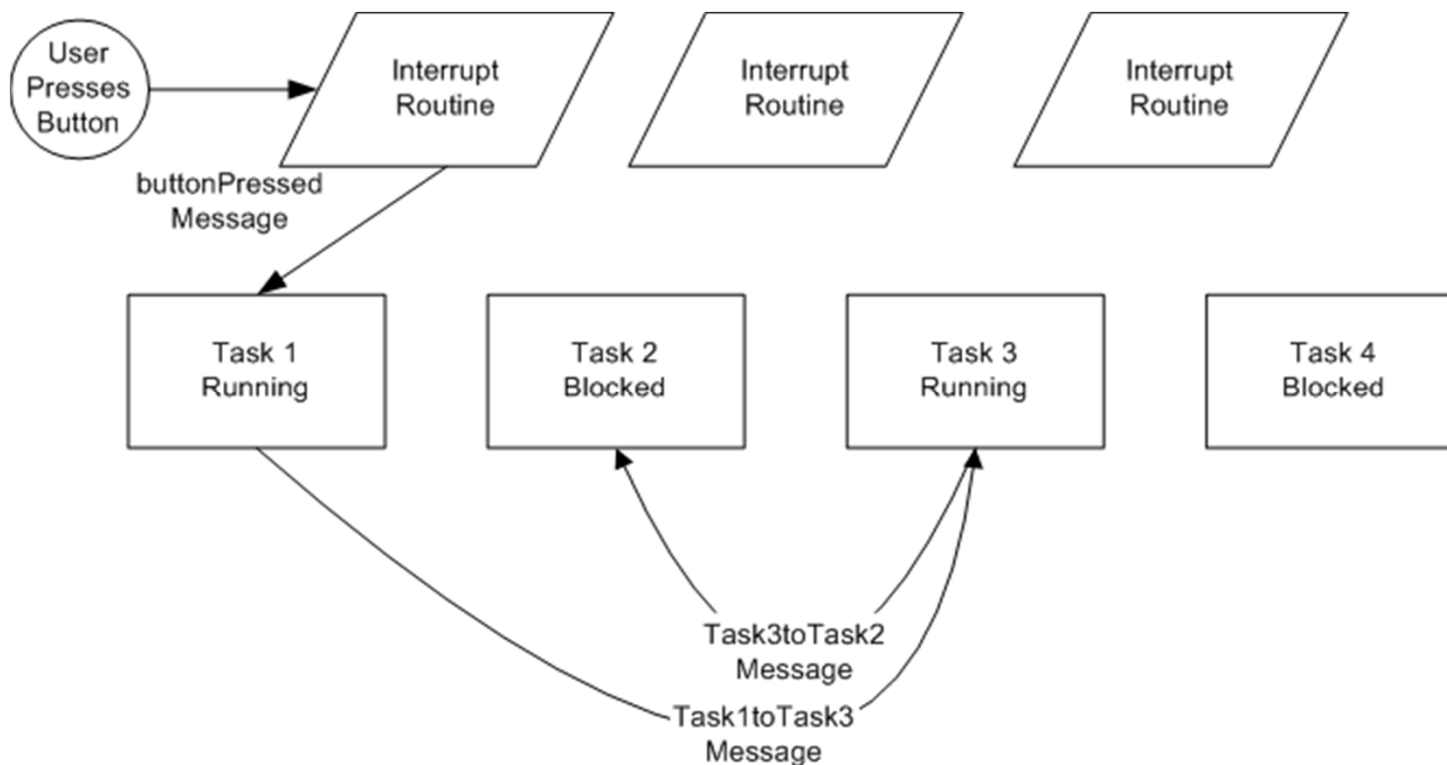
# General Operation of an Embedded Application

Task 3 starts running, sending messages to other tasks, that then start up and send more messages.



# General Operation of an Embedded Application

- **Thus a single interrupt generates a flurry of task and message activity.**





## Task Management

# INTERRUPT ROUTINES

# Interrupt Routines

- **Interrupt routines are thus the backbone of a Real Time System.**
- **Always write short interrupt handlers:**
  - Even the lowest priority interrupt handler will interrupt the highest priority tasks!

# Interrupt Routines

- **Example System:**
  - System responds to commands coming in from a serial port.
  - Commands always end with a carriage return.
  - Commands arrive one at a time. Next command will not come till system responds to previous command.

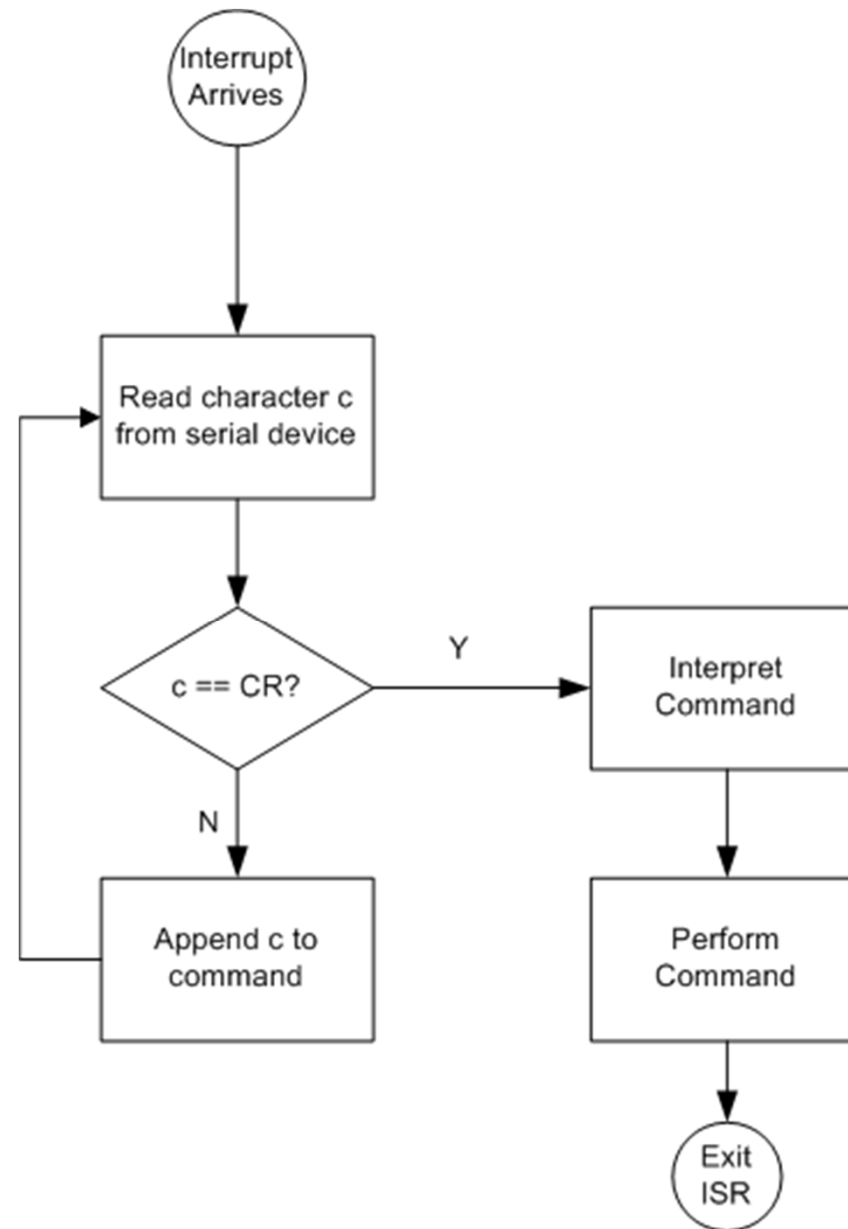
# Interrupt Routines

- Serial port can only store one character at a time.
- Characters may come quickly.
- **We can have three solutions!**

# Solution I

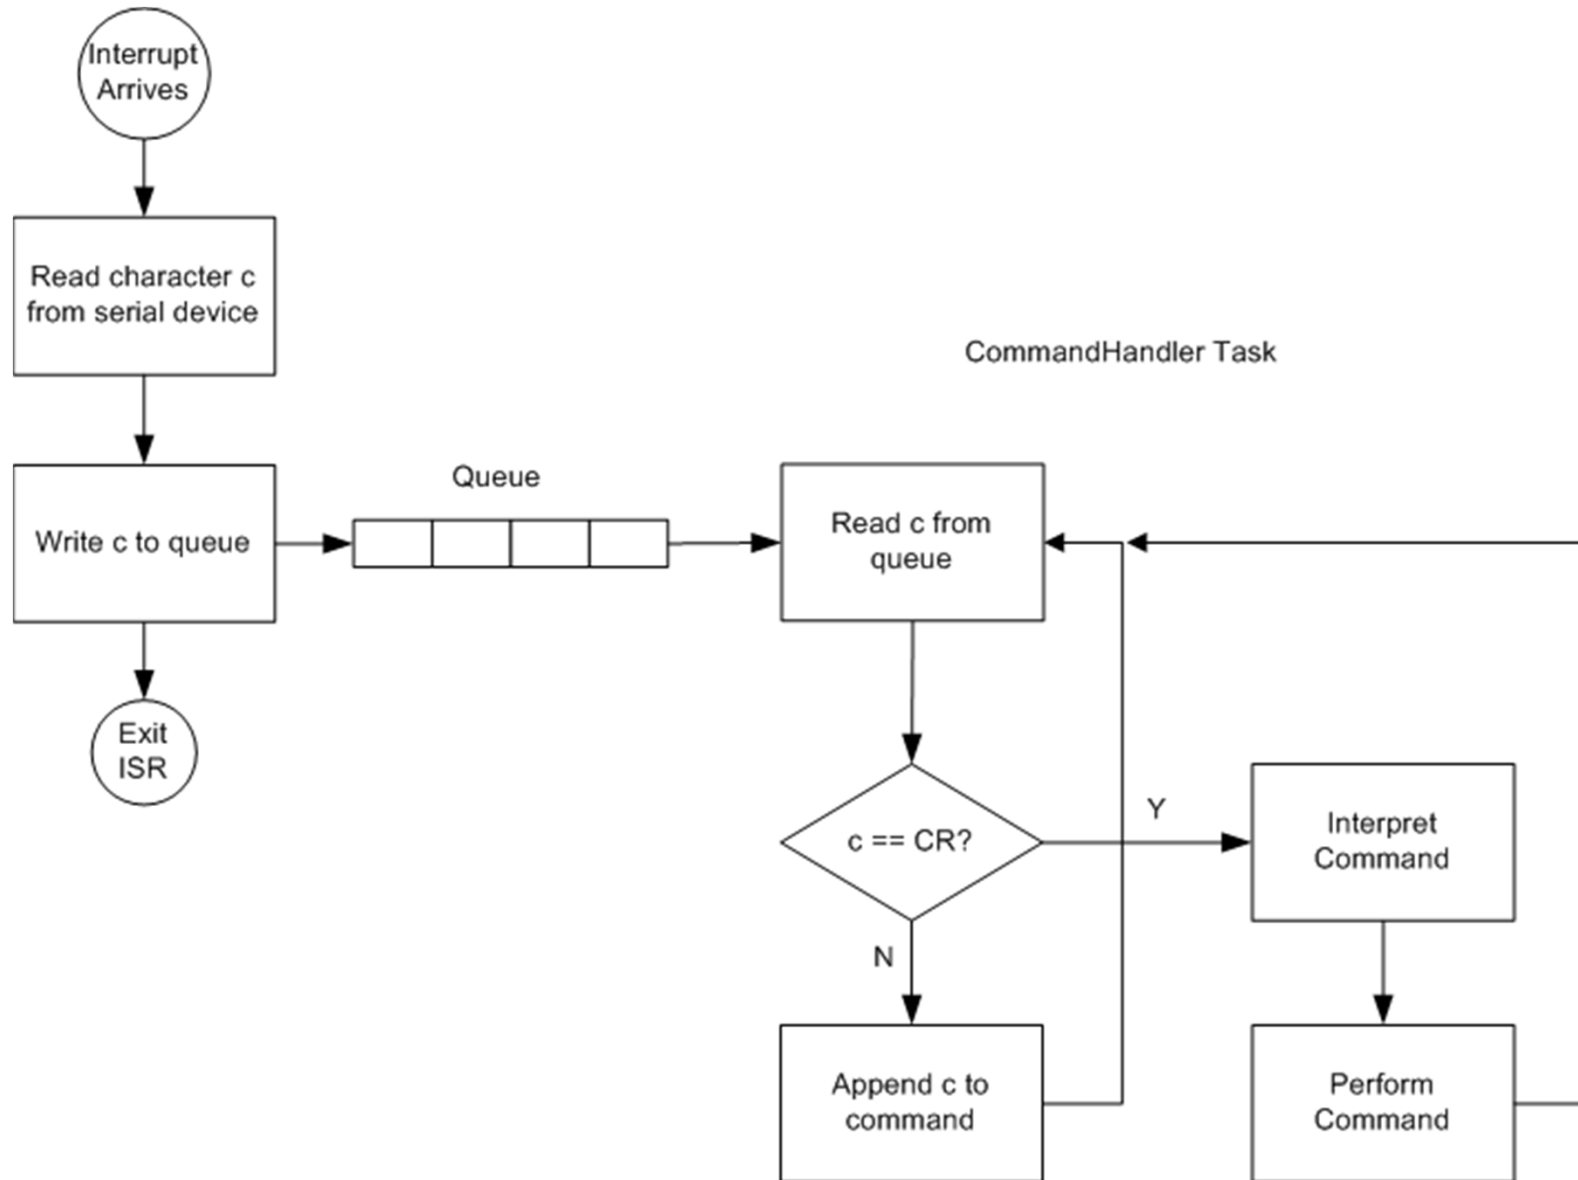
## ISR Does Everything

- **Do everything in the ISR.**
- **BAD solution!**
  - This ISR will take a really long time to complete!
  - High priority tasks are meanwhile suspended.



## Solution II


### ISR Does Almost Nothing



## Solution II

### ISR Does Almost Nothing

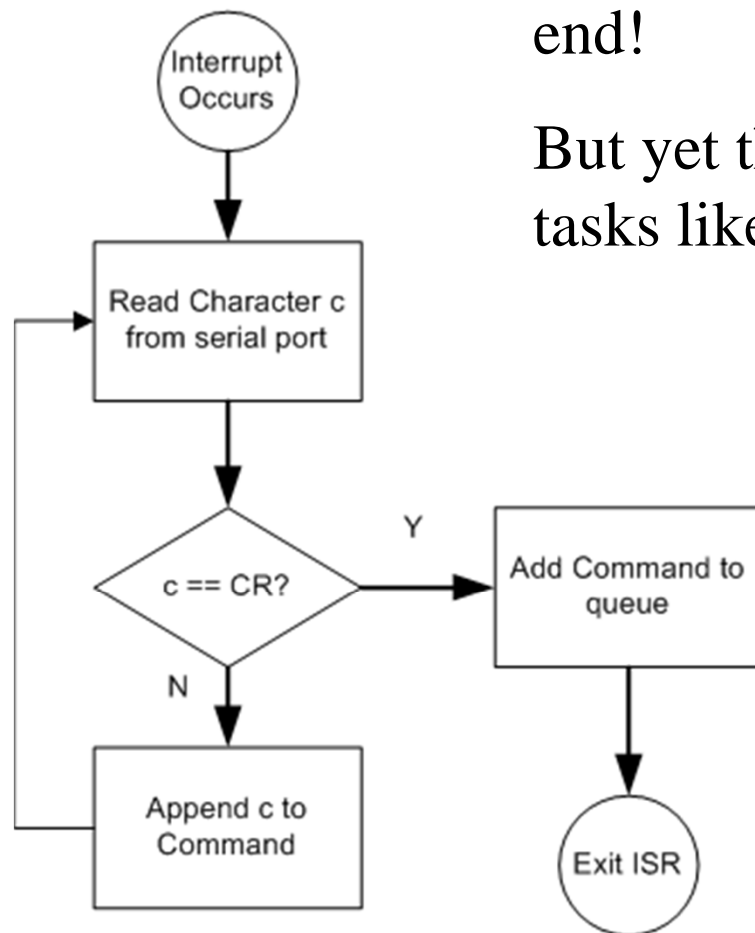
- **This looks like a good idea:**
  - ISR kept *really* short.
- **Trouble is:**
  - ISR sends lots of messages to the task
  - Sending messages takes time!  
✓ See flow chart on the side.
- **This can cause a lot of overheads.**
  - ISR may even start to miss characters.



# Solution III

## ISR Sends Message only at End

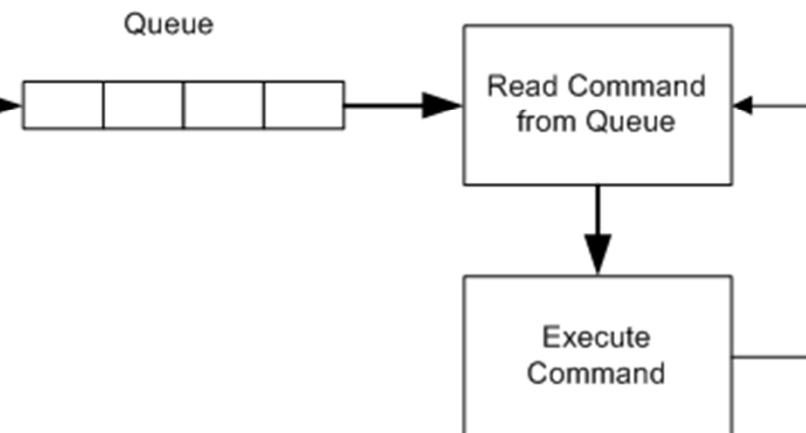
Character ISR



Queue overheads are incurred only at the end!

But yet the ISR does not do unnecessary tasks like interpret commands

Command Handler





## Task Management

# BASIC IDEAS

## Some Basic Ideas

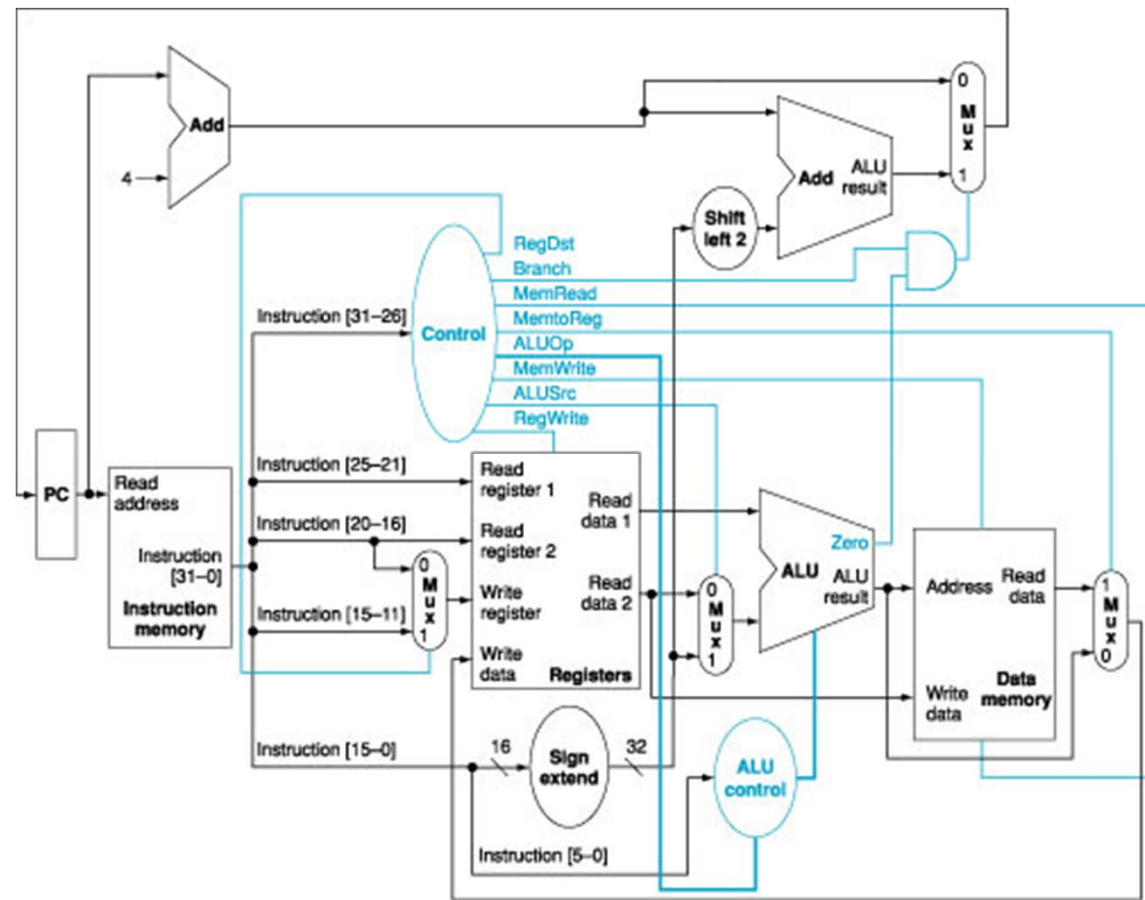
- **Since tasks are the work horses of a real-time system, we will now look at tasks in more detail.**
- **Before we do that however, let's look at some basic ideas first:**
  - CPU Registers.
  - Kernel and User Spaces.

# CPU Registers

- **You have already seen registers.**
  - E.g. the DDRB and PORTB registers that control direction and data on the Port B GPIO pins on the Atmega328.
    - ✓ These are known as “hardware” or “I/O” registers, and are in actual fact memory locations.
    - ✓ These are located outside of the CPU core, and are used to control or configure hardware.
  - Here we will look at a different type of registers.
    - ✓ CPU registers are temporary locations inside the CPU core, and are used to store intermediate computation results.

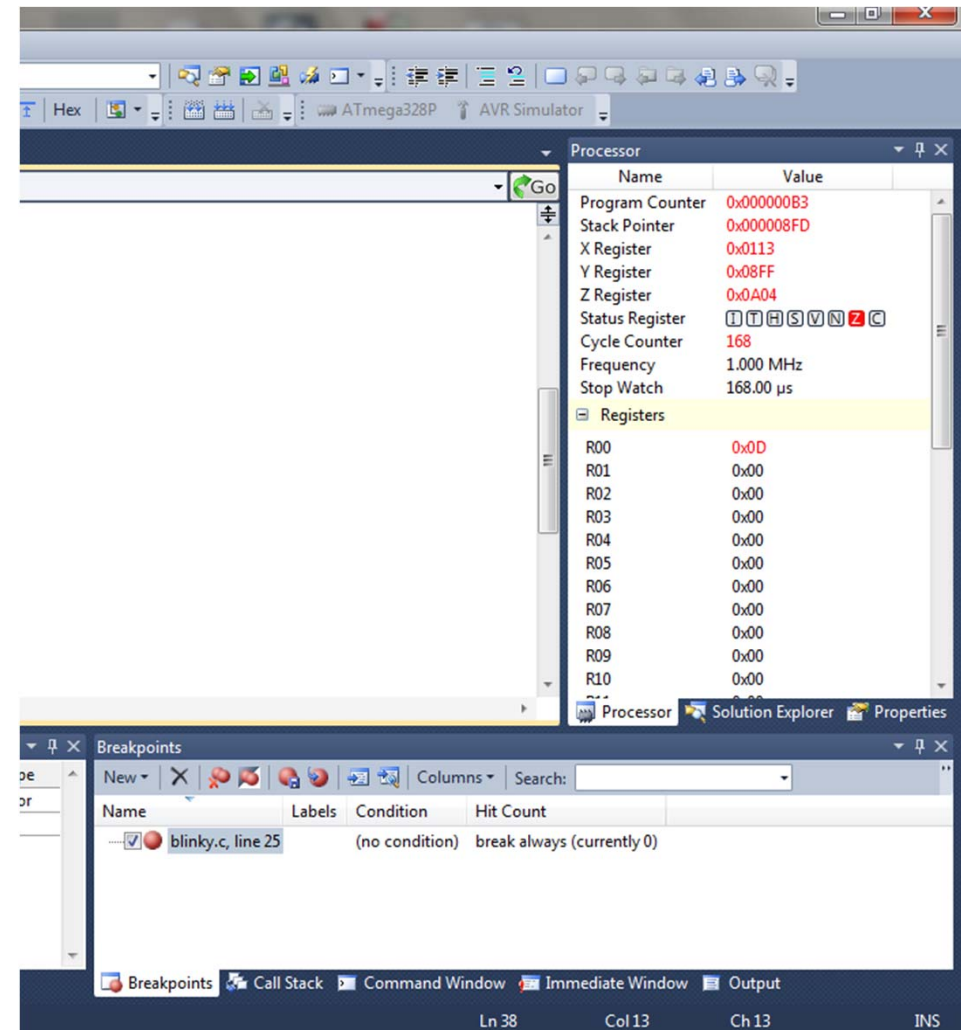
# CPU Registers

- **Take ADD R0, R1, R2 as an example.**
  - This is assembly for  $R0 = R1 + R2$ .
    - ✓ Values are fetched from registers R1 and R2.
    - ✓ The Arithmetic-Logic Unit (ALU) adds these values together.
    - ✓ The answer is written to R0.
  - R0, R1, and R2 are registers holding intermediate results!



# CPU Registers

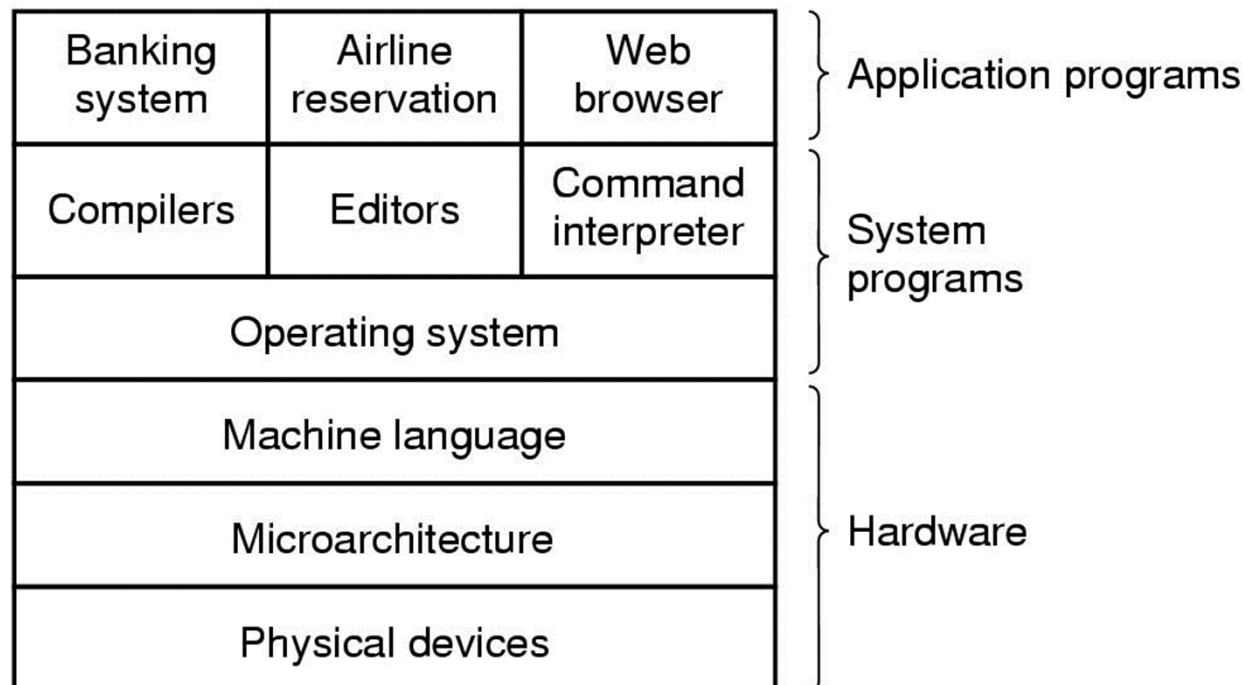
- Registers are often labeled R0 to R31.
  - Alternative schemes:
    - ✓ **AX, BX, CX and DX** in Intel processors.
  - Some registers have aliases.
    - ✓ **R26 and R27** are called “X”, **R28 and R29** are called “Y” and **R30, R31** are called “Z” in the **Atmega328**.
- You can see the register contents in the “Processor” pane in the AVR Studio IDE.



# Kernel and User Spaces

- Kernel vs. User Spaces.**

- The diagram below shows how the hardware, operating system and user programs are typically organized.




# Kernel and User Spaces

- The basic services in the operating system (e.g. to manage memory, manage multiple processes, access hardware, etc), are held in a “kernel”.
  - ✓ All instructions and data used by the operating system operate within the “kernel space”.
- User programs interact with the kernel through an application program interface (API).
  - ✓ All instructions and data used by user programs operate in the “user space”.
- Calls to the OS can involve switching between user and kernel spaces.
  - ✓ This means switching CPU privilege levels.

# Kernel and User Spaces

- **The diagram shows “privilege rings” of a the Intel x86 CPU.**
  - The innermost ring has the highest privileges, allowing access to all hardware and memory resources.
  - The kernel typically operates at the inner rings while user programs operate at the outermost ring.
- **Switching between ring levels takes a lot of time!**
  - This means switching from user space to kernel space is SLOW.

✓ **Avoid if possible!**





## Task Management

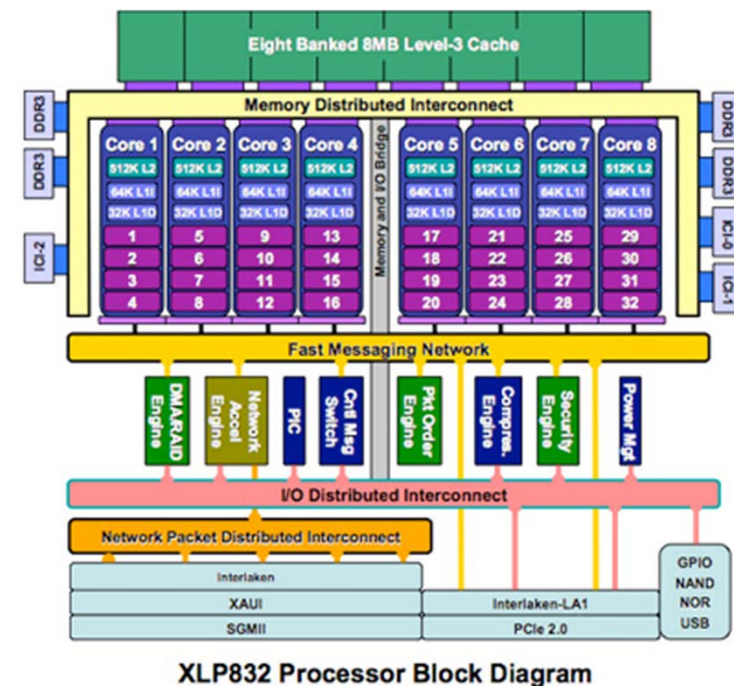
# TASKS AND TASK MANAGEMENT

# Tasks

- **While interrupts are the backbones of real-time systems, tasks are the workhorses.**
  - ISRs get information from sensors, trigger when user presses a button etc.
  - It is tasks that process this information or act on a button press.

# The Process Model


- In this lecture we will assume a single processor with a single core.
  - This is very typical of a microcontroller used in a real-time system.
    - ✓ Some exceptions include the ARM Cortex-A9 used in the iPhone.
  - Most modern desktops however have “multi-core” processors:
    - ✓ Each microprocessor actually consists of multiple execution units.



# The Process Model


- **The materials for this lecture come from Modern Operating Systems.**
  - This book uses the term “process” instead of “task”. We will assume that they mean the same thing.
- **Since we have only a single-core single processor:**
  - At any one time, at most one process can execute.

One program counter




(a)

Four program counters



(b)



(c)

# The Process Model

- **Figure (b) shows what “appears” to be happening in a single processor system running multiple processes:**
  - There are 4 processes each with its own program counter (PC) and registers.
  - All 4 processes run independently of each other at the same time.

# The Process Model


- **Figure (a) shows what actually happens.**
  - There is only a single PC and a single set of registers.
  - When one process ends, there is a “context switch” or “process switch”:
    - ✓ PC, all registers and other process data for Process A is copied to memory.
    - ✓ PC, register and process data for Process B is loaded and B starts executing, etc.
  - Figure (c) illustrates how processes A to D share CPU time.

# Process States

- **A process can be in one of 3 possible states:**
  - **Running**
    - ✓ The process is actually being executed on the CPU.
  - **Ready**
    - ✓ The process is ready to run but not currently running.
    - ✓ A “scheduling algorithm” is used to pick the next process for running.
  - **Blocked.**
    - ✓ The process is waiting for “something” to happen so it is not ready to run yet.
    - ✓ E.g. include waiting for inputs from another process.

# Process States

- The diagram below shows the 3 possible states and the transitions between them.**




1. Process blocks for input
2. Scheduler picks another process
3. Scheduler picks this process
4. Input becomes available



# Process States

- The figure below shows how the processes are organized.
  - The lowest layer selects (schedules) which process to run next.
    - ✓ This is subject to “scheduling policies” which we will look at in a later lecture.



# Switching between Processes

- **When a process runs, the CPU needs to maintain a lot of information about it. This is called the “process context”.**
  - CPU register values.
  - Stack pointers.
  - CPU Status Word Register.
    - ✓ This maintains information about whether the previous instruction resulted in an overflow or a “zero”, whether interrupts are enabled, etc.
    - ✓ This is needed for branch instructions – assembly equivalents of “if” statements.

The AVR Status Register – SREG – is defined as:

Bit	7	6	5	4	3	2	1	0
0x3F (0x5F)	I	T	H	S	V	N	Z	C
Read/Write	R/W	R/W	R/W	R/W	R/W	R/W	R/W	R/W
Initial Value	0	0	0	0	0	0	0	0

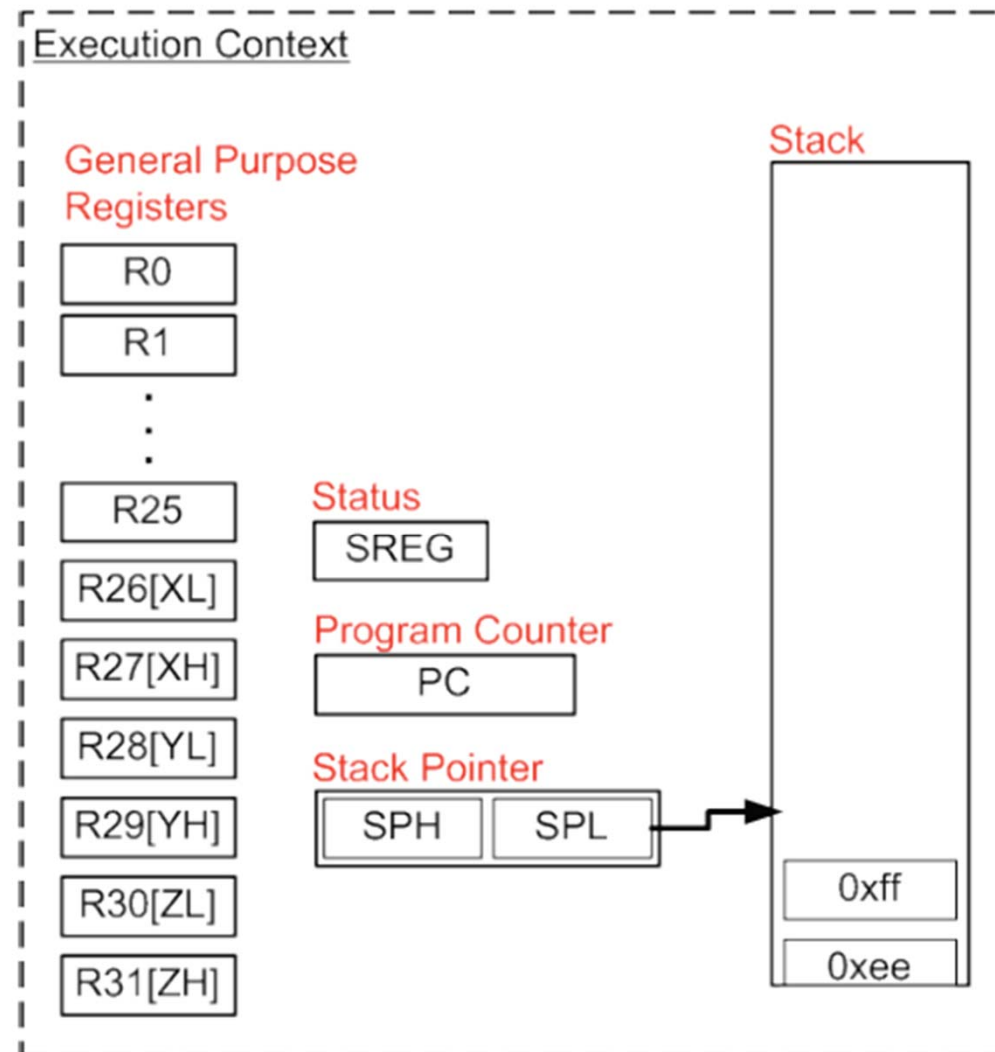
SREG

# Switching between Processes

- **All of these values change as a process runs.**
- **When a process is blocked or put into a READY state, a new process will be picked to take control of the CPU.**
  - All the information for the current process must be saved!
  - The information for the new process must be loaded into the registers, stack pointer and status registers!
    - ✓ **This is to allow the new process to run like as though it was never interrupted!**
- **This process is known as “context switching”.**

# Context Switching on the FreeRTOS Atmega Port

- **Each process is allocated a stack.**
  - Exactly what you learnt in CG1103.
  - The stack pointer are two 8-bit registers SPH and SPL, together forming a single 16-bit SP.
- **The diagram shows the complete Atmega context.**
  - Registers R0-R31, PC.
  - Status register SREG.
  - Stack pointer SPH/SPL.



# Context Switching on the FreeRTOS Atmega Port

- FreeRTOS implements context saving in a macro called “portSAVE\_CONTEXT”.**

```
#define portSAVE_CONTEXT()\
asm volatile (\
    "push r0 \n\t"\           // Save R0
    "in r0, __SREG__ \n\t"\   // Read in status register SREG to R0
    "cli \n\t"\               // Disable all interrupts for atomicity
    "push r0 \n\t"\           // Save SREG
    "push r1 \n\t"\           // Save R1
    "clr r1 \n\t"\            // AVR C expects R1 to be 0, so clear it.
    "push r2 \n\t"\           // Save R2 to R31
    ...
    "push r31 \n\t"\
    "in r26, __SP_L__ \n\t"\   // Read in stack pointer low byte
    "in r27, __SP_H__ \n\t"\   // and high byte
    "sts pxCurrentTCB+1, r27 \n\t"\ // And save it to pxCurrentTCB
    "sts pxCurrentTCB, r26 \n\t"\
    "sei \n\t" : : : \       // Re-enable interrupts
);
```

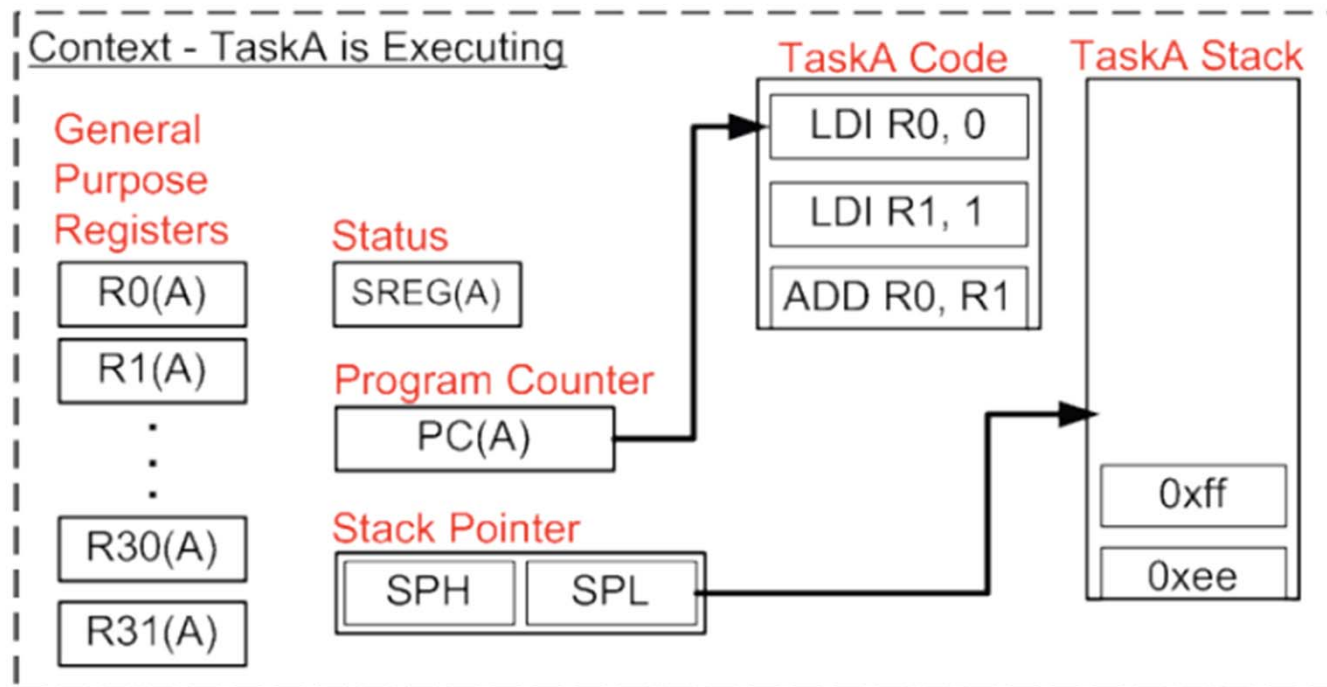
# Context Switching on the FreeRTOS Atmega Port

- **The reverse operation is portRESTORE\_CONTEXT. The stack pointer for the process being restored must be in pxCurrentTCB.**

```
#define portRESTORE_CONTEXT()\nasm volatile (\n    "out __SP_L__, %A0 \\n\\t"\\        // Copy SP_L and SP_H from the\n    "out __SP_H__, %B0 \\n\\t"\\        // pxCurrentTCB variable.\n    "pop r31 \\n\\t"\\                // Restore registers r31 to r1.\n    ...\n    "pop r0 \\n\\t"\\                    // Pop out SREG\n    "out __SREG__, r0\\n\\t"\\          // And restore it.\n    "pop r0 \\n\\t": : "r" (pxCurrentTCB):\\        // Restore R0\n);\n
```

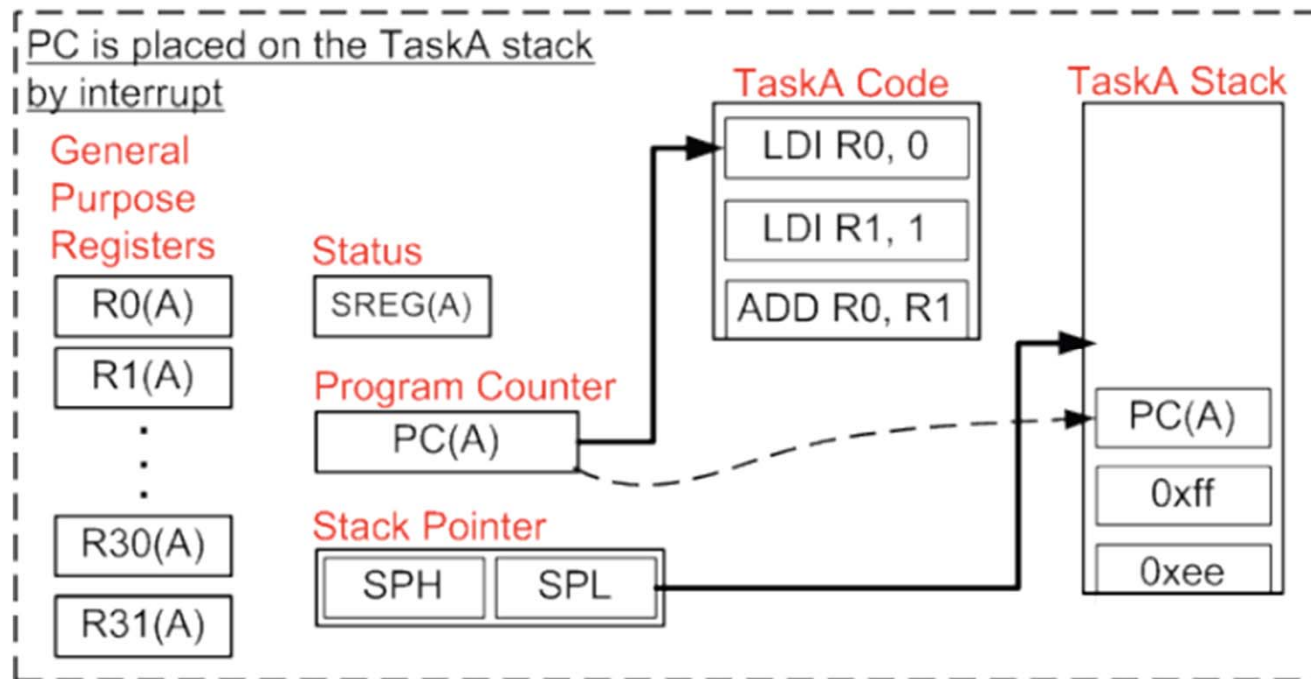
# Context Switching on the FreeRTOS Atmega Port

- We will now see step-by-step how this works.
- Assume that at first Task A is executing.
  - PC would be pointing at Task A code, SPH/SPL pointing at Task A stack, Registers R0-R31 contain Task A data.



# Context Switching on the FreeRTOS Atmega Port

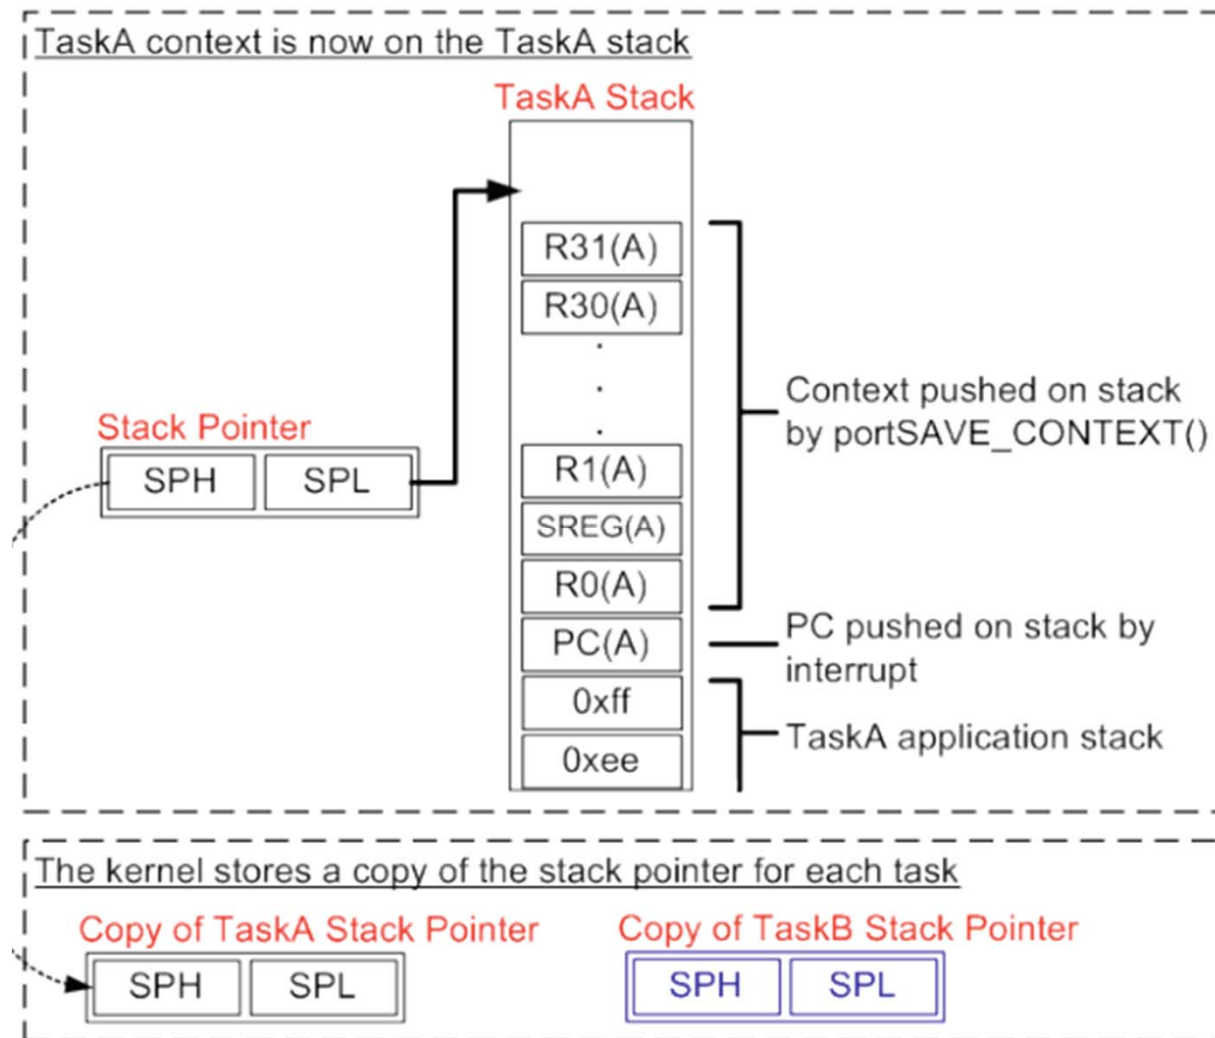
- FreeRTOS relies on regular interrupts from Timer 0 to switch between tasks. When the interrupt triggers, PC is placed onto Task A's stack.**





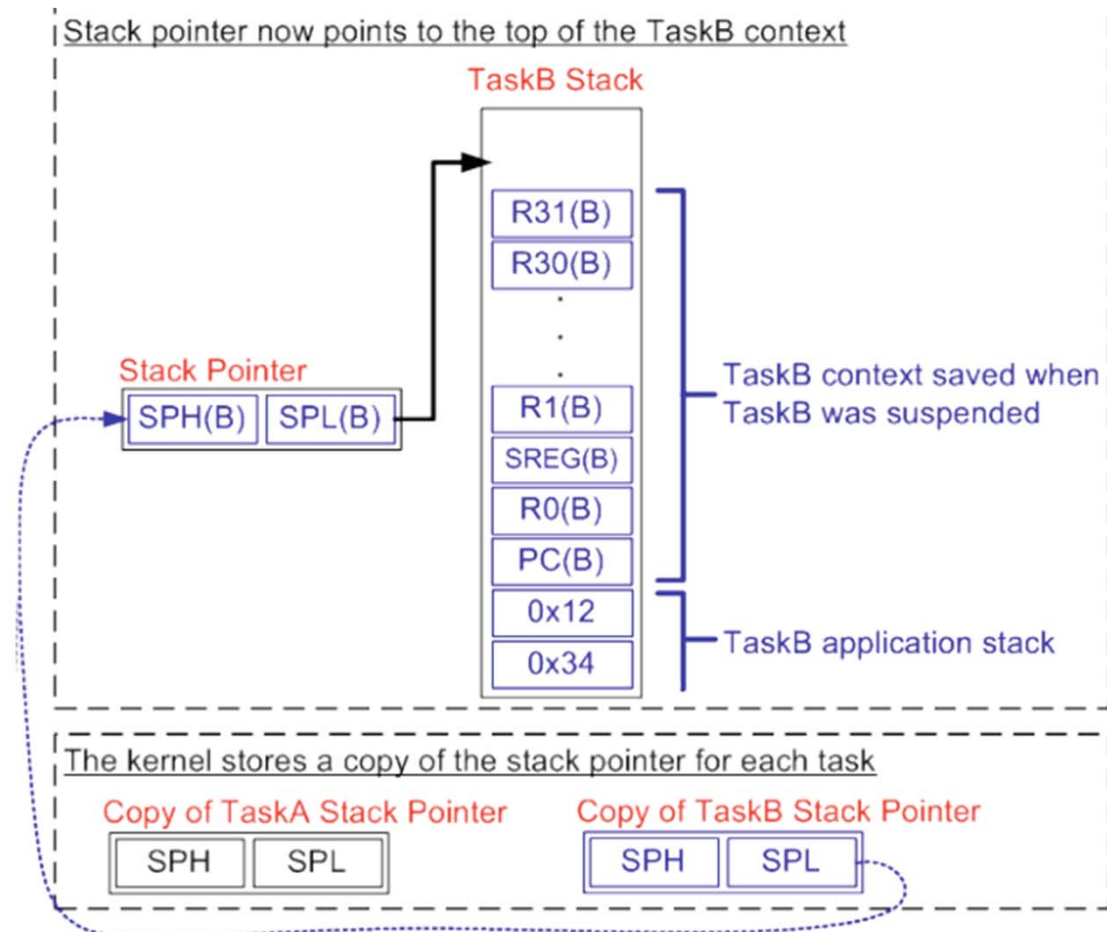
# Context Switching on the FreeRTOS Atmega Port

- The ISR calls `portSAVECONTEXT`, resulting in Task A's context being pushed onto the stack.
- `pxCurrentTCB` will also hold SPH/SPL after the context save.
  - This must be saved by the kernel.



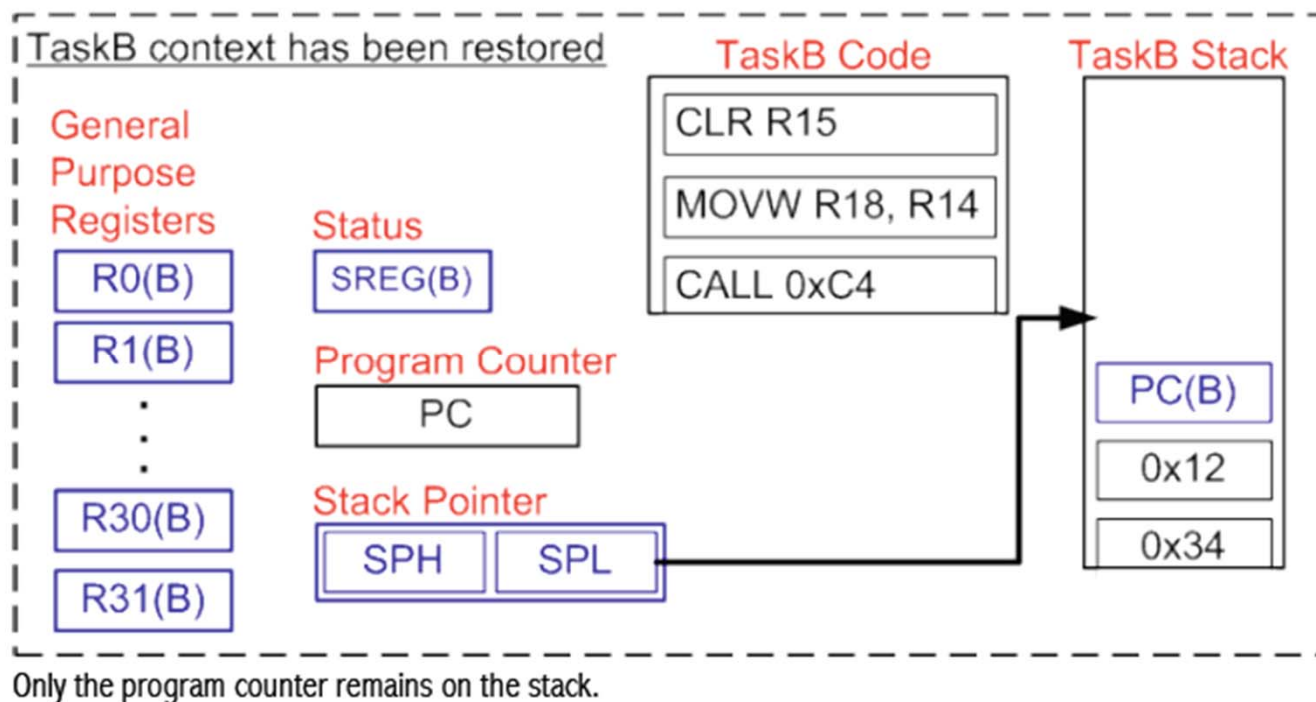
# Context Switching on the FreeRTOS Atmega Port

- **The kernel then selects Task B to run, and copies its SPH/SPL values into pxCurrentTCB and calls portRESTORE\_CONTEXT.**
  - The first two lines will copy pxCurrentTCB into SPH/SPL, causing SP to point to Task B's stack.



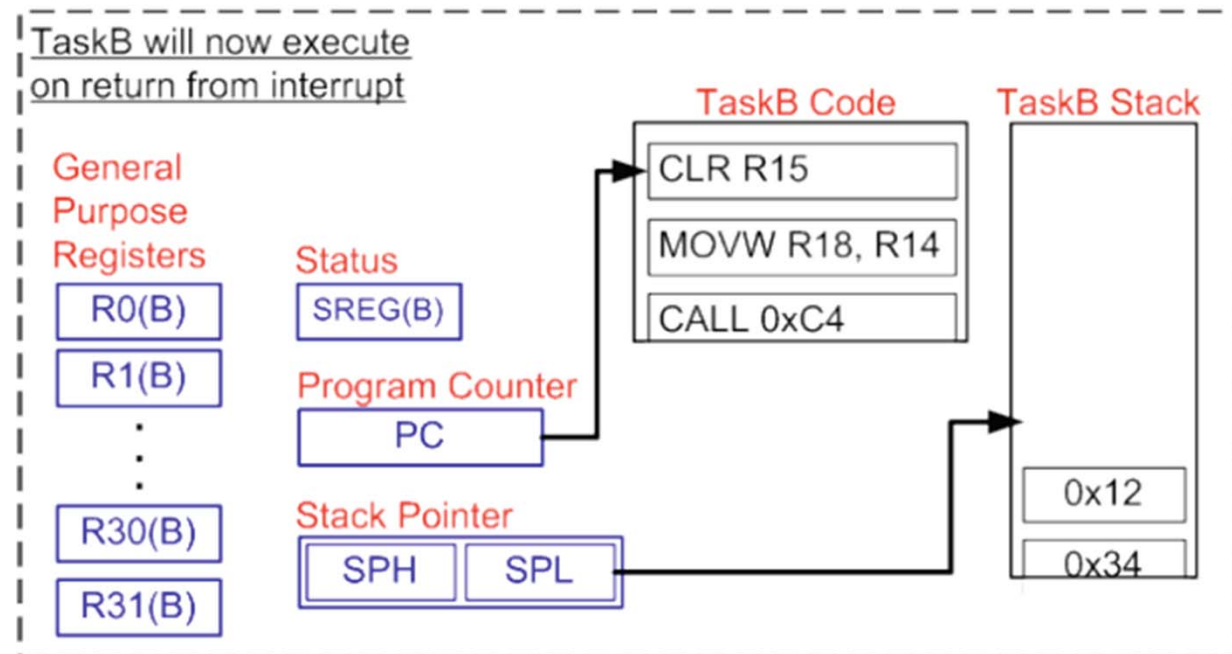
# Context Switching on the FreeRTOS Atmega Port

- The rest of `portRESTORE_CONTEXT` is executed, causing Task B's data to be loaded into R31-R0 and SREG.
  - Now Task B can resume like as though nothing happened!



# Context Switching on the FreeRTOS Atmega Port

- **Only Task B's PC remains on the stack. Now the ISR exits, causing this value to be popped off onto the AVR's PC.**
  - PC points to the next instruction to be executed.
  - End result: Task B resumes execution, with all its data and SREG intact!



# Context Switching on the FreeRTOS Atmega Port

- **Here we looked at context switching controlled by a timer.**
  - **It can also be triggered by other things:**
    - ✓ **Currently running processed waiting for input.**
    - ✓ **Currently running task blocking on a synchronization mechanism (see next lecture).**
    - ✓ **Currently running task wants to sleep for a fixed period.**
    - ✓ **Higher priority task becoming “READY”.**
    - ✓ **...**

## Task Management

# MODELING MULTI-PROGRAMMING

# Modeling Multiprogramming

- **In a system running just one process, suppose the CPU is only 20% utilized.**
- **Then logically, in a system running 5 processes, the CPU will be 100% utilized!**
- **Sadly this is not the case.**
  - Much of the time the task might spend on I/O.
  - While the task is waiting for I/O it will be in the BLOCKED state and cannot run.
- **A better model would be to assume that the CPU spends  $p$  percent of the time doing I/O.**

# Modeling Multiprogramming

- The probability of  $n$  processes all doing I/O at the same time is therefore  $p^n$
- The CPU utilization will therefore be  $1-p^n$
- The figure below shows CPU utilization for differing  $p$  and  $n$ .
  - Diagram omitted. Reserved for your tutorial question. 😊



# Modeling Multiprogramming

- **The model given here is overly simplistic. It assumes that the tasks are independent. This cannot be the case:**
  - Since there is only 1 CPU, even when a task becomes READY, it cannot run if the CPU is BUSY.
- **A more accurate model can be derived using queuing theory.**
- **This model is nonetheless useful for giving “first-cut” guesses at CPU utilization.**

# Summary

- **In this lecture we looked at:**
  - How real-time applications run.
  - How multiple tasks are managed.