

CG2271

Real-Time Operating Systems

Lecture 7

Inter-Task Communication

colintan@nus.edu.sg



Learning Objectives

- **By the end of this lecture you will be able to:**
 - Understand what race conditions are, and why they are bad.
 - Understand the various ways to prevent race conditions.
 - Understand how to pass messages between tasks.

Introduction

- **In the “previous, previous” lecture we looked at how multiple tasks can run on a single CPU.**
- **In the previous lecture, we then looked at how a kernel chooses which task to run next.**
 - We basically assumed that tasks are independent!
- **In real-world applications, there are “dependencies” between tasks.**
 - Task B cannot proceed because it is waiting for Task A’s result.
 - Task B and Task A update the same shared variable, which can result in errors.
 - ...

Introduction

- **If both Task A and Task B are allowed to run freely, errors will occur.**
 - Task B proceeds before Task A completes, resulting in B using stale results.
 - Task A and B update a variable at the same time, causing one task to over-write the results of the other task.
 - Etc.
 - **Some form of coordination is therefore required!**
 - **This lecture comes from Modern Operating Systems.**
 - MOS uses the term “process” instead of task.
 - To be consistent with the book, we will therefore also say “process” instead of task.
- ✓ **Just remember that they mean the same thing!!**

Inter-Task Communications

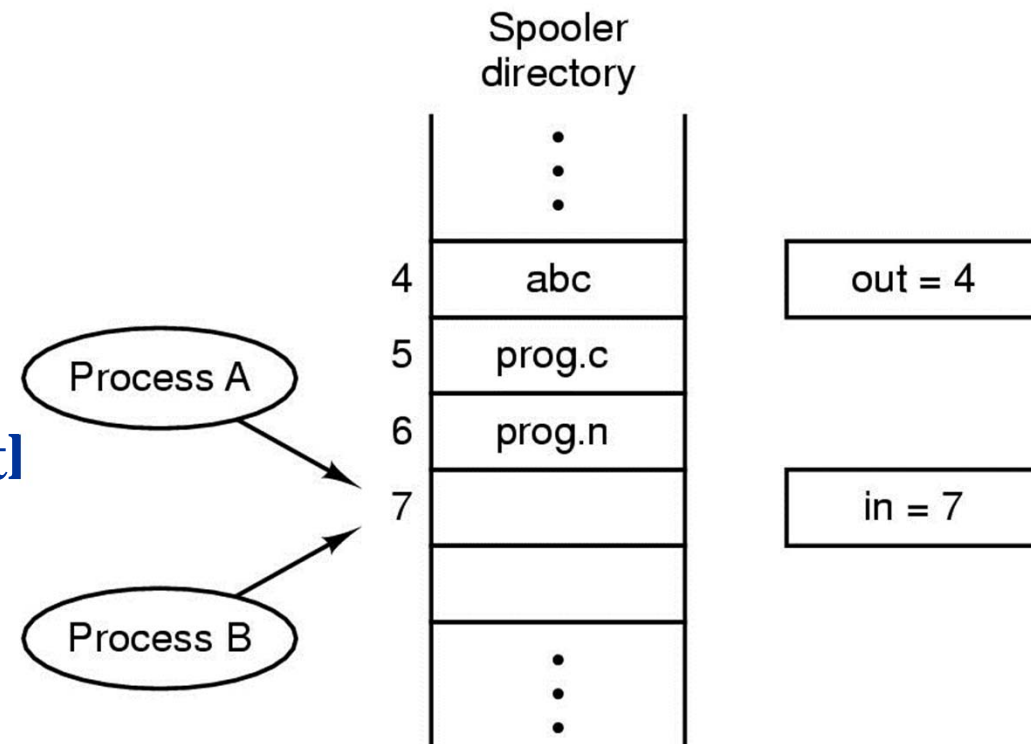
RACE CONDITIONS AND CRITICAL SECTIONS

Race Conditions

- **Race conditions occur when two or more processes attempt to access shared storage.**
 - This causes the final outcome to depend on who runs first.
 - “Shared storage” can mean:
 - ✓ **Global variables.**
 - ✓ **Memory locations.**
 - ✓ **Hardware registers.**
 - *This refers to configuration registers rather than CPU registers.*
 - ✓ **Files.**
 - To understand race conditions, we will consider the example of a queue in a print spooler.

Race Conditions

- A process that wants to print enters the name of the file into a print directory.
- A print daemon (printd) periodically checks the directory, prints out the next file and removes its entry.
- Two variables keep track of the queue:
 - IN: Next available slot.
 - OUT: Next file for printing.



Race Conditions

- **The following can happen:**
 - Process A reads IN as 7 and stores it in a local variable *next*.
 - The OS pre-empts A and starts B.
 - B reads IN as 7 and stores it in a local variable *next*.
 - B inserts its file into slot 7 and updates IN to 8.
 - B is pre-empted and A restarts.
 - A still thinks slot 7 is available and inserts its file there, overwriting B's file, then updates IN to 8.
- **The daemon is unaware of the mistake, and B never gets a printout. ☹**

Critical Sections

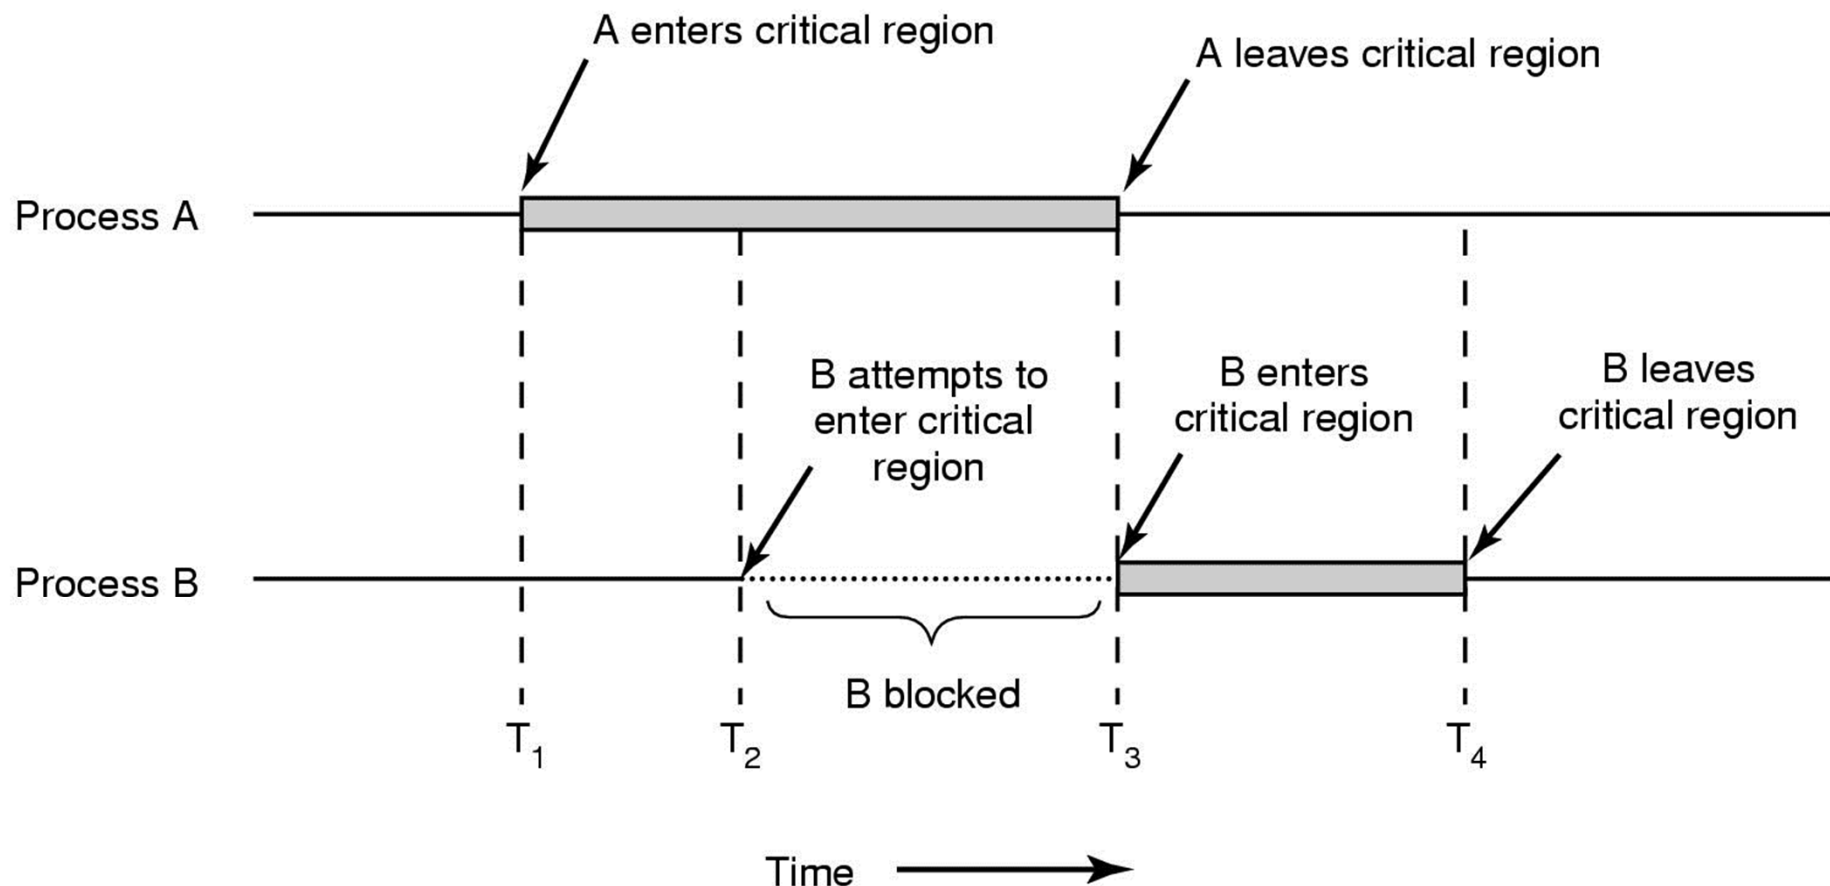
- **To prevent race conditions, we must prevent two processes from reading/writing shared resources at the same time.**
- **This is known as a “mutual exclusion”, shortened to “mutex”.**
- **Conceptually, a RUNNING process is always in one of two possible states:**
 - It is performing local computation. This does not involve global storage, hence to race condition is possible.
 - It is reading/updating global variables. This can lead to race conditions.
- **When a RUNNING process is in the second state, it is within its “critical section”.**

Critical Sections

- **To prevent race conditions, 4 rules must be followed:**
 - No two processes can simultaneously be in their critical section.
 - No assumptions may be made about speeds or # of CPUs.
 - ✓ **Note: We can relax this assumption for *most* embedded systems since they have single CPUs.**
 - ✓ **May apply to systems using multicore microcontrollers.**
 - No process outside of its critical section can block other processes.
 - No process should wait forever to enter its critical section.

Critical Sections

- In an ideal state, this is how mutual exclusion works:**



Inter-Task Communications

IMPLEMENTING MUTUAL EXCLUSION

Implementing Mutual Exclusion

- **There are several ways of implementing mutexes, each with their own + and – points:**
 - Disabling interrupts.
 - Lock variables.
 - Strict alternation.
 - Perterson's Solution.
 - Test and set lock.
 - Sleep/Wakeup

Implementing Mutual Exclusion

Disabling Interrupts

- **Disabling Interrupts.**
 - This works because:
 - ✓ Time-slicing depends on a timer interrupt. If this is disabled, the scheduler is never activated to switch to another process.
 - ✓ Similarly, processes that are blocked pending an event (e.g. arrival of data from the network), depend on an interrupt to tell the scheduler that the event has taken place.
 - Therefore disabling interrupts will prevent other processes from starting up and entering their critical sections.

Implementing Mutual Exclusion

Disabling Interrupts

- **Disabling Interrupts.**
 - There are several problems with this approach:
 - ✓ Carelessly disabling interrupts can cause the entire system to grind to a halt.
 - ✓ This only works on single-processor, single core systems.
- Violates Rule 2.**

Implementing Mutual Exclusion

Lock Variables

- **Using Lock Variables.**
 - A single global variable “lock” is initially 1.
 - Process A reads this variable and sets it to 0, and enters its critical section.
 - Process B reads “lock” and sees it’s a 0. It doesn’t enter its critical section and waits until “lock” is 1.
 - Process A finishes and sets “lock” to 1, allowing B to enter.

Implementing Mutual Exclusion

Lock Variables

- **This approach obviously doesn't work!!**
 - Process A reads in “lock” and sees a “1”. It gets pre-empted and Process B runs.
 - Process B reads in “lock”, sees a “1”, sets it to 0 and enters its critical section.
 - Before B leaves, A is re-started, and enters the critical section.
- **Now >1 process is in the critical section!**
- **PROBLEM: There's a race condition on “lock” itself!**

Implementing Mutual Exclusion

Strict Alternation

```
while (TRUE) {  
    while (turn != 0)      /* loop */ ;  
    critical_region( );  
    turn = 1;  
    noncritical_region( );  
}
```

(a)

```
while (TRUE) {  
    while (turn != 1)      /* loop */ ;  
    critical_region( );  
    turn = 0;  
    noncritical_region( );  
}
```

(b)

- A “turn” variable keeps track of whose turn it is to enter the critical section.
- If it is currently “0”, then A will enter and B will continuously test until it becomes “1”.
- Once A finishes its critical section, it flips “turn” to 1 allowing B to enter the critical region. B flips it back to 0 when done.

Implementing Mutual Exclusion

Strict Alternation

- **Problems (assume “turn”=0):**
 - Since “turn” is 0, B just loops infinitely waiting for “turn” to be 1.
 - ✓ **This is called “busy-waiting” and burns valuable CPU time for no reason.**
 - If “A” is spending a lot of time in its non-critical section, it will not reach the part of the loop where it flips “turn” to 1.
 - Result is that B cannot enter its critical section even though no one is there!
- **Violates Rule #3.**

Implementing Mutual Exclusion

Peterson's Solution

```
#define FALSE 0
#define TRUE 1
#define N      2          /* number of processes */

int turn;                  /* whose turn is it? */
int interested[N];         /* all values initially 0 (FALSE) */

void enter_region(int process) /* process is 0 or 1 */
{
    int other;              /* number of the other process */

    other = 1 - process;    /* the opposite of process */
    interested[process] = TRUE; /* show that you are interested */
    turn = process;         /* set flag */
    while (turn == process && interested[other] == TRUE) /* null statement */ ;
}

void leave_region(int process) /* process: who is leaving */
{
    interested[process] = FALSE; /* indicate departure from critical region */
}
```

Implementing Mutual Exclusion

Peterson's Solution

- **Suppose Process A (i.e. process 0) wants to enter its critical section.**
 - It calls `enter_region`.
 - ✓ “other” is set to 1, “turn” is set to 0.
 - ✓ `interest[0]` is set to 1.
 - ✓ Since “`interest[other]=interest[1]=0`”, the while loop exits immediately and A enters the critical section.
 - If Process B (process 1) wants to enter and calls “enter region”:
 - ✓ “other” is set to 0, “turn” is set to 1, `interest[1]` is set to 1.
 - ✓ It checks “`interest[other]=interest[0]=1`”, it loops until `interest[0]` is set to 0.
 - When process A ends, it calls `leave_region` which sets `interest[0]` to 0 and allows Process B to exit the while loop.

Test and Set Lock

- **Many microprocessors have an instruction that looks like this:**
 - TSL reg, lock ; lock is a variable in memory
- **How this works:**
 - CPU locks the address and data buses, and reads “lock” from memory.
 - ✓ **The locked address and data buses will block accesses from all other CPUs.**
 - The current value is written into register “reg”.
 - A “1” (or sometimes “0”) value is written to “lock”.
 - CPU unlocks the address and data buses.
- **The TSL is “atomic”.**
 - This means that NOTHING can interrupt execution of this instruction.
 - This is guaranteed in hardware.

Test and Set Lock

- **The TSL instruction is used as follows:**

enter_region:

| | |
|-------------------|---|
| TSL REGISTER,LOCK | copy lock to register and set lock to 1 |
| CMP REGISTER,#0 | was lock zero? |
| JNE enter_region | if it was non zero, lock was set, so loop |
| RET | return to caller; critical region entered |

leave_region:

| | |
|--------------|-------------------|
| MOVE LOCK,#0 | store a 0 in lock |
| RET | return to caller |

Test and Set Lock

- **An alternative is the XCHG instruction, used on Intel machines.**
 - Swaps contents of “lock” and “reg” instead of just writing “1” to lock.

```
enter_region:
    MOVE REGISTER, #1          ; Set REGISTER to 1
    XCHG REGISTER, LOCK        ; Exchange with Lock
    CMP REGISTER, #0           ; Was Lock 0?
    JNE enter_region           ; No, go back and try again.
    RET                         ; Yes, enter critical region.

leave_region:
    MOVE LOCK, #0              ; Clear the lock
    RET                        ; And exit
```


Deadlock

- **Busy-wait approaches like Peterson and TSL/XCHG have a problem called “deadlock”.**
- **Consider two processes H and L, and a scheduler rule that says that H is always run when it is READY. Suppose L is currently in the critical region.**
 - H becomes ready, and L is pre-empted.
 - H tries to obtain a lock, but cannot because L is in the critical region.
 - H loops forever, and CPU control never gets handed to L.
 - As a result L never releases the lock.
- **Note: The book calls this a “priority inversion”, which is incorrect.**

Sleep/Wake

- **One solution to this problem is through the use of “Sleep/Wake” functions.**
 - When a process finds that a lock has been set (i.e. another process in the critical section), it calls “sleep” and is put into the blocked state.
 - When the other process exits the critical section and clears the lock, it can call “wake” which moves the blocked process into the READY queue for eventual execution.
- **While this sounds like an ideal solution, it can create a problem called the “producer-consumer problem”.**

Inter-Task Communications

THE PRODUCER/CONSUMER PROBLEM

The Producer/Consumer problem

```
#define N 100                                /* number of slots in the buffer */
int count = 0;                               /* number of items in the buffer */

void producer(void)
{
    int item;

    while (TRUE) {                           /* repeat forever */
        item = produce_item();               /* generate next item */
        if (count == N) sleep();             /* if buffer is full, go to sleep */
        insert_item(item);                   /* put item in buffer */
        count = count + 1;                   /* increment count of items in buffer */
        if (count == 1) wakeup(consumer);    /* was buffer empty? */
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {                           /* repeat forever */
        if (count == 0) sleep();             /* if buffer is empty, got to sleep */
        item = remove_item();                /* take item out of buffer */
        count = count - 1;                   /* decrement count of items in buffer */
        if (count == N - 1) wakeup(producer); /* was buffer full? */
        consume_item(item);                  /* print item */
    }
}
```

The Producer/Consumer problem

- **Producer and consumer share a fixed-size buffer.**
 - A global variable “count” keeps track of the number of items.
 - ✓ If $\text{count} == N$ (FULL), producer sleeps, if $\text{count} == 0$ (EMPTY) consumer sleeps.
 - After reading from the buffer, if $\text{count} == N - 1$:
 - ✓ Consumer reasons that the buffer was earlier full and wakes the producer.
 - After writing to the buffer, if $\text{count} == 1$
 - ✓ Producer reasons that the buffer was earlier empty and wakes the consumer.

The Producer/Consumer problem

- **Deadlock occurs when:**
 - Consumer checks “count” and finds it is 0.
 - Consumer gets pre-empted and producer starts up.
 - Producer adds an item, increments count to “1”, then sends a WAKE to the consumer.
 - ✓ Since consumer is not technically sleeping yet, the WAKE is lost.
 - Consumer starts up, and since count is 0, goes to SLEEP.
 - Producer starts up, fills buffer until it is full and SLEEPS.
- **Since consumer is also SLEEPing, no one wakes the producer. Deadlock.**

Inter-Task Communications

SEMAPHORES

Semaphores

- **A semaphore is a special lock variable that counts the number of wake-ups saved for future use.**
 - A value of “0” indicates that no wake-ups have been saved.
- **Two ATOMIC operations on semaphores:**
 - **DOWN, PEND or P:**
 - ✓ If the semaphore has a value of >0 , it is decremented and the DOWN operation returns.
 - ✓ If the semaphore is 0, the DOWN operation blocks.
 - **UP, POST or V:**
 - ✓ If there are any processes blocking on a DOWN, one is selected and woken up.
 - ✓ Otherwise UP increments the semaphore and returns.

Using Semaphores in the Producer/Consumer Problem

```

#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;

void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();
        down(&empty);
        down(&mutex);
        insert_item(item);
        up(&mutex);
        up(&full);
    }
}

void consumer(void)
{
    int item;

    while (TRUE) {
        down(&full);
        down(&mutex);
        item = remove_item();
        up(&mutex);
        up(&empty);
        consume_item(item);
    }
}

```

```

/* number of slots in the buffer */
/* semaphores are a special kind of int */
/* controls access to critical region */
/* counts empty buffer slots */
/* counts full buffer slots */

/* TRUE is the constant 1 */
/* generate something to put in buffer */
/* decrement empty count */
/* enter critical region */
/* put new item in buffer */
/* leave critical region */
/* increment count of full slots */

```

- EMPTY - # of empty slots.
- FULL - # of full slots.
- MUTEX – Prevents simultaneous access to the buffer.

Mutual Exclusion with Semaphores

- **When a semaphore's counting ability is not needed, we can use a simplified version called a "mutex".**
 - 1 = Unlocked.
 - 0 = Locked.
- **Two processes can then attempt to do DOWN the semaphore.**
 - Only one will succeed. The other will block.
 - When the successful process exits the critical section, it does an UP, waking the other process up.

Mutual Exclusion with Semaphores

Process A

`sema=1`

...

`non_critical_section()`

`DOWN(sema)`

`critical_section()`

`UP(sema)`

...

Process B

`non_critical_section()`

`DOWN(sema)`

`critical_section()`

`UP(sema)`

...

Mutual Exclusion with TSL/XCHG

- **We can also implement mutexes with TSL or XCHG.**
 - 0 = Unlocked, 1 = Locked.

mutex_lock:

| | |
|--------------------|---|
| TSL REGISTER,MUTEX | copy mutex to register and set mutex to 1 |
| CMP REGISTER,#0 | was mutex zero? |
| JZE ok | if it was zero, mutex was unlocked, so return |
| CALL thread_yield | mutex is busy; schedule another thread |
| JMP mutex_lock | try again later |

ok: RET | return to caller; critical region entered

mutex_unlock:

| | |
|------------------------|--------------------|
| MOVE MUTEX,#0 | store a 0 in mutex |
| RET return to caller | |

Problems with Semaphores

Deadlock

- **Our producer/consumer solution has a problem:**

■ If we swapped the semaphores for empty/full with the mutex semaphore, we have a potential deadlock:

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;

void producer(void)
{
    int item;

    while (TRUE) {
        item = produce_item();
        down(&mutex);
        down(&empty);
        insert_item(item);
        up(&mutex);
        up(&full);
    }
}
```

```
/* number of slots in the buffer */
/* semaphores are a special kind of int */
/* controls access to critical region */
/* counts empty buffer slots */
/* counts full buffer slots */
```

```
/* TRUE is the constant 1 */
/* generate something to put in buffer */
/* decrement empty count */
/* enter critical region */
/* put new item in buffer */
/* leave critical region */
/* increment count of full slots */
```

```
void consumer(void)
{
    int item;

    while (TRUE) {
        down(&mutex);
        down(&full);
        item = remove_item();
        up(&mutex);
        up(&empty);
        consume_item(item);
    }
}
```

```
/* infinite loop */
/* decrement full count */
/* enter critical region */
/* take item from buffer */
/* leave critical region */
/* increment count of empty slots */
/* do something with the item */
```

Problems with Semaphores

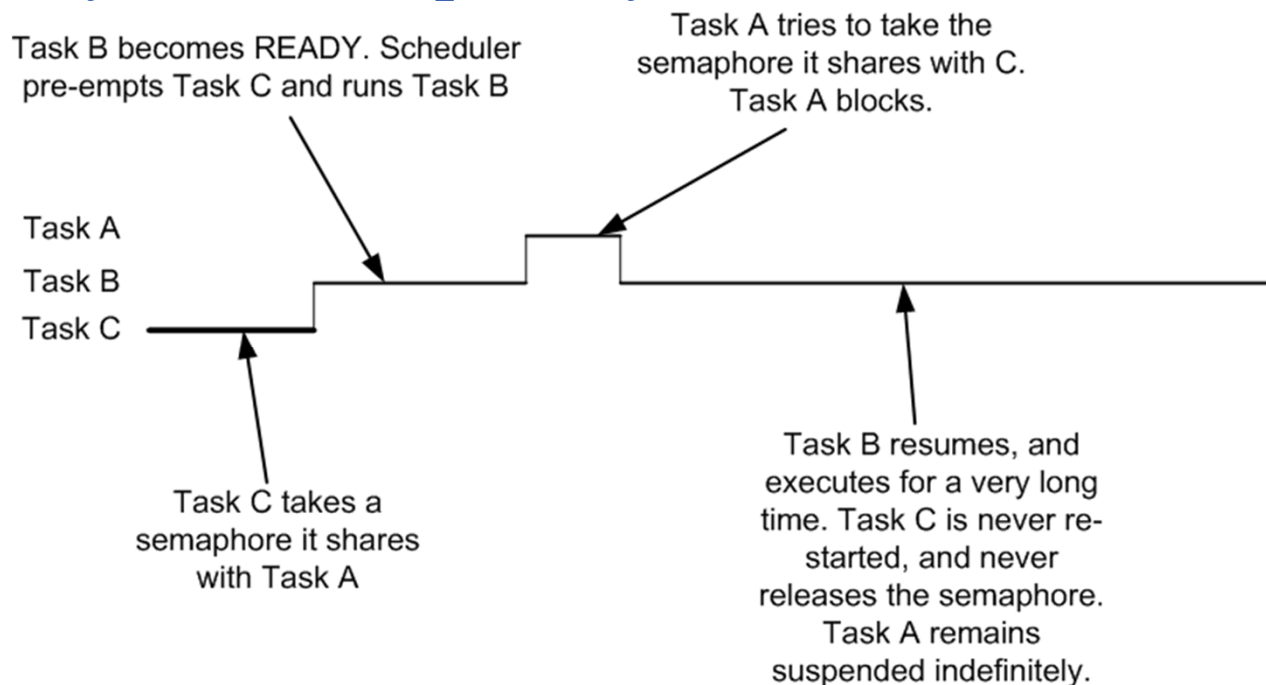
Deadlock

- **This can happen:**
 - Producer successfully DOWNs the mutex.
 - Producer DOWNs “empty”. However the queue is full so this blocks.
 - Consumer DOWNs mutex and blocks.
 - ✓ Consumer now never reaches the UP for “empty” and therefore cannot unblock the producer.
 - ✓ The producer in turn never reaches the UP for mutex and cannot unblock the consumer.
 - ✓ **Deadlock!**

Problems with Semaphores

Priority Inversion

- In the diagram on the following page, $\text{priority}(\text{Task C}) < \text{priority}(\text{Task B}) < \text{priority}(\text{Task A})$.**



- Task B effectively blocks out Task A, although Task A has higher priority!**

Inter-Task Communications

MONITORS

Monitors

- **A monitor is similar to a class or abstract-data type in C++ or JAVA:**
 - Collection of procedures, variables and data structures grouped together in a package.
 - ✓ **Access to variables and data possible only through methods defined in the monitor.**
 - However, only one process can be active in a monitor at any point in time.
 - ✓ **I.e. if any other process tries to call a method within the monitor, it will block until the other process has exited the monitor.**

Monitors

- **Implementation:**

- When a process calls a monitor method, the method first checks to see if any other process is already using it.
- If so, the calling process blocks until the other process has exited the monitor.
 - ✓ This can be achieved using mutexes or binary semaphores.
 - ✓ The mutex/semaphore operations are inserted by the compiler itself rather than by the user, reducing the likelihood of errors.

Monitors and Condition Variables

- **Monitors achieve mutual exclusion, but we also need other mechanisms for coordination.**
 - E.g. in our producer/consumer problem, mutual exclusion alone is not enough to prevent the producer from proceeding when the buffer is full.
- **We introduce “condition variables”.**
 - One process **WAITs** on a condition variable and blocks, until..
 - Another process **SIGNALs** on the same condition variable, unblocking the **WAITing** process.

Monitors and Condition Variables

- Implementing the Producer/Consumer problem with semaphores and condition variables:

- When the buffer is full ($\text{count} == N$), producer will WAIT on a full condition.
- When buffer is empty ($\text{count} == 0$), consumer will WAIT on empty.

```

monitor ProducerConsumer
  condition full, empty;
  integer count;
  procedure insert(item: integer);
  begin
    if count = N then wait(full);
    insert_item(item);
    count := count + 1;
    if count = 1 then signal(empty)
  end;
  function remove: integer;
  begin
    if count = 0 then wait(empty);
    remove = remove_item;
    count := count - 1;
    if count = N - 1 then signal(full)
  end;
  count := 0;
end monitor;

procedure producer;
begin
  while true do
    begin
      item = produce_item;
      ProducerConsumer.insert(item)
    end
  end;
procedure consumer;
begin
  while true do
    begin
      item = ProducerConsumer.remove;
      consume_item(item)
    end
  end;
end;
  
```

Monitors and Condition Variables

- **When a process encounters a WAIT, it is blocked and another process is allowed to enter the monitor.**
- **Problem:**
 - When there's a SIGNAL, the sleeping process is woken up.
 - We will potentially now have two processes in the monitor at the same time:
 - ✓ **The process doing the SIGNAL (the signaler).**
 - ✓ **The process that just woke up because of the SIGNAL (the signaled).**

Monitors and Condition Variables

- **We have 3 ways to resolve this:**
 - We require that the signaler exits immediately after calling SIGNAL.
 - We suspend the signaler immediately and resume the signaled process.
 - We suspend the signaled process until the signaler exits, and resume the signaled process only after that.

Monitors and Condition Variables

- **A condition variable is different from a semaphore.**
 - **Semaphore:**
 - ✓ If Process A UPs a semaphore with no pending DOWN, the UP is saved.
 - ✓ The next DOWN operation will not block because it will match immediately with a preceding UP.
 - **Condition variable:**
 - ✓ If Process A SIGNALs a condition variable with no pending WAIT, the SIGNAL is simply lost.
 - ✓ This is similar to the SLEEP/WAKE problem earlier on.

Monitors and Condition Variables

- **This code looks suspiciously like our original producer/consumer problem on Page 23!**

- Same issues too:

- ✓ **Page 23:**

- Consumer sees $count == 0$, and intends to SLEEP but gets pre-empted.*

- Producer sends a WAKE but the WAKE is lost.*

- In this case, if the consumer gets pre-empted before a WAIT, the corresponding SIGNAL from the producer is also lost!*


- ✓ **However we see that in this case, the mutual exclusion from the monitor prevents the SIGNAL from being lost! (WHY?)**

Inter-Task Communications


BARRIERS

Barriers


- A “**barrier**” is a special form of synchronization mechanism that works with groups of processes rather than single processes.



(a)



(b)



(c)

Barriers

- **The idea of a barrier is that all processes must reach the barrier (signifying the end of one phase of computation) before any of them are allowed to proceed.**
 - Process D reaches the end of the current phase and calls a BARRIER primitive in the OS. It gets blocked.
 - Similarly processes A and B reach the end of the current phase, calls the same BARRIER primitive and is blocked.
 - Finally process C reaches the end of its computation, calls the BARRIER primitive, causing all processes to be unblocked at the same time.

Barriers

- **Example barrier application.**
 - Computing fluid motion across a surface represented by a 1,000,000 x 1,000,000 matrix.
 - ✓ The complete matrix for iteration N must be available before computing for iteration $N+1$.
 - ✓ 1,000 individual processes each computing a part of the entire matrix.
 - ✓ We must wait for all 1,000 processes to complete finding the entire matrix before we can start on the next iteration.
 - A barrier would be very useful in achieving this.

Inter-Task Communications

COMMUNICATION MECHANISMS

Other Communication Mechanisms

- **The mechanisms we've looked at are largely for coordination.**
 - Locks, semaphores, monitors, etc.
 - Any communication takes place through global variables.
- **We will now look at several mechanisms that combine coordination with communication.**
 - Can help eliminate global variables and reduce errors.
- **Mechanisms we will look at:**
 - Queues
 - Mailboxes
 - Pipes

Inter-Task Communications

QUEUES

Queues

- **Queues have four operations:**
 - $q = \text{initq}(\text{ptr}, \text{size})$: Initializes the queue of size “size” at memory location “ptr”. Returns the queue identifier “q”.
 - $\text{enq}(q, \text{data})$: Adds a new item “data” to the tail of the queue “q”.
 - ✓ **Returns an error message when the queue is full.**

Queues

- `data = deq(q, timeout)`: Reads and removes an item from the head of the queue and returns it.
 - ✓ **Blocks if the queue is empty, up to a specific time specified by “timeout”.**
 - ✓ **If timeout is 0, deq blocks forever until there is a message in the queue.**
- `destroy(q)`: Destroys a queue. A queue, once destroyed, cannot be used unless you call `initq` again.

Queues

- **The RTOS guarantees the correctness of the queue operations.**
 - If a task gets pre-empted in the middle of an enq or deq operation, the RTOS ensures that the operation is completed correct.
 - E.g. if the second task calls enq to the same queue, the RTOS can block it until the first task completes enq properly.

Example

- **Consider a case with two tasks Task1 and Task2.**
 - Both are high priority, with urgent things to attend to.
- **If they encounter an error, they need to report this error across a network.**

Example

- **Writing to a network is typically time consuming:**
 - Hundreds of instructions to transfer data to buffers, wait for acknowledgement etc.
 - Latencies in waiting to access the network medium.
- **Makes sense to partition this into another lower priority task.**

Example

- **Task1 and Task2 must therefore share data with this error reporting task *ErrorsTask*.**
- **For our example we will assume that the queue is properly initialized.**

Queue Example

- Task Codes

```
void Task1(void)
{
    .....
    if(!!problem arises)
        vLogError(ERROR_TYPE_X);
    .....
}
```

```
Void Task2(void)
{
    .....
    if(!!problem arises)
        vLogError(ERROR_TYPE_Y);
    .....
}
```

Queue Example

```
void vLogError(int iErrorType)
{
    enq(q, iErrorType);
}
static int cErrors;
void ErrorsTask(int iErrorType)
{
    int iErrorType = deq(q, 0);
    ++cErrors;
    !! send cErrors and iErrorType across network.
}
```

Inter-Task Communications

MAILBOXES

Mailboxes

- **A mailbox is like a queue, except:**
 - The number of mailboxes within a system is typically fixed.
 - ✓ This number is initialized when your application first starts.
 - ✓ Not possible to add more mailboxes.
 - ✓ Unlike queues, which you can create as and when you need them.

Mailboxes

- A mailbox message is also often prioritized.
 - ✓ In a queue, the data is read in the order that it is written.
 - ✓ Higher priority messages in a mailbox are read ahead of lower priority messages, regardless of the order of writing.

Mailbox Operations

- **Several operations are defined on mailboxes.**
 - `initmbbox(num)`: Initializes “num” mailboxes in the system. Can only be done once.
 - `sendmsg(mbID, message, priority)`: Uses mailbox “mbID” to send a message at the given priority.
 - `rcvmsg(mbID)`: Gets the highest priority message from mbID. Blocks if there are no messages.
 - `chkmsg(mbID)`: Similar to `rcvmsg`, but returns NULL if the mailbox is empty instead of blocking.

Inter-Task Communications

PIPES

Pipes

- **A pipe is like a queue:**
 - Not prioritized. Messages are read in the order that they are written.
 - Can be created and deleted as and when required.
- **However a pipe is generally byte-oriented.**
 - A queue, for example, reads/writes an entire integer, which may consist of between 2 and 8 bytes.
 - A pipe reads/writes in units of 1 byte.
- **This can give more flexible message passing.**

Pipes

- **Pipe Operations:**
 - Standard C *fread* and *fwrite* functions.
- **Pipes are usually ready-supported by desktop operating systems.**
 - Using pipes allows your programs to be ported to desktop OSes that are adapted for RTOS use.
 - ✓ **E.g. BlueCat, a Linux derivative.**

Inter-Task Communications

SELECTING MECHANISMS

Choosing Between Message Passing Methods

- **Each method has its advantages and drawbacks:**

- **Semaphores:**

GOOD: Fast, lightweight.

BAD: Can be difficult to use correctly.

For example, you may find yourself having to set up multiple semaphores: One to coordinate between processes, another to protect data updates.

A semaphore that is taken too late or released too early can be worse than not having a semaphore at all.

Choosing Between Message Passing Methods

■ Queues:

GOOD: Simple to use.

Initialize, enqueue data, dequeue data.

All task coordination is achieved through the queue.

RTOS guarantees queue data integrity.

BAD: No message priorities, heavier than semaphores.

Need to manage queue “head” and “tail” pointers.

Need to guarantee correctness of enq and deq operations.

Choosing Between Message Passing Methods

■ Mailboxes:

GOOD: Intuitively appealing.

Each task has its own mailbox to pass it messages.

GOOD: Prioritized messaging.

Tasks can be read by priority instead of FIFO.

BAD: Even heavier than queues.

Extra code to ensure that mailbox messages are sorted by priority.

Can be very expensive for large mailboxes.

Choosing Between Message Passing Methods

■ Pipes

✓ **GOOD: Byte level mechanisms. Very flexible for passing data.**

✓ **GOOD: Supported by almost all desktop OSes**

Allows you to port your software to such OSes that have been adapted to RTOS.

✓ **BAD: Byte-level operations are often more error prone.**

Choosing Between Message Passing Methods

■ Pipes

✓ **BAD: Byte-level operations are slow.**

CPUs transfer in fixed numbers of bytes called the “word size”.

Byte level operations would mean chopping up and joining words.

✓ **BAD: No priorities.**

Choosing Between Message Passing Methods

- **The RTOS manuals will list the sizes and expected execution times for each mechanism.**
- **It is critical that you choose the best mechanism for your application.**
 - Do you need maintainability, priority, execution speed or small footprint?
- **RTOSs sometimes allow you to remove mechanisms that you don't use.**
 - Might think about sticking with just one mechanism.

Inter-Task Communications

POTENTIAL PITFALLS

Potential Pitfalls

- **While queues, mailboxes and pipes simplify data sharing, they also make it easy to insert bugs into your system**
- **For readability, the pitfalls listed talk about queues. Unless otherwise stated, they apply to mailboxes and pipes as well.**

Potential Pitfalls

1. Reading/Writing to the wrong queue.

- RTOSs don't have the ability to restrict a queue to a particular task.
- If one task writes temperature data to the error messages queue, then the error handling task will get confused.

2. Wrong Queue Semantics

- Task A writes integers to a queue.
- Task B reads the queue and interprets the data as a pointer to an integer, instead of the integer itself.

Potential Pitfalls

```
OSEVENT *pq;
// Prototype for the enq and deq functions
void enq(OSEVENT *q, void *data);
void *deq(OSEVENT *q);

void taskA(void)
{
    int datum;
    ... some processing on datum ...
    // Enqueue an integer "datum", casting it as
    // void * to fit the enq prototype.
    enq(pq, (void *) datum);
}
```

Potential Pitfalls

```
void taskB(void)
{
    int *dataptr;

    dataptr = (int *) deq(pq);
    *dataptr = !! Result of some computation.
}
```

- This code is a disaster because taskB is going to (almost) randomly overwrite memory that is used by other tasks.
 - Remember that there is no memory protection in most RTOS!

Potential Pitfalls

3. Running out of space.

- This is a disaster because the data being passed is not usually optional.

4. Inadvertently creating unprotected shared variables.

- This happens when you pass pointers instead of data.
- If you pass data, the RTOS ensures that the enq and deq operations place and remove the data correctly from the queue.
- If you pass pointers, the RTOS ensures that the enq and deq operations place and remove the pointers correctly from the queue.
 - BUT: The data the pointers point to are themselves not protected!

Potential Pitfalls

```
OSEVENT *pq;
// Prototype for the enq and deq functions
void enq(OSEVENT *q, void *data);
void *deq(OSEVENT *q);

void vReadTemperaturesTask(void)
{
    static int iTemperatures[2];
    while(1)
    {
        iTemperatures[0] = !! Read in temp1 from hw
        iTemperatures[1] = !! Read in temp2 from hw

        // Add address of iTemperatures to queue
        enq(pq, (void *) iTemperatures);
    }
}
```

Potential Pitfalls

```
void vMainTask(void)
{
    int *pTemp;

    while(1)
    {
        pTemp = (int *) deq(pq);
        if(pTemp[0] != pTemp[1])
            !! Set of howling alarm
    }
}
```

Potential Pitfalls

- i. *vMainTask* will get the address of *iTemperatures* from the queue, assigning it to *pTemp*.
- ii. It then reads *pTemp[0]* (which is *iTemperatures[0]* since *pTemp* points to the same *iTemperatures* array).
- iii. *vMainTask* gets pre-empted, and *vReadTemperaturesTask* updates *iTemperatures[0]* and *iTemperatures[1]*.

Potential Pitfalls

- iv. *vMainTask* resumes, and reads in *pTemp[1]*.
- v. The comparison will fail, causing the alarm to be triggered erroneously.

Summary

- **In this lecture we looked at:**
 - What race conditions and critical sections are.
 - Mechanisms to prevent race conditions.
 - Mechanisms that provide both coordination and communication functions.