

CG2271 Real Time Operating Systems

Tutorial 9

Question 1

Assume that messages in your RTOS consist of void pointers, that sendmsg places the void pointer passed to it on a queue, and that rcvmsg returns the void pointer it retrieved from the queue. What is wrong with the following code?

```
void vLookForInputTask(void)
{
    ...
    if(!! A key has been pressed on the keyboard)
        vGetKey();
}

void vGetKey(void)
{
    char ch;
    ch = !! Get key from the keyboard;
    /* Send key to keyboard command handler task */
    sndmsg(KEY_MBOX, &ch, PRIORITY_NORMAL);
}

void vHandleKeyCommandsTask(void)
{
    char *p_chLine;
    char ch;
    while(1)
    {
        /* Wait for key to be received */
        p_chLine = rcvmsg(KEY_MBOX, WAIT_FOREVER);
        ch = *p_chLine;
        !! Do stuff with ch
    } /* while */
}
```

1. vGetKey passes the address of ch to the mailbox. But ch is a local variable that is destroyed once vGetKey exits.

3. As a result this de-reference might access rubbish!

2. When vHandleKeyCommandsTask gets the address to ch, the memory location used to hold ch might already be overwritten by something else

Question 2

This code uses the following AMX functions to create and use events:

```
ajevcre(AMXID *p_amxidGroup, unsigned int uValueInit,  
        char *p_chTag);
```

- Creates a group of 16 events, and returns a handle to the events in *p_amxidGroup*. The events group is also given a four-character tag given in *p_chTag*. The events are initialized to “set” or “reset” depending on the value of *uValueInit*.

```
ajevsig(AMXID amxidGroup, unsigned int uMask,  
        unsigned int uValueNew);
```

- Sets and resets the events in the group indicated by *amxidGroup*. The *uMask* parameter indicates which of the 16 events in the group should be set or reset, depending on the parameter *uValueNew*.

```
ajevwat(AMXID amxidGroup, unsigned int uMask, unsigned int uValue,  
        int iMatch, long lTimeOut);
```

- Causes the task to wait for one or more events in the group indicated by *amxidGroup*. The *umask* parameter indicates which of the 16 events in the group the task should wait for. *uValue* indicates whether the task wants to wait for the selected events to be “set” or “reset”, while *iMatch* specifies whether the task should unblock when ALL or AT LEAST ONE of the indicated events is set or reset. *lTimeOut* indicates how long the task is willing to wait.

Given the information above, rewrite the following code with semaphores.

```
/* Handle for trigger group of events */  
AMXID amxidTrigger;  
  
/* Constants for use in the group */  
#define TRIGGER_MASK      0x0001  
#define TRIGGER_SET       0x0001  
#define TRIGGER_RESET     0x0000  
  
void main(void)  
{  
    .....  
    /* Create an event group with the trigger and keyboard events set */  
    ajevcre(&amxidTrigger, 0, "EVTR");  
    .....  
}  
  
void interruptvTriggerISR(void)  
{  
    /* User pulled trigger. Set the event */  
    ajevsig(amxidTrigger, TRIGGER_MASK, TRIGGER_SET);  
}
```

```

void vScanTask(void)
{
    .....
    while(1)
    {
        /* Wait for user to pull trigger */
        ajevwait(amxidTrigger, TRIGGER_MASK, TRIGGER_SET,
                WAIT_FOR_ANY, WAIT_FOREVER);

        /* Reset the trigger event */
        ajevsig(amxidTrigger, TRIGGER_MASK, TRIGGER_RESET);
        !! turn on hardware scanner
    } /* while */
}

```

Re-written Code:

```

OSSemaphore sem;
void main(void)
{
    .....
    OSSemCreate(&sem, 0); /* Initsemaphore to
    0 */
    .....
}

void interrupt interruptvTriggerISR(void)
{
    /* User pulled trigger. Set the event */

    OSSemRelease(&sem); /* Release the semaphore */
}

void vScanTask(void)
{
    .....
    while(1)
    {
        /* Wait for user to pull trigger */
        OSSemTake(&sem); /* Try to take a semaphore */
        !! turn on hardware scanner
    } /* while */
}
.....

```

Question 3

What does the term “re-entrancy” mean? In particular, what does it mean when we say that a routine is “re-entrant”? Demonstrate, with an example, the problems associated with a non re-entrant routine, and how the routine can be made re-entrant. GIYF. ☺

Consider a routine AddOne:

```
void AddOne()
{
    statusCount+=1;
}
```

In assembly this would be broken down to:

```
LOAD R0, statusCount      ; Load contents of statusCount to R0
ADD R0, R0, 1             ; Increment R0 by 1
STORE statusCount, R0      ; Write updated value to statusCount
```

This routine is called by both Task1 and Task2. It is possible for Task 1 to get pre-empted while it is still inside this routine, and Task 2 gets started up. Furthermore it is possible that Task 2 also calls AddOne before Task1 completely exits it (because Task1 got pre-empted). A routine is “re-entrant” if, in this situation, it behaves exactly as though Task1 had fully exited AddOne before Task2 entered.

We can demonstrate that the example above is not re-entrant. Suppose statusCount is currently 3:

Description	Task1 code exec	Task2 code exec	Value
Task 1 calls AddOne	LOAD R0, statusCount	-	statusCount=3, R0=3
	ADD R0, R0, 1	-	statusCount=3, R0=4
Task1 gets pre-empted.			
Task2 calls AddOne	-	LOAD R0, statusCount	statusCount=3, R0=3
	-	ADD R0, R0, 1	statusCount=3, R0=4
	-	STORE statusCount, R0	statusCount=4, R0=4
Execution returns to Task1			
	STORE statusCount, R0	-	statusCount=4, R0=4

Note that the final result of statusCount=4 is wrong, because Task1 would’ve updated statusCount to 4, and then Task2 should’ve updated it to 5.

We can make addOne re-entrant by requiring a task to acquire a semaphore before entering the task. Rewrite AddOne to:

```
OSSemaphore addsema=1;
```

```
void AddOne
{
    OSGetSema(&addsema);
    statusCount+=1;
    OSReleaseSema(&addsema);
}
```

Suppose this compiles to:

```
!! Get semaphore addsema
LOAD R0, statusCount          ; Load contents of statusCount to R0
ADD R0, R0, 1                 ; Increment R0 by 1
STORE statusCount, R0         ; Write updated value to statusCount
!! Release semaphore addsema
```

The execution trace is now:

Description	Task1 code exec	Task2 code exec	Value
Execution begins			addsema=1
Task 1 calls AddOne	!! Get semaphore	-	addsema=0
	LOAD R0, statusCount	-	statusCount=3, R0=3
	ADD R0, R0, 1	-	statusCount=3, R0=4
Task1 gets pre-empted.			
Task2 calls AddOne	-	!! Get semaphore	addsema=0
Task2 is blocked. Control returns to Task1			
	STORE statusCount,R0	-	statusCount=4, R0=4
	!! Release semaphore	-	addsema=1
Control is handed back to Task2, addsema=0			
	-	LOAD R0, statusCount	statusCount=4, R0=4
	-	ADD R0, R0, 1	statusCount=4, R0=5
	-	STORE statusCount, R0	statusCount=5 R0=5
	-	!!Release semaphore	addsema=1

We now get the correct result of statusCount=5, and AddOne is now re-entrant.

Question 4

Is this function re-entrant?

```
int strlen(char *p_sz)
{
    int iLength;

    iLength = 0;
```

```

while(*p_sz != '\0')
{
    ++iLength;
    ++p_sz;
}

return iLength;
}

```

Yes and no.

YES:

On its own, the variables *iLength* and *p_sz* are local, and hence even if *strlen* gets pre-empted halfway, and a second task calls *strlen*, it works fine. Each call to *strlen* will have its own copy of *iLength* and *p_sz* on it's the calling task's stack. Since the RTOS manages each task's stack separately, there will therefore not be a conflict.

NO:

The problem is that *p_sz* may point to a buffer that is shared between *Task1* and *Task2*. When *Task1* is pre-empted halfway through *strlen* and *Task2* is started, it's possible that *Task2* changes the buffer that *p_sz* points to. When *Task1* resumes, *strlen* is really now counting the characters of a different string than when it first started.

Question 5

Consider the statement: “In a non-preemptive RTOS, tasks cannot interrupt each other. Therefore there are no data sharing problems amongst task.” Would you agree with this?

No. Consider two tasks Task A and Task B. Task B depends on Task A computing a particular value in *sharedData* before Task B can start its job. Task A meanwhile depends on an event occurring before it can compute *sharedData*:

Task A:

!!Does some stuff:

!! Wait on Event E

!! Compute *sharedData*

If Event E has not occurred, Task A BLOCKs while Task B starts.

Task B:

!! Use *sharedData*; /* Wrong! Task A hasn't computed *sharedData* yet! */

Hence there is still some need to coordinate between two tasks. There is also an issue of interrupt handlers corrupting data that Task A and/or Task B are sharing with it.