

CS2020

# Data Structures and Algorithms

**Welcome!**

# Introductions

---

Seth Gilbert

- Assistant Professor in the CS Department
- Office: COM2-3-23
- Office hours: TBA
- Interests: distributed algorithms

# Introductions

---

Steven Halim

- Lecturer in the CS Department
- Office: COM2-2-60
- Office hours: TBA
- Interests: competitive programming, visualization techniques, ICPC, IOI

# Introductions

---

Tutors: James Yong Kim Leng

- Foodie (Live To Eat)
- Loves:  
engineering/reverse-engineering/hacking
- Very approachable  
although don't usually  
appear to be
- Game programmer
- Favorite IDE: Eclipse



# Introductions

---

Tutors: Jon Hay

- Year 4 undergrad
- Double degree (!) in CS and applied math
- Game development
- Soft spot for: hamsters



# Introductions

---

Tutors: Ngo Minh Duc

- Year 4 undergrad in computer science
- Took CS1101s
- Likes: algorithms and mathematics



# Introductions

---

Tutors: Nguyen Hoanh Tien

- Year 3 undergrad in computer engineering
- Took CS1101s
- Likes: swimming and kite flying



# Introductions

---

## Tutors: Zhao Cong

- Competitive programmer since senior middle school
- Just finished internship at PayPal
- Hobbies: reading, blogging, F1, tennis, movies





# Introductions

---

Tutors: Koh Zi Han

- Mysterious and secretive



# Introductions



# Algorithms

---

What is an *algorithm*?

- Set of instructions for solving a problem
  - "First, wash the tomatoes."
  - "Second, peel and cut the carrots."
  - "Third, mix the olive oil and vinegar."
  - "Finally, combine everything in a bowl."
- *Finite* sequence of steps
- Unambiguous
- English, Chinese, pseudocode, Java, etc.

# Algorithms

---

## History

- Named for al-Khwārizmī (780-850)
  - Persian mathematician
- Many ancient algorithms
  - Multiplication: Rhind Papyrus
    - Babylon and Egypt: ~1800BC
  - Euclidean Algorithm: Elements
    - Greece: ~300BC
  - Sieve of Eratosthenes
    - Greece: ~200BC



# Algorithms

---

“If you need your software to run twice as fast, hire better programmers.

But if you need your software to run more than twice as fast, use a better **algorithm**.”

-- *Software Lead at Microsoft*

# Software

---

Desirable features?

- Speed / Performance
- ???

# Software

---

## Desirable features?

- Speed / Performance
- Correctness / lack of bugs
- Memory usage
- Easy to maintain / easy to read
- Modular
- Completed on schedule
- Elegant
- Portable

# Algorithms

---

## Goals of this course:

- How to organize and manipulate data?
  - Efficiency
    - **Time**: *How long does it take?*
    - **Space**: *How much memory? How much disk?*
    - **Others**: Energy, power, heat, parallelism, etc.
  - Scalability
    - Inputs are *large* : e.g., the internet.
    - Bigger problems consume more resources.
- Solve real (fun!) problems



# Algorithms

---

## Goals of this course:

- Discover existing “toolbox” of algorithms and data structures that you can use to solve real world problems.
- Learn how to choose the right algorithm for the right problem.
- Learn how to design and analyze new algorithms and data structures when needed.

# Algorithms

---

## How to solve a problem:

- Identify the problem
  - Ex: what's the fastest way to get to NUS?
- Abstract irrelevant details
  - Ex: traffic+lights+merging+speed=time
- Find good algorithms
- Implement (in Java)
- Evaluate
  - How fast? How does it scale?

# Semester Overview

---

- Topic 1: Linked data structures
  - Arrays
  - Searching
  - Sorting
  - Lists, Stacks, Queues
  - Divide-and-Conquer
- Example problems: document distance, peak finding

# Semester Overview

---

- Topic 2: Trees
  - Binary Search Trees
  - Balanced Trees
  - Priority Queues
  - Heaps
- Example problems: simple scheduling

# Semester Overview

---

- Topic 3: Hash Tables
  - Dictionaries
  - Hash functions
  - Chaining
  - Amortized Analysis
- Example problems: DNA similarity

# Semester Overview

---

- Topic 4: Graphs
  - Searching in a graph
  - Spanning trees
  - Shortest paths
- Example problems: Google map routes

# Semester Overview

---

- Topic 5: Dynamic Programming
  - All-pairs shortest paths
  - Floyd-Warshall
  - Travelling Salesman Problem

# Java

---

- Goal: Learn Java
  - Quick overview in lectures...
  - More details in recitations / discussion groups.
  - Learn on your own...
  - Solve problem sets.



# Why Java?

---

- Hardware: Assembly language
- Procedural (imperative) languages:
  - Fortran, COBOL, BASIC, Pascal , C
- Functional languages:
  - IPL, Lisp, Scheme, Haskell
- Declarative languages:
  - SQL, Lex/Yacc
- Object-oriented languages:
  - Simula 67, Smalltalk 80, C++, Java, C#, Python?

# Why Java?

---

- Good aspects:
  - Common in industry / real-world / web
  - Modularity / Abstraction via OOP
  - Avoids memory leak issues of C/C++
- Less good aspects:
  - Performance?? (compare to: C++)
  - Elegance?? (compare to: Scheme)

# Language Does Not Matter

---

- Algorithms are more important:
  - Fact: C can be 20x as fast as Python!  
(Source: MIT 6.006, Spring 2008, Lecture 2)
  - Fast sorting in Python (merge-sort):
    - Time:  $T(n) = 2n \log(n) \mu s$
    - Total time for 10,000 elements: 0.266s
  - Slow sorting in C (insertion-sort):
    - Time:  $T(n) = 0.01n^2$
    - Total time for 10,000 elements: 1s

# Administrative Details

---

- Weekly schedule:
  - Two lectures: Tues/Fri 10am-12pm
  - One recitation: Friday
  - One discussion group: Thursday
- Weekly work:
  - Problem set
  - Discussion group problems
  - Occasional bonus problems

# Administrative Details

---

## Discussion Groups:

- Register via CORS (known as: labs)
- Fill out preference form here!

## Tutorials:

- Register via CORS
- Three slots (in COM2-105):
  - 2pm,
  - 3pm
  - 4pm

# Administrative Details

---

- Quizzes:
  - Quiz 1 - Feb. 11 (15%)
  - Practical Programming Quiz – Mar. 3 (10%)
  - Quiz 2 – Mar. 25 (15%)
- Final Exam:
  - Apr. 25 (40%)
- Remainder: problem sets and participation

# Administrative Details

---

Course website:

[cs2020.ddns.nus.edu.sg](http://cs2020.ddns.nus.edu.sg)


Go register!

- Same infrastructure as CS1101s
  - Experience points...
  - Levels...
  - Facebook connect...
- But no comic strip, no achievements, etc.

# Administrative Details

---

- Textbook: Introduction to Algorithms
  - Cormen, Leiserson, Rivest, Stein




- Recommended...



# Administrative Details

---

- Textbook: Java Foundations
  - Lewis, DePasquale, Chase




- Optional...

# Administrative Details

---

- Textbook: Competitive Programming
  - Halim



- Optional...

# Problem Sets

---

- Different types of questions:
  - Give an algorithm...
  - Show that an algorithm is correct...
  - Analyze the performance of an algorithm...
  - Implement an algorithm...
  - Measure the performance of a solution...

# Problem Sets

---

“Give an algorithm...”

“Implement an algorithm...”

Five parts to each answer:

1. State the problem being solved.
2. Describe the solution in words. (If it is a Java algorithm, describe the algorithm being implemented.)
3. Give the algorithm/Java.
4. Explain why it works.
5. (Optional) Analyze its performance.

# Problem Sets

---

If the tutor does not understand your solution, then it will not be graded.

The tutor may ask you to explain your code/algorithm better.

The tutor is not a compiler.

# Problem Sets

---

## Collaboration Policy

- Working together is strongly encouraged!
- You must write-up your problems sets alone.
- You must list on your submission the name of everyone you worked with, and all sources used.
- Cheating / plagiarism will be dealt with harshly.

# Clickers...

Who is your favorite author?

- a) Shakespeare
- b) J.K. Rowling
- c) Confucius
- d) Homer Simpson



# Clickers...

Today:

Sign up for a clicker.

If your clicker is missing:

\$89 / clicker.





# Clickers...

1. Simply choose your response from the keypad buttons.
  2. The light will go **GREEN** to confirm your response has been received.
  3. You can **change your answer** by simply keying in your new choice.
- (The system will only count the last vote.)

## NOTE:

Please DO NOT press the **GO** button, this will change the Radio Frequency of the Keypad.



# Break time

---

Questions?

Sign up for a clicker...

Sign up for a “Discussion Group”...


# Today

---

- Problem: Document Distance
  - How similar are two documents?
- Solution:
  - Algorithm idea
  - Java implementation
  - Performance measurement

# Where are you from?


1. China
2. Europe
3. India
4. Japan
5. Malaysia
6. Singapore
7. United States
8. Vietnam
9. Other...



Are you:

1. Male
2. Female
3. Other


0 of 5



# Who's your favorite Harry Potter character?

- 0% 1. Harry Potter
- 0% 2. Hermione
- 0% 3. Ron
- 0% 4. Snape
- 0% 5. Voldemort
- 0% 6. Gandolf
- 0% 7. What?? Who's that?

# Who's your favorite Harry Potter character?




# Today

---

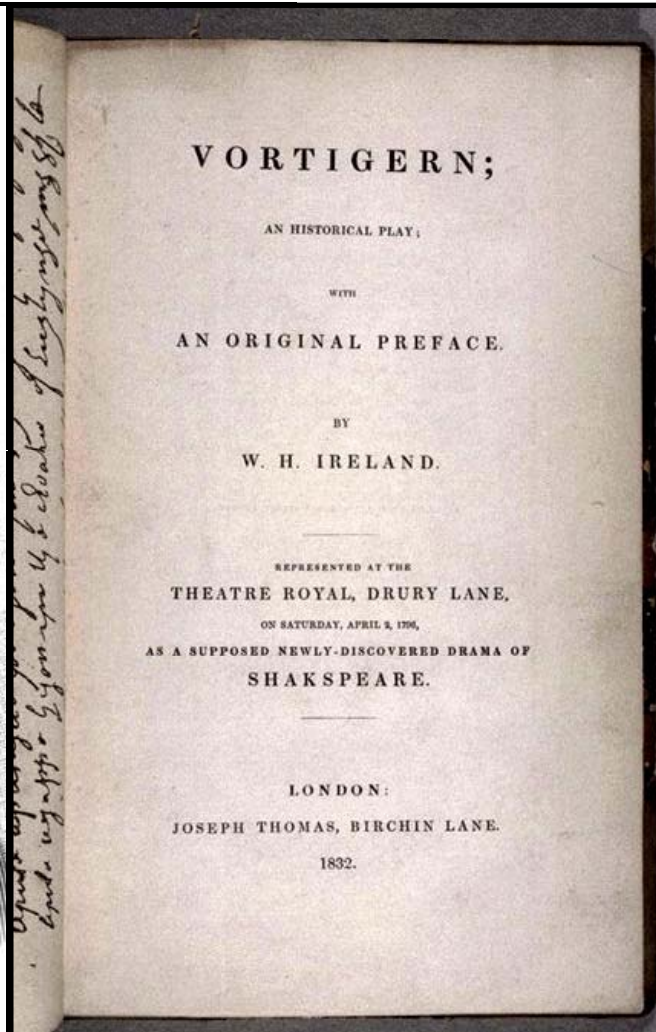
- Problem: Document Distance
  - How similar are two documents?
- Solution:
  - Algorithm idea
  - Java implementation
  - Performance measurement



# Who wrote this?



**William  
Shakespeare??**



***mystery* play  
“found” in 1796**



**William Henry  
Ireland??**

# Document distance

---

- How similar are two documents?
  - Are two documents written by the same author?
  - Detect forgeries
  - Find plagiarism / cheating
  - Was Homer one author or many?
- What does “similar” mean?

# Metrics of similarity

---


- Binary: (e.g., detect plagiarism)
  - Exactly same words in same order
- Scalar:
  - Number of words in the same order
  - Number of shared *uncommon* words
  - Same # of words per sentence
  - Same ratio of adjectives / nouns
  - Written on similar paper / using similar ink

# Vector Space Model

---

## Strategy:

- View each document as a high-dimensional vector.
- The *metric of similarity* is the angle between the two vectors.



- Identical:  $\Phi = 0$
- No words in common:  $\Phi = \pi/2$

[Salton, Wang, Yang '75]

# Vector Space Model

---

Document as vector:

Example 1:

“to be or not to be” = [2,1,1,2]

be	not	or	to
2	1	1	2

# Vector Space Model

---

Document as vector:

Example 1:

“to be or not to be” =  $[0, 2, 0, 1, 0, 1, 2]$

<b>afraid</b>	<b>be</b>	<b>greatness</b>	<b>not</b>	<b>of</b>	<b>or</b>	<b>to</b>
<b>0</b>	<b>2</b>	<b>0</b>	<b>1</b>	<b>0</b>	<b>1</b>	<b>2</b>

# Vector Space Model

---

Example 1:

“to be or not to be” =  $[0, 2, 0, 1, 0, 1, 2]$

Example 2:

“be not afraid of greatness” =  $[1, 1, 1, 1, 1, 0, 0]$

afraid	be	greatness	not	of	or	to
1	1	1	1	1	0	0

## Example 3: “to be afraid, to be not afraid”

0 of 5

0% a. 1, 1, 1, 3, 0, 0, 2

0% b. 2, 2, 0, 1, 0, 0, 2

0% c. 3, 2, 3, 1, 1, 1, 1

0% d. I have no idea.

afraid	be	greatness	not	of	or	to
?	?	?	?	?	?	?



# Vector Space Model

---

Example 3: “to be afraid, to be not afraid”

1. [1, 1, 1, 3, 0, 0, 2]
2. [2, 2, 0, 1, 0, 0, 2]
3. [3, 2, 3, 1, 1, 1, 1]
4. I have no idea.

afraid	be	greatness	not	of	or	to
?	?	?	?	?	?	?

# Vector Space Model

---

Dot Product:

$$v = [v_1, v_2, v_3, v_4]$$

$$w = [w_1, w_2, w_3, w_4]$$

$$v \cdot w = v_1 w_1 + v_2 w_2 + v_3 w_3 + v_4 w_4$$

# Vector Space Model

---

Dot Product:

$$v = [v_1, v_2, \dots, v_n]$$

$$w = [w_1, w_2, \dots, w_n]$$

$$v \cdot w = \sum v_i w_i$$

## Dot Product Question:

$$v = [0, 2, 0, 1]$$

$$w = [1, 1, 1, 1]$$

$$(v \cdot w) =$$

0% a. 1

0% b. 2

0%  3

0% d. 4

0% e. 5

# Vector Space Model

---

Norm of a vector (L2 norm):

$$|v| = \text{SQRT}(v \bullet v)$$

Example: distance between two points

$$|(x_1, y_1) - (x_2, y_2)| = \sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$$

# Vector Space Model

---

Norm of a vector (L2 norm):

$$|v| = \sqrt{v \cdot v}$$

$$|v| = \sqrt{\sum_{i=1}^n v_i \cdot v_i}$$

Example:  $\text{NORM}(3, 0, 4, 0) =$   
 $\text{SQRT}(3*3 + 0*0 + 4*4 + 0*0) = 5$

# Vector Space Model

---

Law of cosines:

$$\Theta(v, w) = \cos^{-1} \left( \frac{v \cdot w}{\|v\| \cdot \|w\|} \right)$$

Notes:

- $\Phi$  is an angle between  $(0, \pi)$
- If  $(v=w)$ , then  $\Phi=0$ .
- If  $(v \cdot w) = 0$ , then  $\Phi=\pi$ .

# Compare Two Documents

Given: documents A and B

1. Create vectors  $v_A$  and  $v_B$
2. Calculate norm:  $|v_A|$
3. Calculate norm:  $|v_B|$
4. Calculate dot product:  $(v_A \cdot v_B)$
5. Calculate angle  $\Phi(v_A, v_B)$



# Performance Profiling

---

*(Dracula vs. Lewis & Clark)*

Step	Function	Running Time
Create vectors:	Read each file	1,824.00s
	Parse each file	0.20s
	Sort words in each file	328.00s
	Count word frequencies	0.31s
Dot product:		6.12s
Norm:		3.81s
Angle:		6.56s
<b>Total:</b>		<b>72minutes <math>\approx</math> 4,311.00s</b>

# Eclipse-TPTP

Profiling and Logging - CS2020 Test/src/sg/edu/nus/cs2020/DocumentDistanceMain.java - Eclipse Platform

File Edit Source Refactor Navigate Search Project Run Window Help

Navigator Profiling Monitor

Execution Statistics

Execution Statistics - sg.edu.nus.cs2020.DocumentDistanceMain at gilbert-d960 [ PID: 7180 ]

### Session summary

#### Highest 10 base time

Package	Base Time (seconds)	Average Base Time (seconds)	Cumulative Time (seconds)	Calls
sg.edu.nus.cs2020	5,003.303381	0.012981	5,003.303381	38...
VectorTextFile	4,311.153603	215.557680	4,311.986191	20
ReadFile(java.lang.String) java.lang.String	3,648.053534	1,824.026767	3,648.053534	2
InsertionSortWords() void	656.330413	328.165206	656.330413	2
DotProduct(sg.edu.nus.cs2020.VectorTextFile,	6.134406	2.044802	6.533115	3
SplitString(java.lang.String) void	0.390386	0.195193	0.390386	2
CountWordFrequencies() void	0.185574	0.092787	0.619453	2
VerifySort() void	0.034114	0.017057	0.034114	2
Angle(sg.edu.nus.cs2020.VectorTextFile, sg.ec	0.024622	0.024622	6.557860	1
VectorTextFile(java.lang.String)	0.000273	0.000136	4,305.428330	2
ParseFile(java.lang.String) void	0.000158	0.000079	3,648.444078	2
Norm() double	0.000123	0.000062	3.814559	2
VectorTextFile2	676.500618	33.825031	680.487998	20
WordCountPair	6.706844	0.000021	6.706844	31...
VectorTextFile3	6.594781	0.000101	8.481658	65...

Session summary Execution Statistics Call Tree Method Invocation Details Method Invocation

Console Problems

<terminated> DocumentDistanceMain [Java Application] java.exe (January 5, 2011 3:59:34 PM)

The angle between A and B is: 0.5708476330610679

The angle between A and B is: 0.5708476276825866

The angle between A and B is: 0.5708476276825866

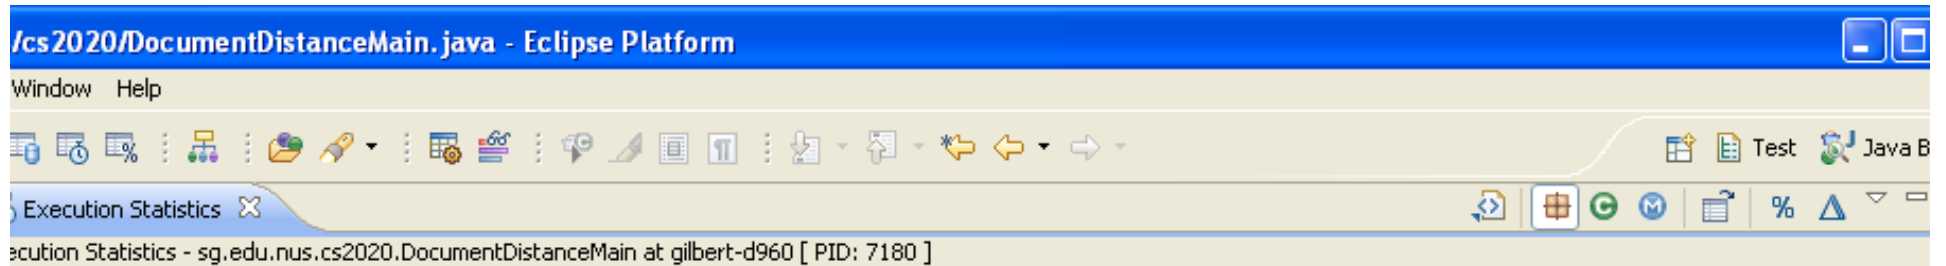
Test.java VectorTextFile2.java VectorTextFile.java DocumentDistanceMain hamlet.txt midsummer.txt Tom Sawyer.txt JFK.txt verne.txt 13

```
package sg.edu.nus.cs2020;

import java.io.IOException;

public class DocumentDistanceMain
```

# Eclipse-TPTP



## Session summary

### Highest 10 base time

Package	<	Base Time (seconds)	Average Base Time (seconds)	Cumulative Time (seconds)	Calls
sg.edu.nus.cs2020		5,003.303381	0.012981	5,003.303381	38...
VectorTextFile		4,311.153603	215.557680	4,311.986191	20
ReadFile(java.lang.String) java.lang.String		3,648.053534	1,824.026767	3,648.053534	2
InsertionSortWords() void		656.330413	328.165206	656.330413	2
DotProduct(sg.edu.nus.cs2020.VectorTextFile, sg.edu.nus.cs2020.VectorTextFile) double		6.134406	2.044802	6.533115	3
SplitString(java.lang.String) void		0.390386	0.195193	0.390386	2
CountWordFrequencies() void		0.185574	0.092787	0.619453	2
VerifySort() void		0.034114	0.017057	0.034114	2
Angle(sg.edu.nus.cs2020.VectorTextFile, sg.edu.nus.cs2020.VectorTextFile) double		0.024622	0.024622	6.557860	1
VectorTextFile(java.lang.String)		0.000273	0.000136	4,305.428330	2
ParseFile(java.lang.String) void		0.000158	0.000079	3,648.444078	2
Norm() double		0.000123	0.000062	3.814559	2
VectorTextFile2		676.500618	33.825031	680.487998	20
WordCountPair		6.706844	0.000021	6.706844	31...
VectorTextFile3		6.594781	0.000101	8.481658	65,...
VectorTextFile4		0.000000	0.000000	0.000000	0

# Performance Profiling

---

*(Dracula vs. Lewis & Clark)*

Step	Function	Running Time
Create vectors:	Read each file	1,824.00s
	Parse each file	0.20s
	Sort words in each file	328.00s
	Count word frequencies	0.31s
Dot product:		6.12s
Norm:		3.81s
Angle:		6.56s
<b>Total:</b>		<b>72minutes <math>\approx</math> 4,311.00s</b>

# Performance Profiling

---

*(Dracula vs. Lewis & Clark)*

Version	Change	Running Time
Version 1		4,311.00s
Version 2	Better file handling	676.50s
Version 3	Faster sorting	6.59s
Version 4	No sorting!	2.35s

- Version 4 will be released later in the semester...

# Performance Profiling

---

*(Dracula vs. Lewis & Clark)*

Step	Function	Running Time
Create vectors:	Read each file	1,824.00s
	Parse each file	0.20s
	Sort words in each file	328.00s
	Count word frequencies	0.31s
Dot product:		6.12s
Norm:		3.81s
Angle:		6.56s
<b>Total:</b>		<b>72minutes <math>\approx</math> 4,311.00s</b>

# ReadFile (excerpt)

```
// Open the file as a stream and find its size
InputStream = new FileInputStream(fileName);
iSize = inputStream.available();

// Read in the file, one character at a time, normalizing as we go.
for (int i=0; i<iSize; i++)
{
    // Read a character
    char c = (char)inputStream.read();

    // Ensure that the character is lower-case
    c = Character.toLowerCase(c);

    // Check if the character is a letter
    if (Character.isLetter(c))
    {
        strTextFile = strTextFile + c;
    }
    // Check if the character is a space or an end-of-line marker
    else if ((c == ' ') || (c == '\n')) && (!strTextFile.endsWith(" "))
    {
        strTextFile = strTextFile + ' ';
    }
}
```

# String Problem!

---

What happens when:

➤ `strTextFile = strTextFile + c`

1. Creates new temporary string.
2. Copies `strTextFile` to the new string.
3. Adds the new character `c`.
4. Reassigns `strTextFile` to point to the new string.



# String Problem!

---

What happens when:

➤ `strTextFile = strTextFile + c`

1. Creates new temporary string.
- 2. Copies `strTextFile` to the new string.**
3. Adds the new character `c`.
4. Reassigns `strTextFile` to point to the new string.

Copying a string of `k` characters takes time  $O(k)$ !

# How long does it take to read a file containing $n$ characters?

1.  $O(n)$
2.  $O(n \log n)$
3.  $O(n^2)$
4.  $O(2^n)$
5. Big-O notation?

```
// Open the file as a stream and find its size
InputStream = new FileInputStream(fileName);
iSize = inputStream.available();

// Read in the file, one character at a time, normalizing as
for (int i=0; i<iSize; i++)
{
    // Read a character
    char c = (char)inputStream.read();

    // Ensure that the character is lower-case
    c = Character.toLowerCase(c);

    // Check if the character is a letter
    if (Character.isLetter(c))
    {
        strTextFile = strTextFile + c;
    }
    // Check if the character is a space or an end-of-line ma
    else if ((c == ' ') || (c == '\n')) && (!strTextFile.end
    {
        strTextFile = strTextFile + ' ';
    }
}
```

0%

0%

0%

0%

0%

1

2

3

4

5

# String Problem!

---

How long to read in a file of  $n$  characters?.

$$1 + 2 + 3 + 4 + \dots + n = n(n+1)/2 = \Theta(n^2)$$

Very, very, very slow!

# Fix the string problem!

```
// Open the file as a stream and find its size
InputStream = new FileInputStream(fileName);
iSize = inputStream.available();

// Initialize the char buffer to be arrays of the appropriate size.
charBuffer = new char[iSize];

// Read in the file, one character at a time, normalizing as we go.
for (int i=0; i<iSize; i++)
{
    // Read a character
    char c = (char)inputStream.read();

    // Ensure that the character is lower-case
    c = Character.toLowerCase(c);

    // Check if the character is a letter
    if (Character.isLetter(c))
    {
        charBuffer[iCharCount] = c;
        iCharCount++;
    }
    // Check if the character is a space or an end-of-line marker
    else if ((c == ' ') || (c == '\n')) && (!strTextFile.endsWith(" "))
    {
        charBuffer[iCharCount] = ' ';
        iCharCount++;
    }
}
```

# Performance Profiling, V2

---

*(Dracula vs. Lewis & Clark)*

Step	Function	Running Time
Create vectors:	Read each file	1.09s
	Parse each file	3.68s
	Sort words in each file	332.13s
	Count word frequencies	0.30s
Dot product:		6.06s
Norm:		3.80s
Angle:		6.06s
<b>Total:</b>		<b>11minutes <math>\approx</math> 680.49s</b>

# Goals for the Semester

---

## Algorithms:

- Design of efficient algorithms
- Analysis of algorithms

## Implementation:

- Solve real problems
- Analyze and profile performance
- Improve performance via better algorithms

# Document Distance

---

*(Dracula vs. Lewis & Clark)*

Version	Change	Running Time
Version 1		4,311.00s
Version 2	Better file handling	676.50s
Version 3	Faster sorting	6.59s
Version 4	No sorting!	2.35s

# For next time...

---

## Friday lecture:

- Object-oriented programming
- Doc. Distance V3: Sorting

## Friday tutorial:

- Details of Document Distance implementation

## Discussion Groups:

- None this week. Sign up in CORS.

## Problem Set 1:

- Released. Due next week.



# Administrative Details

---

## Registration:

1. If you are not currently registered (via CORS), send me an e-mail.
2. Go to [cs2020.ddns.nus.edu.sg](https://cs2020.ddns.nus.edu.sg) and register.
3. Register for "Tutorial" session on CORS.
4. Register for "Lab" (Discussion Group) on CORS.
5. Fill out Discussion Group Preference form.