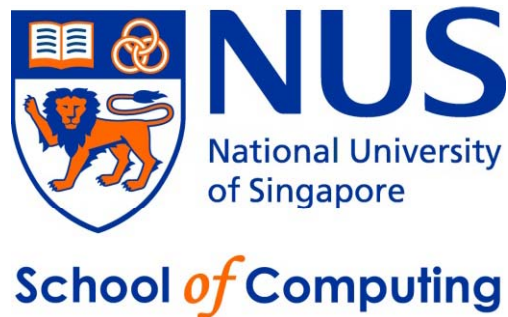


# CS2020 – Data Structures and Algorithms Accelerated

## Lecture 23 – Finale

[stevenhalim@gmail.com](mailto:stevenhalim@gmail.com)




# Outline

- What are we going to learn in this lecture?
  - Review of the entire CS2020 (the last time we use clickers)
    - And introducing the world beyond polynomial algorithms
  - Admin (15-20 minutes): please return your clicker!
  - Future
    - Modules in SoC beyond CS2020
    - Part time TA jobs
    - Final Exam Tips

# Do you bring your clicker?

0 marks for your CS2020 final exam if you do not/cannot click yes ☹️

1. Yes
2. No (do NOT click me)




# Important

- For those who click “yes” (or “no” 😊)
  - You will need to return your clicker today
  - Sign (again) the loan form that you have signed at the start of this semester
  - We will do this **later** during lecture break (10.50-11.10am)
- For those who cannot click “yes”
  - You will not get 0 marks for final if you:
    - Meet Steven or Seth during reading week and return your clicker (if you just forgot to bring it today)
    - Pay 89 SGD replacement fee (we will talk with SoC UG office for some compensation) if you really lost / broke it ...

# PS10

1. I have submitted my codes few days ago 😊
2. I have JUST submitted my code 1 hour++ ago, I haven't sleep >.<
3. I need more time, is it possible to extend the deadline again?



# Semester Review: Topic 1

- Linear Data Structures
  - Arrays/Vectors
  - Searching
    - $O(n)$  linear search,  $O(\log n)$  binary search
  - Sorting
    - $O(n^2)$  sort,  $O(n \log n)$  sort,  
 $\Omega(n \log n)$  lower bound of sorting,  $O(n)$  special case sort
  - Lists/Linked Lists/Skip List
  - Stacks, Queues
- Divide and Conquer Paradigm
  - Notable Example: Document Distance Problem
  - In red: not covered in CS1020

What is the best sorting algorithm to sort this particular sequence?

$X = \{ 1, 1, 1, 2, 1, 3, 3, 4, 5, 6, 7, 10, 1M \}$

1. Bubble Sort



2. Insertion Sort

3. Quick Sort


4. Merge Sort

5. Heap Sort

6. Counting Sort

7. Radix Sort

8. Others: \_\_\_\_\_



$$g(n) = 5000 \log n * \log n + 5n - n + \text{sqrt}(n^2)$$

$$g(n) =$$

😊 1.  $O(\text{sqrt}(n^2))$

2.  $O(n^2)$

3.  $O(n \log n)$


😊 4.  $O(n)$

5.  $O(n \log \log n)$

6.  $O(\log n)$

7.  $O(\log \log n)$


8.  $O(1)$






# What is the worst possible input for the merge sort algorithm?

1. A nearly sorted array
2. Reverse sorted array
3. A very unsorted array
4. Array where all elements are the same
5. 1 & 2 only
6. 1 & 4 only
7. All of the above



# What is the time complexity of Strassen's Algorithm for Matrix Multiplication?

1.  $O(n^3)$
2.  $O(n^{2.81})$
3.  $O(n^{2.376})$
4.  $O(n^2)$
5.  $O(n \log n)$
6.  $O(n)$
7. Eh, did we ever discuss this algorithm?



# Semester Review: Topic 2

- Trees
  - Binary Search Trees
  - Balanced BSTs
  - Heaps and Priority Queues
- Notable Example: Air Traffic Controller Problem
- Green: I will likely cover them in CS2010

What is the right child of the root of BST if the following sequence of items are inserted to an initially empty BST: {8, 7, 2, 10, 4}


1. 2

2. 7

3. 8


😊 4. 10

5. 4




# Which Statement About Heap is False?

1. Heap must be a complete binary tree
2. Heap can be implemented with array
3. Heap can be used for sorting
4. Building a heap from an unsorted array can only be done in  $O(n \log n)$
5. Heap can be used in Prim's algorithm



Which one is a **valid** partition of an array  
 $X = \{8, 2, 7, 9, 10, 1, 5\}$  around value 7?  
(you can select up to 5 choices)

- 😊 1.  $X' = \{1, 2, 5, 7, 9, 10, 8\}$
- 😊 2.  $X' = \{1, 2, 5, 7, 8, 9, 10\}$
- 😊 3.  $X' = \{5, 2, 1, 7, 8, 9, 10\}$
- 4.  $X' = \{5, 2, 1, 8, 7, 9, 10\}$
- 5.  $X' = \{2, 1, 5, 8, 7, 9, 10\}$



# Semester Review: Topic 3

- Hash Tables
  - Dictionaries
  - Hash functions
  - Chaining
  - Amortized Analysis
- Notable Example: DNA Analysis


Give me a pair of numbers  $x_1$  and  $x_2$  so that  
 $h(x_1) = h(x_2)$  for  $h(x) = (x * x) \% 7$   
(you can select up to 4 options)

😊 1.  $x_1 = 71, x_2 = 55$

2.  $x_1 = 77, x_2 = 66$

😊 3.  $x_1 = 7, x_2 = 14$


4.  $x_2 = 9, x_3 = 147$





# In which scenario that hash table is a worse choice than BST?

1. When  $n$  is big, there will be lots of collision, so hash table can be slower than BST
2. When  $n$  is very small, it is faster to use BST
3. When you want to list down the items in sorted order




**2<sup>ND</sup> HALF**

# Semester Review: Topic 4

- Graphs
  - Graph Data Structures: AdjMat/List/EdgeList/Implicit Graph, etc
  - Graph Traversal: DFS/BFS
  - Minimum Spanning Trees: Kruskal's/Prim's
  - Single Source Shortest Paths: Bellman Ford's/Dijkstra's
- Notable Examples: McDonalds, Parity Networks


# Select graph terminologies that you know **now...** (can select up to 8/clicker)

1. AdjMatrix/List/EdgeList
2. DFS/BFS
3. Topological Sort
4. MST/Prim's
5. MST/Kruskal's
6. SSSP/Bellman Ford's
7. SSSP/Dijkstra's
8. APSP/Floyd Warshall's




Select DS/algorithms that you have already **implement now...** (can select up to 8/clicker)

1. AdjMatrix/List/EdgeList
2. DFS/BFS
3. Topological Sort
4. MST/Prim's
5. MST/Kruskal's
6. SSSP/Bellman Ford's
7. SSSP/Dijkstra's
8. APSP/Floyd Warshall's





# DFS/BFS can run in $O(V^2)$ on

1. Tree
2. Directed Acyclic Graph
3. Bipartite Graph
4. General Graph
- 😊 5. Complete Graph
6. Impossible, it is  $O(V+E)$



# Hm... Can I use BFS to find toposort?

1. Yes, why not?
2. No, I think BFS will have a problem because \_\_\_\_\_



# Semester Review: Topic 5

- Dynamic Programming (mainly on Graph)
  - Algorithms on DAG
  - Algorithms on Tree
  - Algorithms on General Graph
  - All Pairs Shortest Paths: Floyd Warshall's
  - DP, the true form
- Notable Example: Traveling Salesman Problem (TSP)




# Larger TSP Instances

N	N!	$N^2 * 2^N$	Improvement
9	362,880	$81 * 512 = 41,472$	8.75 x
10	3,628,800	$100 * 1024 = 102,400$	35.44 x
11	39,916,800	$121 * 2048 = 247,808$	161.07 x
12	479,001,600	$144 * 4096 = 589,824$	812.11 x
13	6,227,020,800	$169 * 8192 = 1,384,448$	4,497.84 x
14	87,178,291,200	$196 * 16384 = 3,211,264$	27,147.66 x
15	1,307,674,368,000	$225 * 32,768 = 7,372,800$	177,364.69 x
16	20,922,789,888,000	$256 * 65,536 = 16,777,216$	1,247,095.46 x
17	355,687,428,096,000	$289 * 131,072 = 37,879,808$	9,389,895.22 x
51	Hm...	5,856,931,315,395,330,048 ☹	...
105?		4.47227131760519332848036e+35 ☹	
		TSP is NP-Complete	

# How to solve this?

- <http://www.tsp.gatech.edu/world/images/anim1a.html>



# Research on Optimization Problems

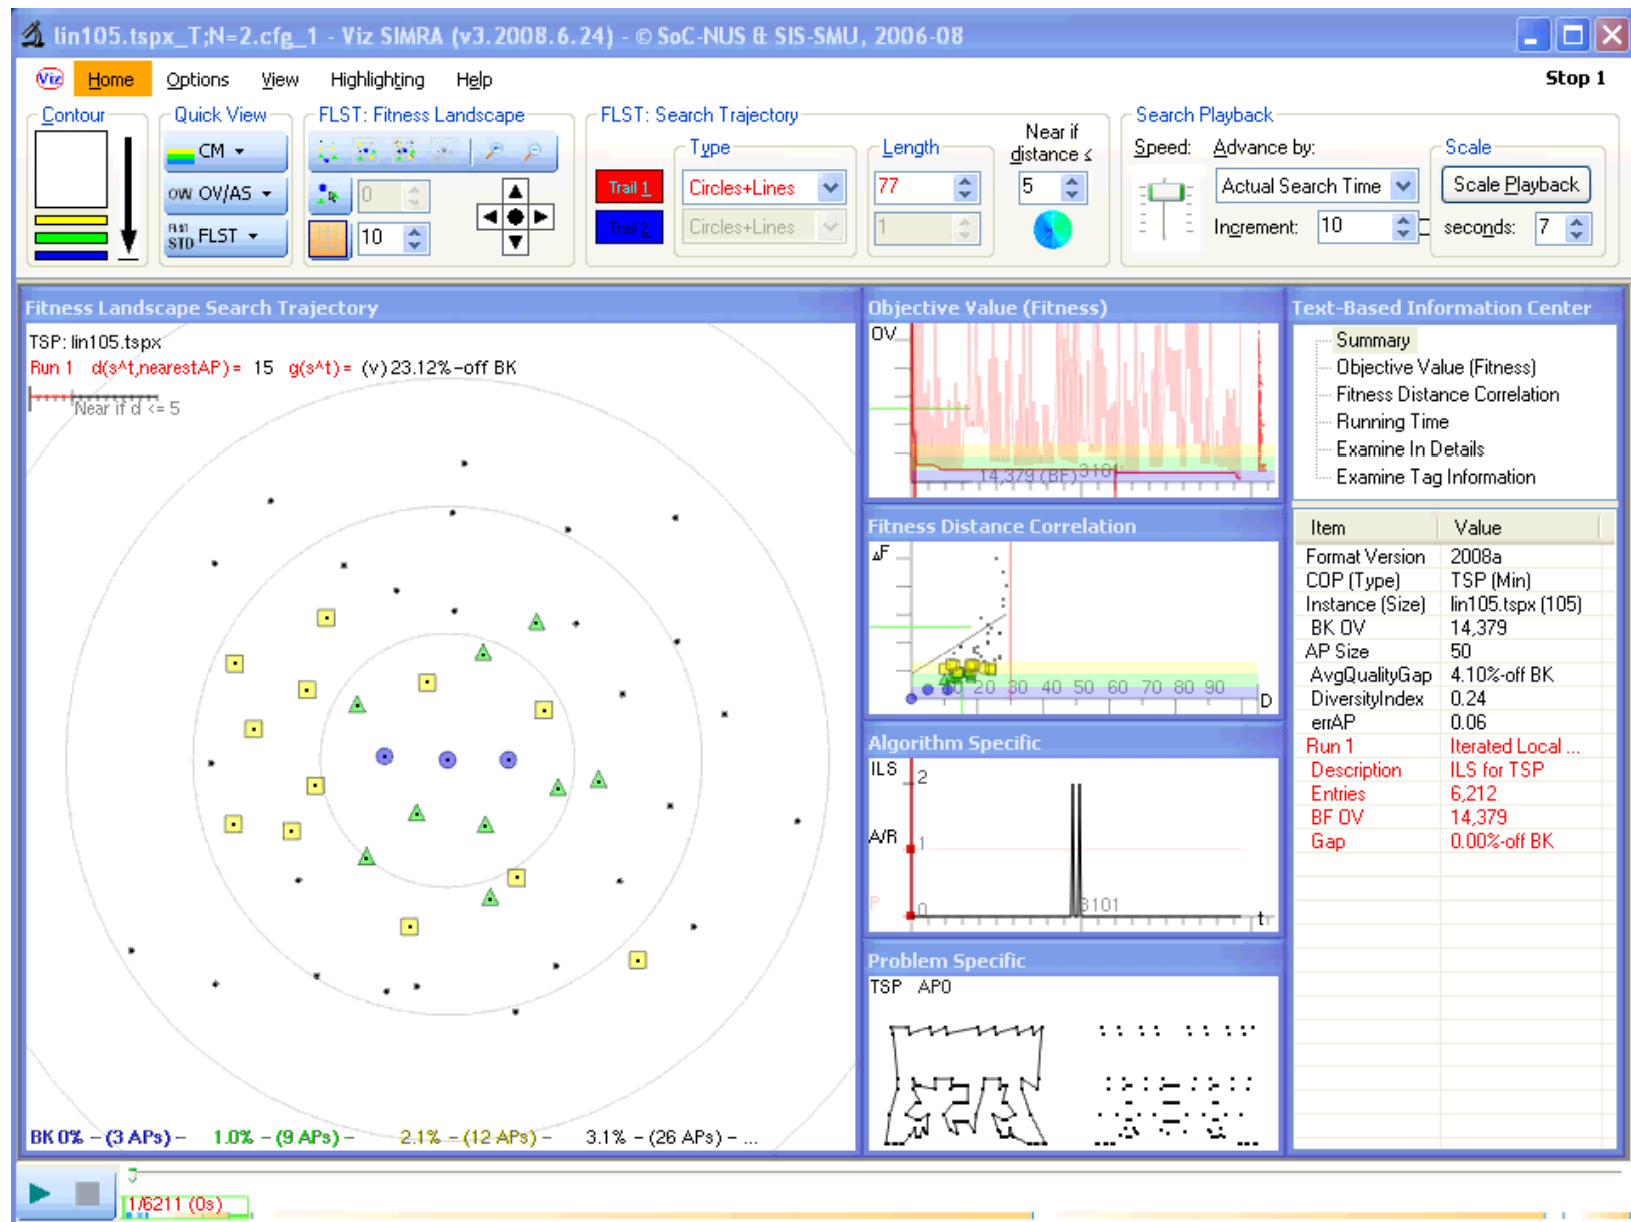
- What if we can accept **near optimal** solution?
- Other options are now available:
  - Approximate/Heuristics Solution
  - Local Search

# Euclidean TSP Approximation with MST

- <http://www.personal.kent.edu/~rmuhamma/Algorithms/MyAlgorithms/AproxAlgor/TSP/tsp.htm>
- We can get MST of a complete graph with  $N$  vertices ( $N^2$  edges) in  $O(N^2 \log N)$
- We can then get a TSP tour that is at worst twice the length of the optimal tour...


# Local Search Solution

## From Steven's PhD thesis (2004-2010)




# Local Search Solution

- Start from any permutation that cycles back (a tour)
- Use (2-Opt) Swap-Edge Heuristics
- Locally optimize it
- Iteratively repeat this process with other random tours



Solution: (0, 3, 4, 5, 9, 8, 2, 1, 6, 7, 0)

AFTER TWO-SWAP((4,5),(8,2))



# Some Demo 😊

- <http://comopt.ifl.uni-heidelberg.de/software/TSPLIB95/STSP.html>

# **SUMMARY OF 2<sup>ND</sup> HALF**



# Graph DS

- L13: AdjMatrix, AdjList
- RecitationWk8: Implicit Graph
- L15: EdgeList
- L20: special graph that has 2 edges/vertex
- Space Complexity
- Benefits/Disadvantages of using them

# MST

- L15: Kruskal's Algorithm
- L15: Union Find Data Structure
- L15: Prim's Algorithm (brief)
- Uses L6: Heap Data Structure
- Model problem as an MST problem

# SSSP, Special Cases

- L16: General Shortest Paths Algorithm:  $O(\infty)$
- L16: Bellman Ford's:  $O(VE)$  – general
- RecWk9: SSSP (Graph) Problem Modeling
- L17: Dijkstra's:  $O((V + E) \log V)$  – binary heap
- L17: BFS:  $O(V + E)$ , unweighted graph
- L17: DFS/BFS:  $O(V)$ , Tree
- L17: Toposort + DP:  $O(V + E)$ , DAG
- RecWk10/Quiz2: Graph Problem Modeling

# DP (have we test you on this?)

- L18: Algorithms on DAG
- L20: Algorithms on General Graph → implicit DAG
- RecWk11: Algorithms on Tree → implicit DAG
- L19: APSP: Floyd Warshall's DP algorithm
- L21: DP, true form
- Space Complexity: Vertices in implicit DAG/states
- Time Complexity: Edges in implicit DAG/transitions

# Traveling Salesman Problem

- L20: Recursive Backtracking Algorithm:  $O(N!)$
- L20: Dynamic Programming Algorithm:  $O(N^2 \times 2^N)$
- PS10: TOUR
- L23: Approximation Algorithm/MST:  $O(N^2 \log N)$ 
  - Only for Euclidean MST
  - The results are not guaranteed to be optimal
- L23: Local Search / Heuristics:  $O(\text{Iterations} * C)$ 
  - The results are not guaranteed to be optimal

# Things That Can Be Done Better

- ii.java → IntegerPair.java
- iii.java → IntegerTriple.java
- L6: Heap\_DecreaseKey() → heapDecreaseKey()
- In several demos: Reduce the usage of global variables
- More Object Oriented examples
- Less typos in our lecture notes/Pses
- More balanced DGs/PSes
- Etc (we will look at your NUS Online Teaching Feedback regarding this module 😊)

# Admin

- Now, let's return your clicker
- The procedure:
  - I have sorted the loan forms based on clicker ID
  - I will call your name one by one
  - Come forward, return your clicker, and sign again with red pen in your own loan form to indicate that you have returned the clicker
  - Then you can have 10-15 minutes break  
(until this process is over, we need about 54 iterations 😊)

**SOME FINAL SLIDES**



Have you count your own marks so far? (15% PS, 5% DG, 15% Quiz1, 10% Coding Quiz (please just use your own estimation), 15% Quiz2)

1. I have scored  
[55-60]%
2. [50-55]%
3. [45-50]%
4. [40-45]%
5. I haven't count
6. I have count but don't  
want to tell

# Relevant Modules After CS2020 (1)

- CS3233: Competitive Programming
  - Pre-req: **at least** A- from CS2020, **(very)** strong interest in algorithms, programming (especially implementation), and problem solving
  - Must free-up your schedule every Wednesday 6-9pm on the 2<sup>nd</sup> semester...
  - CS3233 syllabus will be altered **a bit** because of CS2020/10
- CS3230: Design and Analysis of Algorithms
  - Discussion: take soon (next semester) or next next AY (after “distance propagation” reaches the **new** CS3230)
- CS5206: Foundations in Algorithms

# CS3233: Competitive Programming

1. I will **definitely** take your challenging module next year (S2, AY2011-2012)
2. I will consider 😊, maybe I will take it...
3. I will consider, but at the moment I do not have plan to take it...
4. I know for sure that this module is not for me, thanks

# CS3230: Design and Analysis of Algorithms

1. I will definitely take this module next sem (S1 AY2011-2012) to increase my chance for A before the syllabus gets harder
2. I will strategically take other modules first and will take the **updated** and proper CS3230 only next next AY
3. I do not know what to do at the moment...

# Relevant Modules After CS2020 (2)

- CS3210: Parallel Computing
- CS3211: Parallel and Concurrent Programming
- CS4231: Parallel and Distributed Algorithms
  - If you like Dr Seth's mystery lecture last Tuesday
- CS2103: Software Engineering
- CS3215: Software Engineering Project
  - Larger scale software creation modules
- CS3243: Artificial Intelligence
- Enjoy 😊

# Final Exam (1)

- Not a fully open book test
- Your only “data structure”: 2 sheets/4 sides A4 paper
- Either handwritten or printed
  - Use whatever font size convenient for you
  - NO magnifying glass is allowed in the exam hall 😊
- Purpose:
  - To improve your search performance during exam time from  $O(N)$  where  $N$  is the number of books/lecture notes that you are planning to bring into exam hall into  $O(1)$  where you can use a kind of “hashing” to identify the required information in your 4 sides A4 paper 😊

# Final Exam (2)

- There will be 4 questions
  - Not of the same length
  - Not of the same marks
  - More similar to Quiz 2 style rather than Quiz 1
    - No MCQ
  - Topics: from the entire semester with focus on things that are not yet tested in Quiz 1, Coding Test, & Quiz 2 :O
    - Your job is to see those quizzes and test and see which part in the syllabus that we have not test yet 😊

# Final Exam (3)

- Time Management
  - You have 20 more minutes than quiz 1 or 2
  - We assume that you have 4 sides A4 paper with  $O(1)$  performance 😊
  - We estimate that you will likely spend more time thinking than writing answers
- Study Preparation
  - In case you are not aware, few years ago I have compiled solutions for a lot of past exam papers of the (old) CS1102
    - [http://www.comp.nus.edu.sg/~stevenha/myteaching/exam\\_hints.pdf](http://www.comp.nus.edu.sg/~stevenha/myteaching/exam_hints.pdf)



# Consultation Slots @ Study Week

- Dr Seth Gilbert, COM2-3-23
  - Preferred time slot: Wednesday 20 Afternoon
  - Not available on study week: Monday 18/Tuesday 19
  - Email first to book a slot: [seth.gilbert@comp.nus.edu.sg](mailto:seth.gilbert@comp.nus.edu.sg)
- Dr Steven Halim, COM2-3-37
  - Really just opposite of Seth's office
  - Preferred time slot: Tuesday 19/Thursday 21 Afternoon
  - Not available on study week: Monday 18
  - Email first to book a slot: [stevenhalim@gmail.com](mailto:stevenhalim@gmail.com)
- Note: next Friday is Good Friday

# Part-time Teaching Assistant Job (1)

- FYI, I will teach CS2010 next semester
  - S1, AY2011/2012 (Aug-Nov 2011)
  - Estimated enrollment (if everyone pass CS1020): ~228 :O
- I estimate that if I want 1:20 staff:students ratio
  - Then I need to hire **at least**  $\text{ceil}(228/20.0) = 12$  different TAs
  - Note: This is worse than CS2020 or even CS1101S ratio ☹...
- I prefer TAs who:
  - Like algorithms and problem solving
  - Eventually score at least A- in CS2020
    - You will know this by end of May 2011
  - Like to talk with students (to be precise: your friends)
    - FYI: these CS2010 students will be from YOUR BATCH

# Part-time Teaching Assistant Job (2)

- If you are interested, drop me an email as soon as you know your CS2020 grade
  - Or if you are sure that you will get at least A-, then you can email me earlier
  - We will do some interviews, maybe during July 2011
- FYI, Dr Seth will teach CS2020 next year
  - S2, AY2011/2012 (Jan-Apr 2012)
  - Estimated enrollment next year: ~40-50
  - So if you prefer to TA CS2020 next year, drop Seth an email

# Thank You 😊

- For being a good pioneer batch of CS2020
- We need your constructive feedback to further improve this module next year
- Please write those comments in your NUS Online Teaching Feedback exercise 😊