

CS2020

# Data Structures and Algorithms (Recitation)

**Welcome!**

# Today

---

## 2-3-4 Trees:

- New type of balanced search tree

## Cache Aware Algorithms

- Modeling cache performance
- B-trees

# Balanced Search Trees


---

Many types:

- AVL trees
- Red-Black trees
- Splay trees
- ...

All use rotations

- Complicated and unintuitive
- Concurrency is hard
- Messy: lots of case analysis




# 2-3-4 Trees


---

## Rule # 1:


- Every non-leaf node has either:  
2 or 3 or 4 children



Binary Tree:  
2 children




3 children



4 children

# 2-3-4 Trees




# 2-3-4 Trees


---

## Rule # 2:


- Satisfies binary search property



Binary Tree:  
2 children  
1 key



3 children  
2 keys



4 children  
3 keys

# 2-3-4 Trees




---

## Rule # 2:


- Satisfies binary search property

Node contains keys  $\{k_1, \dots, k_r\}$ ,  $r \in \{1, 2, 3\}$

Node contains subtrees  $\{T_1, \dots, T_{r+1}\}$ ,  $r \in \{1, 2, 3\}$




# 2-3-4 Trees





# 2-3-4 Trees



# 2-3-4 Trees

---


## Rule # 2:

- Satisfies search property

Node contains keys  $\{k_1, \dots, k_r\}$ ,  $r \in \{1, 2, 3\}$

Node contains subtrees  $\{T_1, \dots, T_{r+1}\}$ ,  $r \in \{1, 2, 3\}$

$\text{keys}(T_1) < k_1 < \text{keys}(T_2) < k_1 < \text{keys}(T_3) < k_3 < \text{keys}(T_4)$



# 2-3-4 Trees

---


## Rule # 2:

- Satisfies search property


Node contains keys  $\{k_1, \dots, k_r\}$ ,  $r \in \{1, 2, 3\}$

Node contains subtrees  $\{T_1, \dots, T_{r+1}\}$ ,  $r \in \{1, 2, 3\}$

$$\forall r \in \{1, 2, 3\} : keys(T_r) < k_r < keys(T_{r+1})$$




# 2-3-4 Trees




# 2-3-4 Trees: Searching

search(82)




# 2-3-4 Trees: Searching

search(82)



# 2-3-4 Trees: Searching

search(82)



# 2-3-4 Trees

---

search(k):


$j = 1;$

**while** ( $j \leq \text{num\_keys}$ ) **and** ( $k > k_j$ ) **do**  $j++;$

**if** ( $j \leq \text{num\_keys}$ ) **and** ( $k == k_j$ ) **then return** true;

**else if** ( $T_j \neq \text{null}$ ) **then return**  $T_j.\text{search}(k);$


**else return** false;





# 2-3-4 Trees: in-order-traversal

in-order-traversal()




traverse T1:

15, 23, 30, 32, 38

# 2-3-4 Trees: in-order-traversal

in-order-traversal()




print k1:

15, 23, 30, 32, 38, 42

# 2-3-4 Trees: in-order-traversal

in-order-traversal()




traverse T2:

15, 23, 30, 32, 38, 42, 48, 60, 65, 70

# 2-3-4 Trees: in-order-traversal

in-order-traversal()




print k2:

15, 23, 30, 32, 38, 42, 48, 60, 65, 70, 73

# 2-3-4 Trees: in-order-traversal

in-order-traversal()



traverse T3:


15, 23, 30, 32, 38, 42, 48, 60, 65, 70, 73, 77, 80,

# 2-3-4 Trees

---

Rule #3: Every leaf has the same depth.

- Every path from root->leaf is the same length.



# 2-3-4 Trees

---

**Claim:** a B-tree with  $n$  keys has height  $O(\log n)$

# 2-3-4 Trees

---

**Claim:** a B-tree with  $n$  keys has height  $O(\log n)$

Intuition:

- Every node has at least two children.
- When the height increases by 1, the number of keys doubles.
- After  $\log(n)$  levels, there are at least  $n$  keys.



# 2-3-4 Trees

---

**Claim:**  $h < c \log n$

Proof:

depth = 0: 1

depth = 1: 2


depth = 2: 4

depth = 3: 8

...

depth =  $h$ :  $2^h$

$$n \geq 1 + 2 + 4 + 8 + \dots + 2^h = \sum_{j=0}^h 2^j$$



# Math Intermission

---

Prove:  $\sum_{j=1}^h 2^j = 2^{h+1} - 1$

# Math Intermission

---

Prove:  $\sum_{j=0}^h 2^j = 2^{h+1} - 1$

Induction:

– Base case:  $h=0$

$$\sum_{j=0}^0 2^j = 2^0 = 1 = 2^1 - 1 = 2^{h+1} - 1$$

# Math Intermission

---

Prove:  $\sum_{j=0}^h 2^j = 2^{h+1} - 1$

Inductive step:  $\sum_{j=0}^h 2^j = 2^h + \sum_{j=0}^{h-1} 2^j$

$$= 2^h + (2^h - 1)$$
$$= 2^{h+1} - 1$$

# 2-3-4 Trees

---

**Claim:**  $h = O(\log n)$

Proof:

depth = 0: 1

depth = 1: 2


depth = 2: 4

depth = 3: 8

...

depth =  $h$ :  $2^h$


$$n \geq 1 + 2 + 4 + 8 + \dots + 2^h = \sum_{j=0}^h 2^j = 2^{h+1} - 1$$



# 2-3-4 Trees


---

**search:**  $O(\log n)$




# 2-3-4 Trees: Inserting

insert(70)



# 2-3-4 Trees: Inserting


insert(70)






# 2-3-4 Trees: Inserting

insert(71)




# 2-3-4 Trees: Inserting

insert(72)




# 2-3-4 Trees: Inserting

insert(72)




# 2-3-4 Trees: Inserting

insert(72)



# 2-3-4 Trees: Inserting


insert(72)



**Too many keys!!!**

# 2-3-4 Trees: Inserting


insert(72)



**Split the node in two.**

# 2-3-4 Trees: Inserting


insert(72)



**Insert into parent.**

# 2-3-4 Trees: Inserting

insert(72)

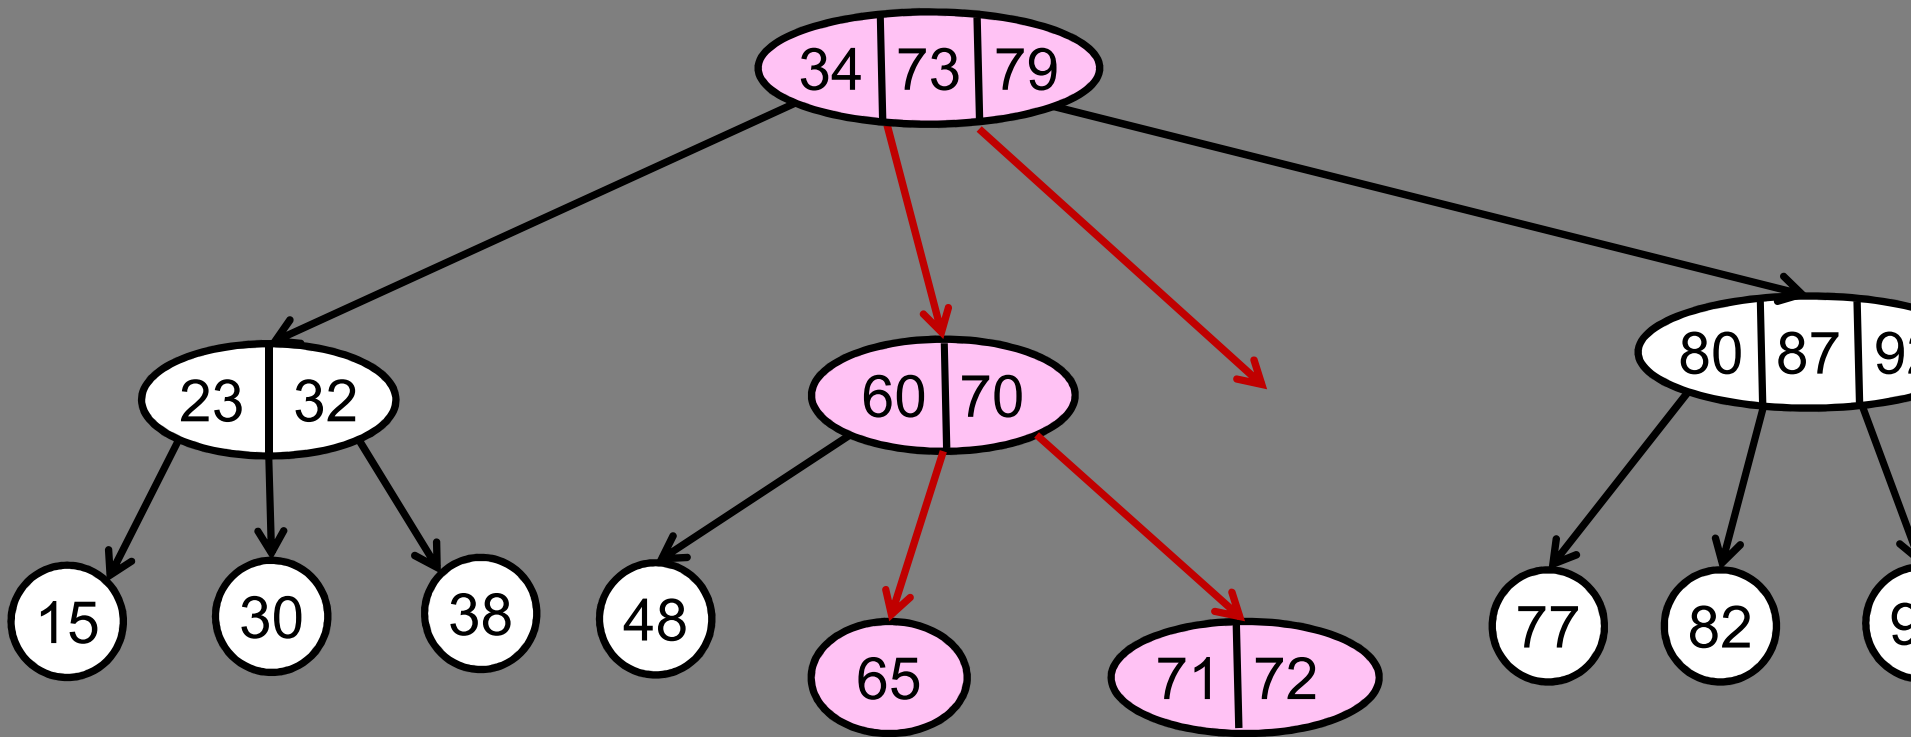


**Recurse (if parent is full).**



# 2-3-4 Trees: Inserting

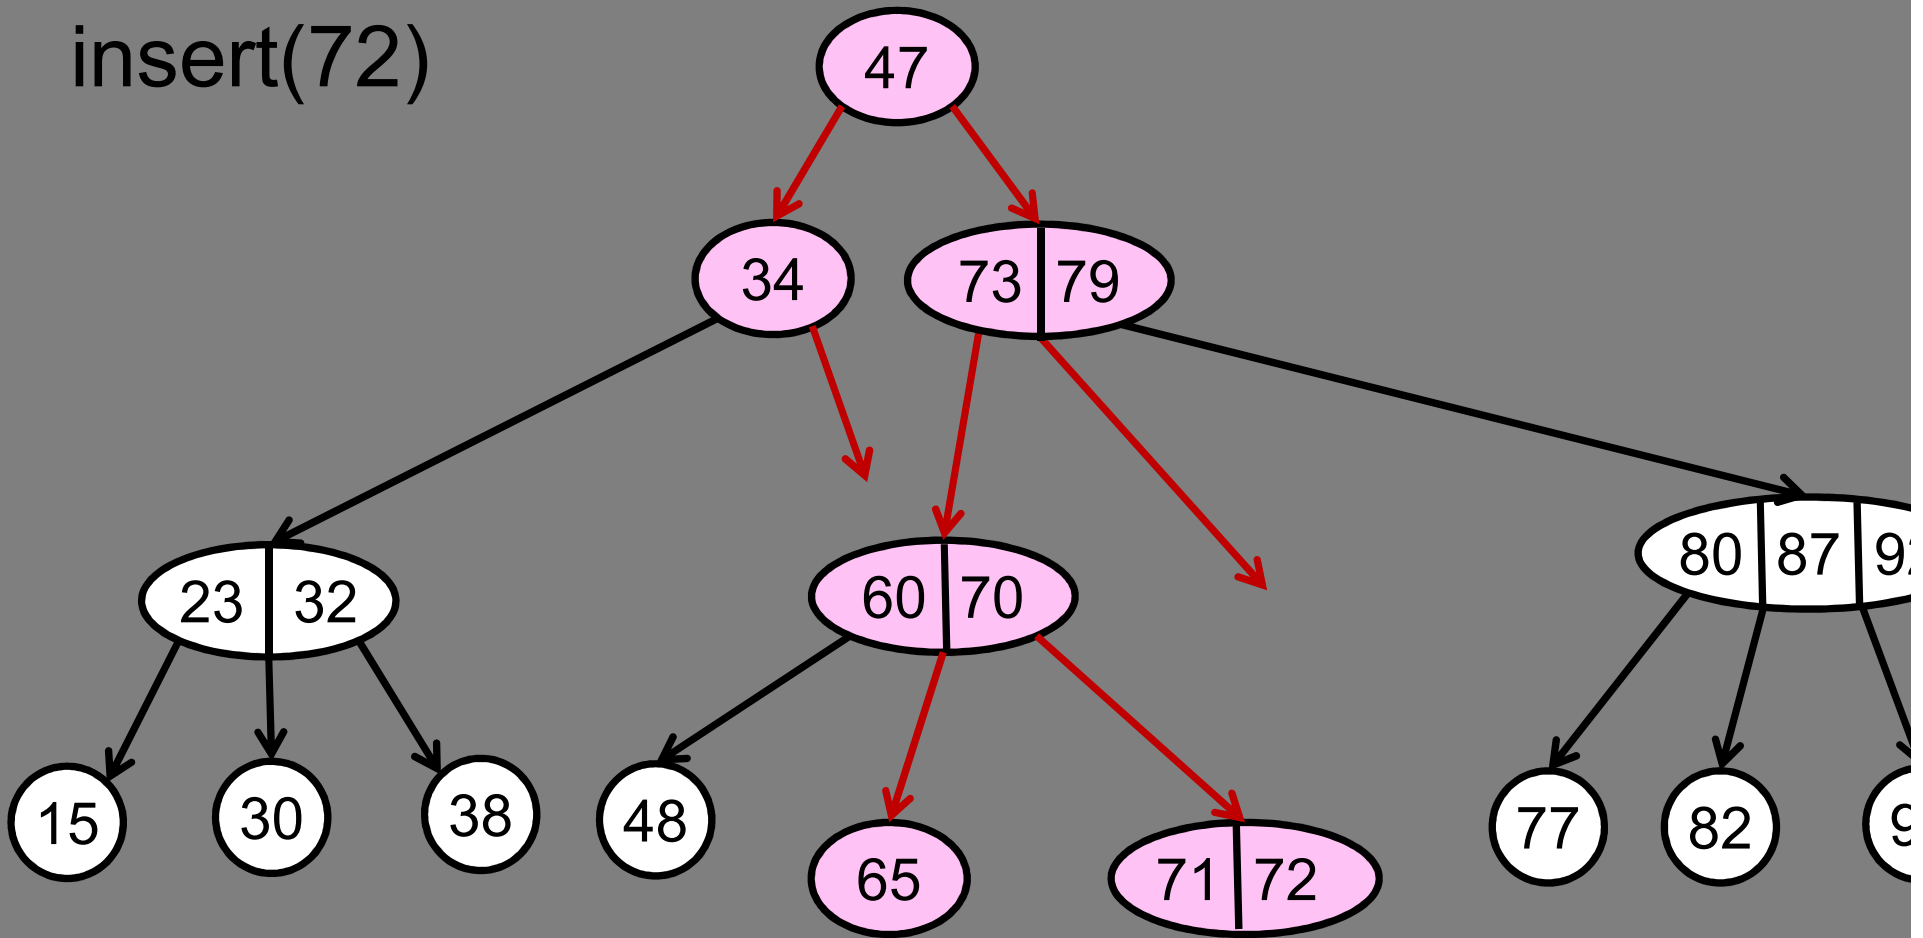
insert(72)



**What if the root is full?**

# 2-3-4 Trees: Inserting

insert(72)



**Split and create a new root.**

# 2-3-4 Trees


---

insert(key)

1. Search for leaf **x** to start insert.
2. Repeat until done:
3.     Insert **key** into **x**.
4.     If **x** overfull:
5.         **key** = median(**x**);
6.         remove **key** from **x** and split(**x**);
7.         If (**x** == root) then create new root.
8.         Else **x** = parent(**x**)
9.     Else done;


# 2-3-4 Trees: Inserting

insert(72)



# 2-3-4 Trees: Inserting

insert(72)



# 2-3-4 Trees

---

Preserves 3 Properties:

1. Every node has 2 or 3 or 4 children.
2. Search tree property.
3. All leaves have the same depth.

# 2-3-4 Trees

---

Lazy option: do splitting when needed.

Proactive option: split in advance.

One pass insertion:

- If root contains 3 keys, split root and create new root (containing one key).
- While searching for the leaf, split any node that is full (i.e., contains 3 keys).
- On arrival at leaf, there is enough space in the leaf to add the key!

# 2-3-4 Trees

---

`x.insert(key):`     *(assume x not full)*

1.    `if leaf()` then add key to array of keys.
2.    `else`
3.        `let j be the insertion point in x;`
4.        `if Tj has 3 keys, then`
5.            `split(x, j, Tj)`
6.            `if (key > kj) then j++;`
7.        `Tj.insert(key)`



# 2-3-4 Trees


---

`x.delete(key):`     (assume `x` has at least 2 keys)

1.    if `leaf()` then remove key
2.    else if (key found at `x` in key `kj`)
3.        if `T(j)` and `T(j+1)` have 1 key each:
4.            merge `T(j)` and `T(j+1)` and delete `kj`
5.        else replace `kj` with successor in `T(j+1)` or predecessor in `T(j)`
6.    else if (key in `Tj` and `Tj` has at least 2 keys)
7.        `Tj.delete(key)`
8.    else ...


# 2-3-4 Trees: Deleting (leaf)

delete(30)




# 2-3-4 Trees: Deleting (leaf)

delete(30)




# 2-3-4 Trees: Deleting (mid-tree)

delete(32)




# 2-3-4 Trees: Deleting (mid-tree)

delete(32)




# 2-3-4 Trees: Deleting (mid-tree)

delete(32)




# 2-3-4 Trees: Deleting (mid-tree)

delete(32)



# 2-3-4 Trees: Deleting (mid-tree)


delete(32)






# 2-3-4 Trees: Deleting (mid-tree)

delete(32)




# 2-3-4 Trees: Deleting (hard case)

delete(70)




# 2-3-4 Trees: Deleting (hard case)

delete(70)




# 2-3-4 Trees: Deleting (hard case)

delete(70)




# 2-3-4 Trees: Deleting (hard case)

delete(70)




# 2-3-4 Trees: Deleting (hard case 2)

delete(70)




# 2-3-4 Trees: Deleting (hard case 2)

delete(70)



# 2-3-4 Trees: Deleting (hard case 2)


delete(70)






# 2-3-4 Trees: Deleting (hard case )

delete(70)




# 2-3-4 Trees: Deleting (hard case 2)

delete(70)



# 2-3-4 Trees: Deleting (hard case 2)

delete(70)



# 2-3-4 Trees

---

## Summary:

- Three rules:
  1. Two, three, or four children
  2. Search Property
  3. All leafs have same depth
- Search: typical tree search
- Insert: split nodes to make room
- Delete: merge and juggle (see CLRS for details)

# B-trees

---

The search tree used by *all* large databases...

# B-trees

---

Where is most data stored? **Hard disk!**

- Magnetic
- Mechanical
- Slow (6000rpm = 10ms)

Two step access:

1. *seek (find right track)*
2. *read track*




Image source: Wikipedia


# B-trees

---

Two step access:

1. *seek (find right track)*
2. *read track*

Solution: Cache entire track




# B-trees

---

Solution: Cache entire track


- Read in new track: costs 1 unit of work
- Read cached data: **free**





# 2-3-4 Trees

---



## Search:


- Each step of the search reads a new track.
- Reading a track takes 10ms
- Each search takes  $10\text{ms} * O(\log n)$  time!

# B Trees

---

## **Bigger nodes:**

- Each node stores at least B keys
- No node stores more than  $2B-1$  keys
- All leaves have same depth.




Time to search node: ???

# B Trees

---

## **Bigger nodes:**

- Each node stores at least B keys
- No node stores more than  $2B-1$  keys
- All leaves have same depth.



Time to search node:  $O(B)$  reads

$O(1)$  disk seeks!

# B Trees

---


## Tree depth:

- Level 1:  $B$
- Level 2:  $B^2$
- Level 3:  $B^3$
- Level 4:  $B^4$
- ...
- Level  $h$ :  $B^h$

$$n > B^h \quad \Rightarrow \quad h < \log_B n = \log(n)/\log(B)$$

# B Trees

---




## **Bigger nodes:**

- Search, insert, delete as before.
- How to find index in big node?
  - Linear search:  $O(B)$
  - Binary search:  $O(\log B)$

# B Trees

---



## **Bigger nodes:**

- Search, insert, delete as before.
- How to find index in big node?
  - Linear search:  $O(B)$
  - Binary search:  $O(\log B)$
- But there are level 2 caches and level 3 caches...
  - Use small B-tree for smaller value of  $B$  to search for key within node!

# B-trees

---

## Summary:

- Caching performance matters.
- B-trees are fast
- B-trees are concurrent
- B-trees are (generally) simpler

## Research: (cache-oblivious algorithms)

- What if you don't know the size of your cache?
- Can you devise an algorithm that performs well for all values of  $B$ ?