

In the Lecture Series Introduction to Database Systems



# Relational Algebra



*Presented by Stéphane Bressan*

*Introduction to Database Systems*

# Relational Query Languages

- Two mathematical Query Languages form the basis for practical languages like SQL:
  - Relational Algebra: Operational, useful for representing execution plans
  - Relational Calculus: Declarative: Describe what you want, rather than how to compute it.
- Query languages are NOT programming languages:
  - Not expected to be Turing complete

# Operations (Operators)

- Operations on a single relation
  - selection  $\sigma$ , projection  $\pi$
- Usual set operations (relations are sets):
  - union  $\cup$ , intersection  $\cap$ , and difference — (non-symmetric)
- Operations combining two or more relations
  - Cartesian product  $\times$ , join  $\bowtie$  and natural join  $\bowtie_n$
- A renaming operation  $\rho$
- A division operation  $/$

## Example: employee

name	salary	eNumber
Clark	150000	1006
Gates	5000000	1005
Jones	50000	1001
Peter	45000	1002
Phillips	25000	1004
Rowe	35000	1003
Warnock	500000	1007

## Example: plane

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9

## Example: canFly

eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC9

## Example: assigned

eNumber	date	fNumber
1001	Nov 1	100
1001	Oct 31	100
1002	Nov 1	100
1002	Oct 31	100
1003	Oct 31	100
1003	Oct 31	337
1004	Oct 31	337
1005	Oct 31	337
1006	Nov 1	991
1006	Oct 31	337

## Projection

Keeps vertical slices of a relation according to a list  $L$  of attributes (i.e. *a list of columns*) of the relation  $R$ :

$$\pi_L(R) = \{t \mid \exists t_1 \\ (t_1 \in R \wedge_{A \in L} t.A = t_1.A)\}$$



## Projection (Example)

$\pi_{\text{eNumber}, \text{fNumber}}$  (assigned)

eNumber	date	fNumber
1001	Nov 1	100
1001	Oct 31	100
1002	Nov 1	100
1002	Oct 31	100
1003	Oct 31	100
1003	Oct 31	337
1004	Oct 31	337
1005	Oct 31	337
1006	Nov 1	991
1006	Oct 31	337

## Projection (Result)

$\pi_{\text{eNumber}, \text{fNumber}}$  (assigned)

eNumber	fNumber
1001	100
1002	100
1003	100
1003	337
1004	337
1005	337
1006	991
1006	337

## Projection (SQL)

```
SELECT DISTINCT eNumber, fNumber  
FROM assigned
```

## Selection

Selects the t-uples of a relation verifying a condition c:

$$\sigma_c(R) = \{t \mid t \in R \wedge c\}$$

c is any Boolean expression ( $\wedge$ ,  $\vee$ ,  $\neg$ ) involving t ( $<$ ,  $=$ ,  $>$ ,  $\neq$ ,  $\leq$ ,  $\geq$ )

---

## Selection (Example)

$\sigma_{\text{salary} < 100000}(\text{employee})$

name	salary	eNumber
Clark	150000	1006
Gates	5000000	1005
Jones	50000	1001
Peter	45000	1002
Phillips	25000	1004
Rowe	35000	1003
Warnock	500000	1007

## Selection (Result)

$\sigma_{\text{salary} < 100000}(\text{employee})$

name	salary	eNumber
Jones	50000	1001
Peter	45000	1002
Phillips	25000	1004
Rowe	35000	1003

## Selection (SQL)

```
SELECT *  
FROM employee  
WHERE salary < 100000
```

## Selection (Example)

$\sigma_{\text{salary} > 100000 \wedge \neg(\text{name} = \text{'Gates'})}(\text{employee})$

name	salary	eNumber
Clark	150000	1006
Gates	5000000	1005
Jones	50000	1001
Peter	45000	1002
Phillips	25000	1004
Rowe	35000	1003
Warnock	500000	1007



## Selection (Result)

$\sigma_{\text{salary} > 100000 \wedge \neg(\text{name} = \text{'Gates'})}(\text{employee})$

name	salary	eNumber
Clark	150000	1006
Warnock	500000	1007

## Selection (SQL)

```
SELECT *  
FROM employee  
WHERE salary > 100000  
AND name <> 'Gates'
```

## Remark: Composability

The result of a query is a relation

$\sigma_{\text{salary} < 50000}(\text{employee})$

$\pi_{\text{name, salary}}(\sigma_{\text{salary} < 50000}(\text{employee}))$

## Remark: Commutativity

$\pi_{\text{name, salary}}(\sigma_{\text{salary} < 50000}(\text{employee}))$

$\sigma_{\text{salary} < 50000}(\pi_{\text{name, salary}}(\text{employee}))$

Can we always do this?

## Remark: SQL

$\pi_{\text{name, salary}}(\sigma_{\text{salary} < 50000}(\text{employee}))$

SELECT DISTINCT name, salary  
FROM employee  
WHERE salary < 50000

Projection



Selection

## Union, Intersection, Set-difference

- $R_1 \cup R_2 = \{ t \mid t \in R_1 \vee t \in R_2 \}$
- $R_1 \cap R_2 = \{ t \mid t \in R_1 \wedge t \in R_2 \}$
- $R_1 - R_2 = \{ t \mid t \in R_1 \text{ and } \neg(t \in R_2) \}$

The relations  $R_1$  and  $R_2$  must be  
**union compatible:**

- Same number of attributes
  - Corresponding attributes have the same type (but not necessarily the same name)
-

## Union (Example)

$\text{plane}_1 \cup \text{plane}_2$

$\text{plane}_1$

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Boeing	B747
Boeing	B757

$\text{plane}_2$

maker	mNumber
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
MD	DC10
MD	DC9

## Union (Result)

$\text{plane}_1 \cup \text{plane}_2$

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9



## Union (SQL)

```
SELECT *  
FROM plane1  
UNION  
SELECT *  
FROM plane2
```

What about duplicates?

## Intersection (Example)

$\text{plane}_1 \cap \text{plane}_2$

Plane<sub>1</sub>

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Boeing	B747
Boeing	B757

Plane<sub>2</sub>

maker	mNumber
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
MD	DC10
MD	DC9

## Intersection (Result)

$\text{plane}_1 \cap \text{plane}_2$

maker	mNumber
Airbus	A330
Boeing	B747

## Intersection (SQL)

```
SELECT *  
FROM plane1  
INTERSECT  
SELECT *  
FROM plane2
```

What about duplicates?

## Difference (Example)

plane<sub>1</sub> — plane<sub>2</sub>

Plane<sub>1</sub>

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Boeing	B747
Boeing	B757

Plane<sub>2</sub>

maker	mNumber
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
MD	DC10
MD	DC9

## Difference (Result)

plane<sub>1</sub> — plane<sub>2</sub>

maker	mNumber
Airbus	A310
Airbus	A320
Boeing	B757

## Difference (SQL)

```
SELECT *  
FROM plane1  
EXCEPT  
SELECT *  
FROM plane2
```

What about duplicates?

## Cartesian Product

Combines the t-uples of two relations in all possible ways

$$R1 \times R2 = \{t \mid \exists t_1 \exists t_2 \\ (t_1 \in R_1 \wedge t_2 \in R_2 \\ \wedge_{A \in R1} t.A = t_1.A \\ \wedge_{A \in R2} t.A = t_2.A)\}$$



## Cartesian Product (Example)

canFly  $\times$  plane

eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC10

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9

# Cartesian Product (Result)

canFly  $\times$  plane

90 t-uples

eNumber	mNumber	maker	mNumber
1001	B727	Airbus	A310
1001	B727	Airbus	A320
1001	B727	Airbus	A330
1001	B727	Airbus	A340
1001	B727	Boeing	B727
1001	B727	Boeing	B747
1001	B727	Boeing	B757
1001	B727	MD	DC10
1001	B727	MD	DC9
1001	B747	Airbus	A310
1001	B747	Airbus	A320
1001	B747	Airbus	A330
1001	B747	Airbus	A340
1001	B747	Boeing	B727
1001	B747	Boeing	B747
1001	B747	Boeing	B757
1001	B747	MD	DC10
1001	B747	MD	DC9
1001	B727	Airbus	A310
1001	B727	Airbus	A320
...	...	...	...

## Cartesian Product (SQL)

```
SELECT *  
FROM canFly, plane
```

## Join ( $\theta$ -Join)

Combines the t-uples of two relations that verify a condition

$$\begin{aligned} R_1 \bowtie_c R_2 &= \{t \mid \exists t_1 \exists t_2 \\ &\quad (t_1 \in R_1 \wedge t_2 \in R_2 \wedge c \\ &\quad \wedge_{A \in R_1} t.A = t_1.A \\ &\quad \wedge_{A \in R_2} t.A = t_2.A)\} \\ &= \sigma_c (R_1 \times R_2) \end{aligned}$$

## Join ( $\theta$ -Join) (Example)

**canFly**  $\bowtie_{\text{canFly.mNnumber=plane.mNumber}}$  **plane**

eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC10

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9

## Join (Result)

canFly ⋈<sub>canFly.mNnumber=plane.mNumber</sub> plane

eNumber	mNumber	maker	mNumber
1001	B727	Boeing	B727
1001	B747	Boeing	B747
1001	DC10	MD	DC10
1002	A320	Airbus	A320
1002	A340	Airbus	A340
1002	B757	Boeing	B757
1002	DC9	MD	DC9
1003	A310	Airbus	A310
1003	DC9	MD	DC9
1003	DC10	MD	DC10

## Join (SQL)

```
SELECT *  
FROM canFly c, plane p  
WHERE c.mNumber = p.mNumber
```

## Equi-join

- Combines two relations on a condition composed only of equalities of attributes of the first and second relation
- Projects only one of the redundant attributes (since they are equal)

$$R1 \bowtie_{E(A1.1=A2.1 \wedge \dots \wedge A1.n = A2.n)} R2$$



## Equi-join (Example)

**canFly** ⋈<sub>E(canFly.mNnumber=plane.mNumber)</sub> **plane**

eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC10

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9

## Equi-join (Result)

canFly  $\bowtie_{E(\text{canFly.mNnumber}=\text{plane.mNumber})}$  plane

eNumber	mNumber	maker
1001	B727	Boeing
1001	B747	Boeing
1001	DC10	MD
1002	A320	Airbus
1002	A340	Airbus
1002	B757	Boeing
1002	DC9	MD
1003	A310	Airbus
1003	DC9	MD
1003	DC10	MD

## Natural Join

- Combines two relations on a condition composed only of equalities of attributes with the same name in the first and second relation
- Projects only one of the redundant attributes (since they are equal)

$$R_1 \bowtie_N R_2$$

## Natural Join (Example)

canFly  $\bowtie_N$  plane

eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC10

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9

## Natural Join (Result)

canFly  $\bowtie_N$  plane

eNumber	mNumber	maker
1001	B727	Boeing
1001	B747	Boeing
1001	DC10	MD
1002	A320	Airbus
1002	A340	Airbus
1002	B757	Boeing
1002	DC9	MD
1003	A310	Airbus
1003	DC9	MD

## Renaming

Renaming a relation or its attributes:

$$\rho(R'(N_1 \rightarrow N'_1, \dots, N_n \rightarrow N'_n), R)$$

The new relation  $R'$  has the same instance as  $R$ , but its schema has attribute  $N'_i$  instead of attribute  $N_i$

## Renaming (Example)

$\rho(\text{staff}(\text{salary} \rightarrow \text{wages}), \text{employee})$

employee

name	salary	eNumber
Clark	150000	1006
Gates	5000000	1005
Jones	50000	1001
Peter	45000	1002
Phillips	25000	1004
Rowe	35000	1003
Warnock	500000	1007

## Renaming (Result)

$\rho(\text{staff}(\text{salary} \rightarrow \text{wages}), \text{employee})$

staff

name	wages	eNumber
Clark	150000	1006
Gates	5000000	1005
Jones	50000	1001
Peter	45000	1002
Phillips	25000	1004
Rowe	35000	1003
Warnock	500000	1007



## Renaming (SQL)

```
SELECT name, salary AS wages, eNumber  
FROM employee
```

## Example

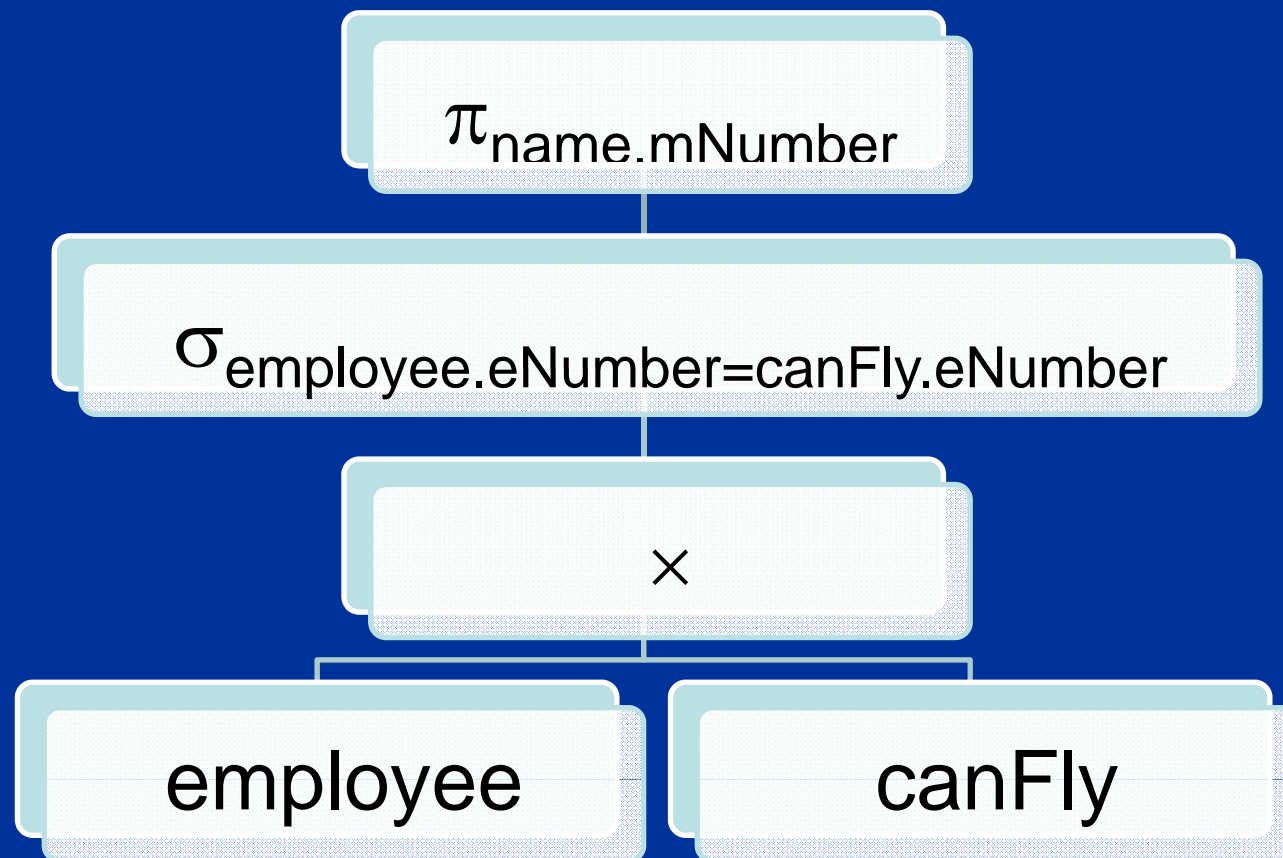
Find for each employee, her name and the model numbers of the planes she can fly

name	salary	eNumber
Clark	150000	1006
Gates	5000000	1005
Jones	50000	1001
Peter	45000	1002
Phillips	25000	1004
Rowe	35000	1003
Warnock	500000	1007

eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9

## Example

Find for each employee, her name and the model numbers of the planes she can fly



## Remark: Project-Select-Join

Project-Select-Join (PSJ) queries  
correspond to simple SQL queries:

SELECT name, mNumber ← Projection  
FROM employee, canFly ← Cartesian  
WHERE Product  
employee.eNumber=canFly.eNumber  
← Selection

## Remark: Rewriting Algebra Expressions

- $\pi_{\text{eNumber}} (\sigma_{\text{canFly.mNumber}=\text{plane.mNumber} \wedge \text{maker}=\text{'Airbus'}} (\text{canFly} \times \text{plane}))$
- $\pi_{\text{eNumber}} (\sigma_{\text{canFly.mNumber}=\text{plane.mNumber}} (\text{canFly} \times \sigma_{\text{maker}=\text{'Airbus'}}(\text{plane})))$   
Which one is more efficient?
- $\pi_{\text{eNumber}} ((\text{canFly} \bowtie_{\text{canFly.mNumber}=\text{plane.mNumber}} \sigma_{\text{maker}=\text{'Airbus'}}(\text{plane})))$
- $\pi_{\text{eNumber}} (\text{canFly} \bowtie_{\text{canFly.mNumber}=\text{plane.mNumber} \wedge \text{maker}=\text{'Airbus'}} \text{plane})$

## Example

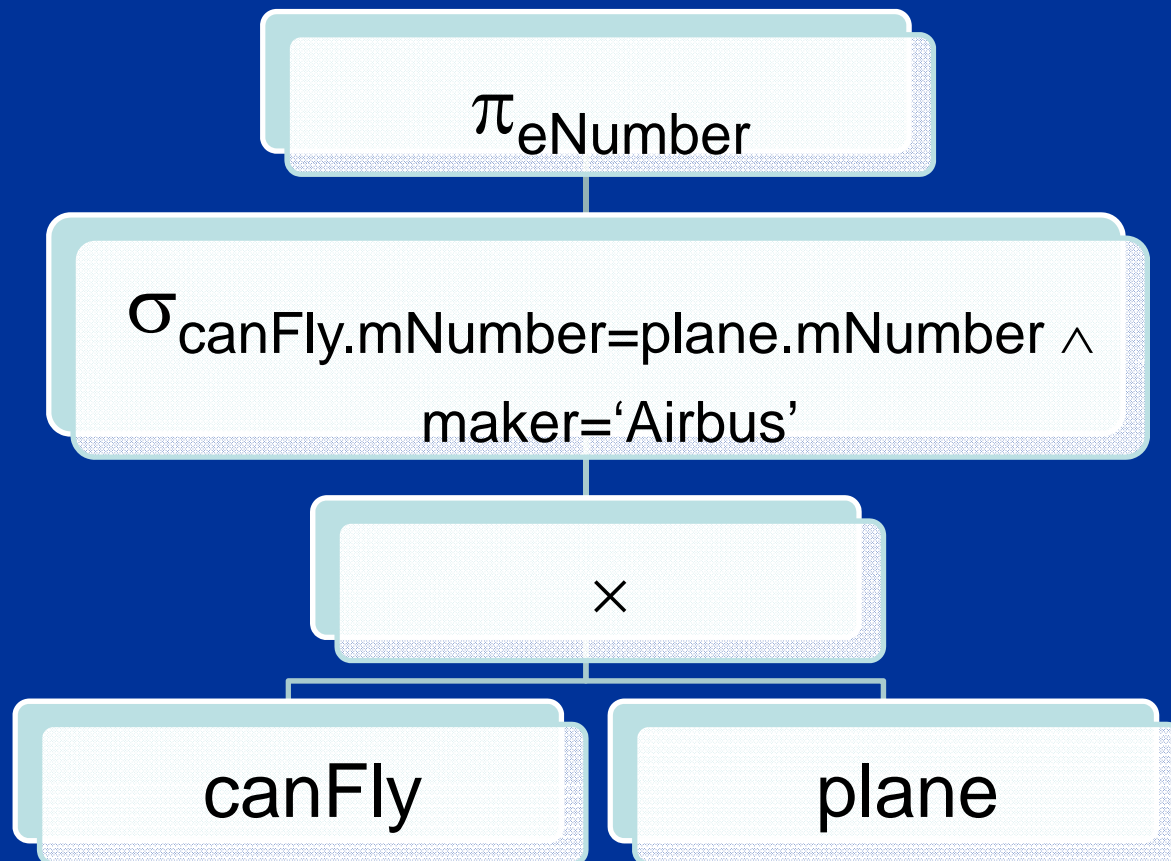
Find the employee numbers of employees who can fly Airbus planes

eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC10

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9

## Example

Find the employee numbers of employees who can fly Airbus planes



## Example

Find the employee numbers of employees who can fly Boeing or MD planes

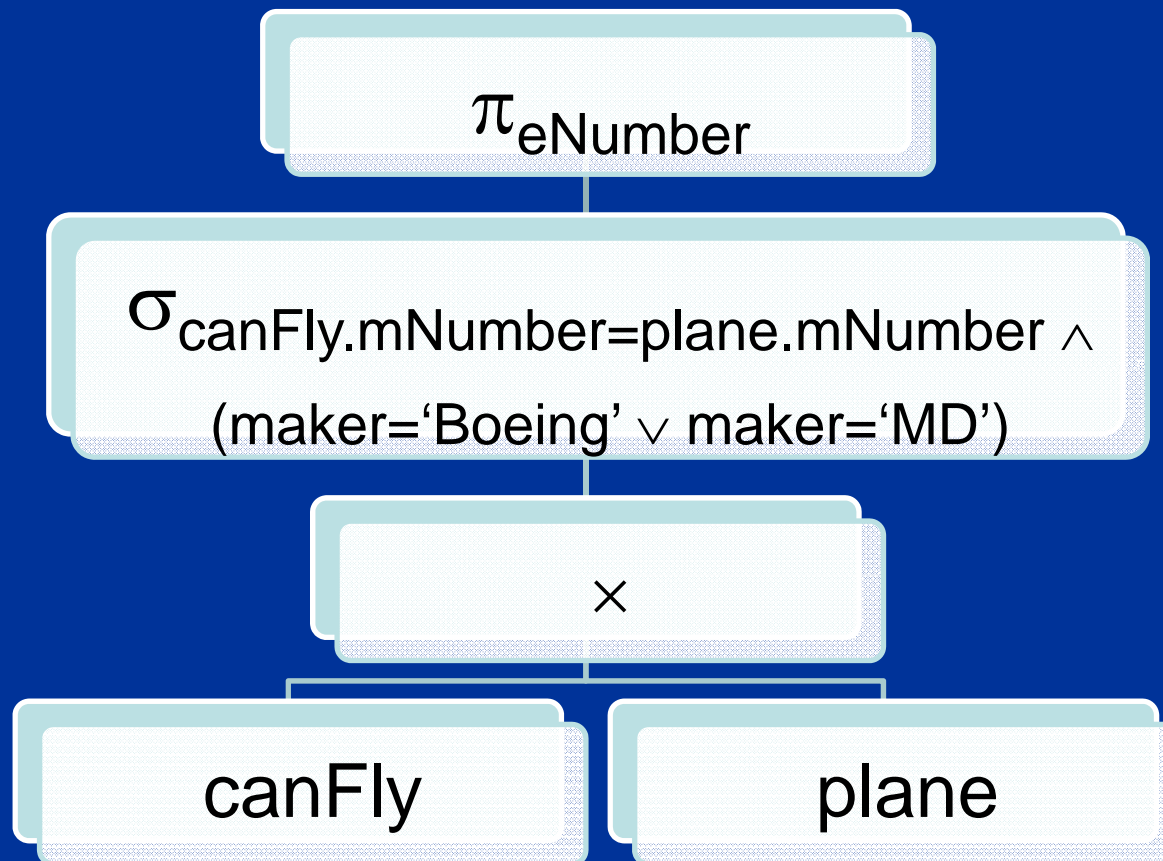
eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC10

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9




## Example

Find the employee numbers of employees who can fly Boeing **or** MD planes



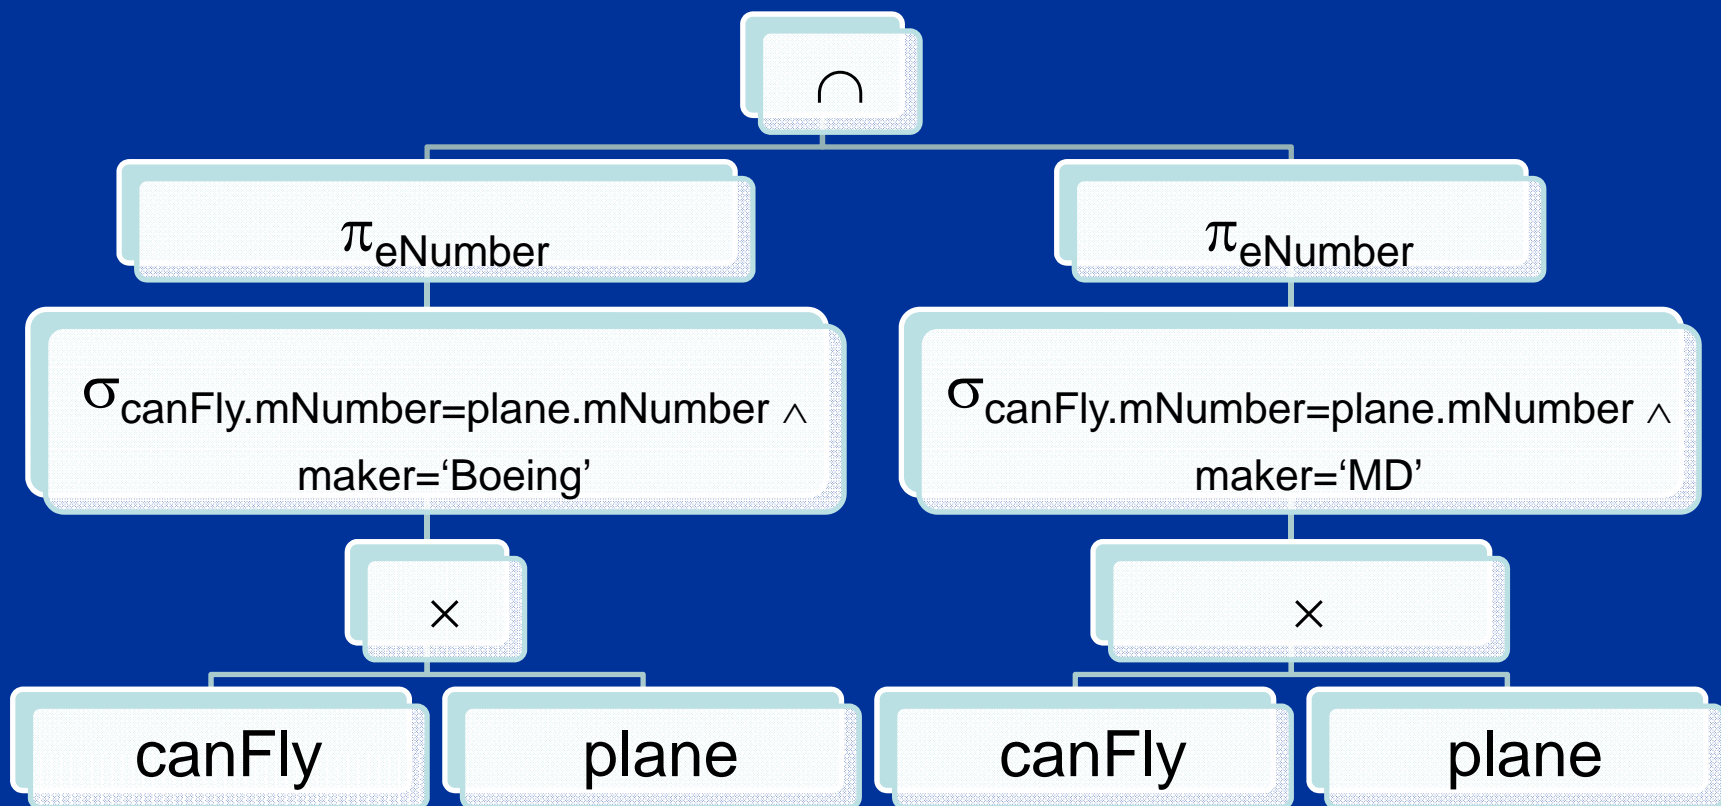
## Example

Find the employee numbers of employees who can fly Boeing **and** MD planes



## Example

Find the employee numbers of employees who can fly Boeing **and** MD planes



## Example

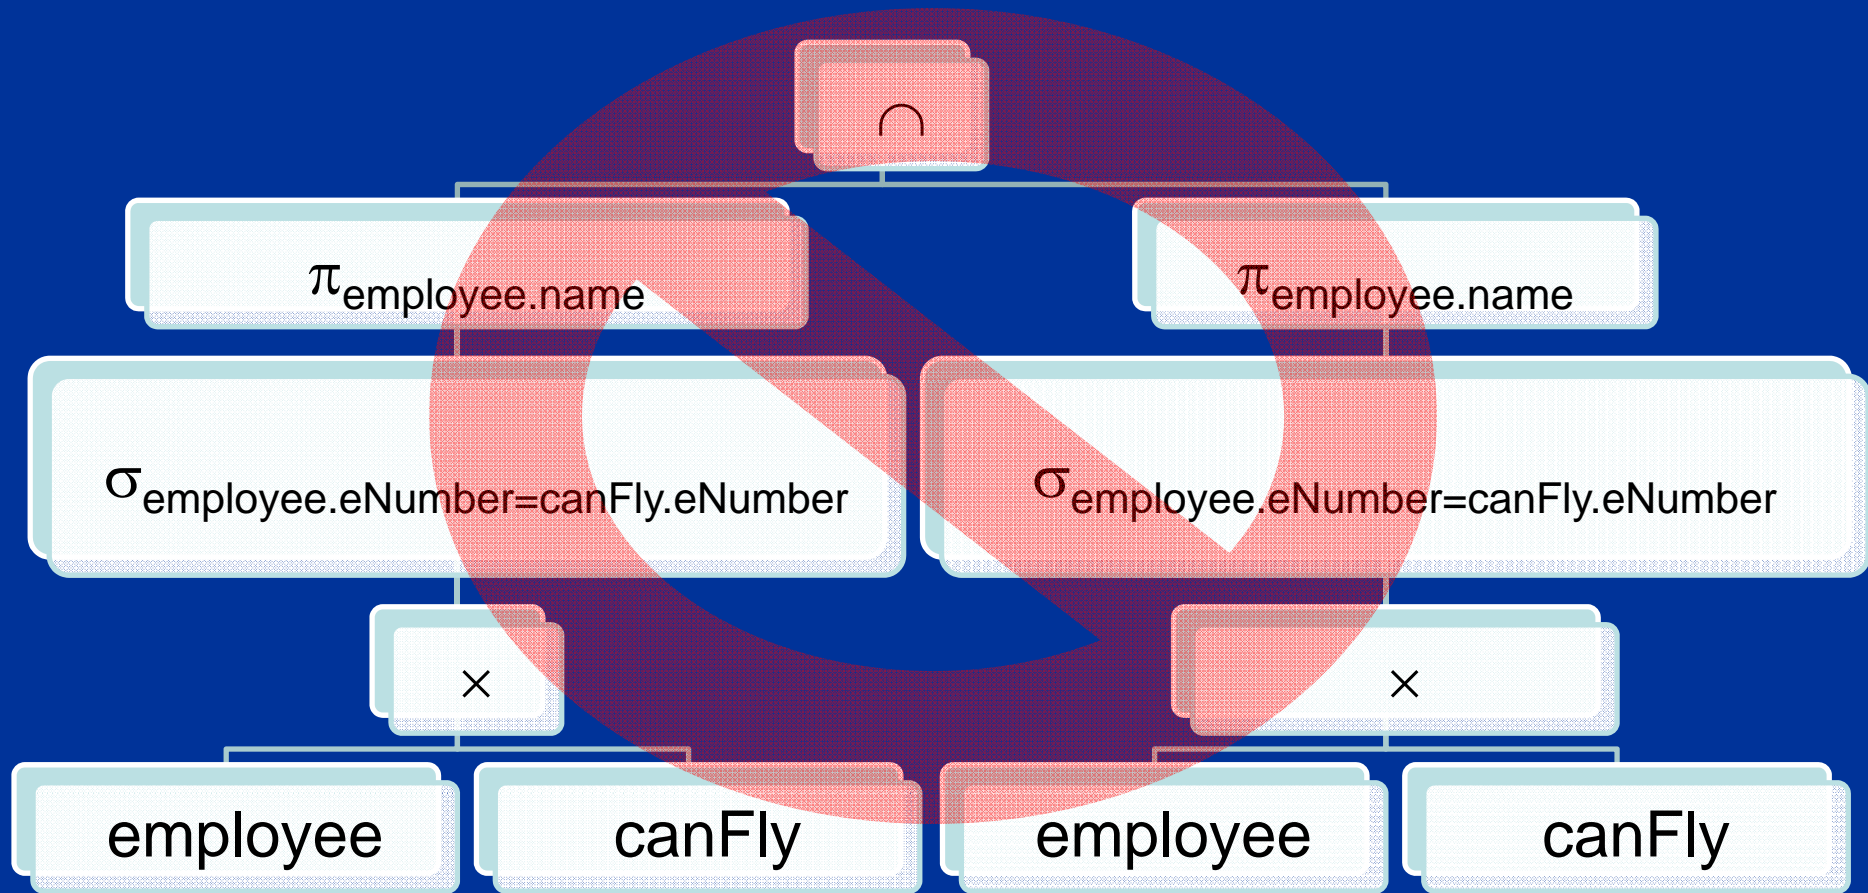
Find the names of employees who can fly two planes or more

name	salary	eNumber
Clark	150000	1006
Gates	5000000	1005
Jones	50000	1001
Peter	45000	1002
Phillips	25000	1004
Rowe	35000	1003
Warnock	500000	1007

eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC10

## Example

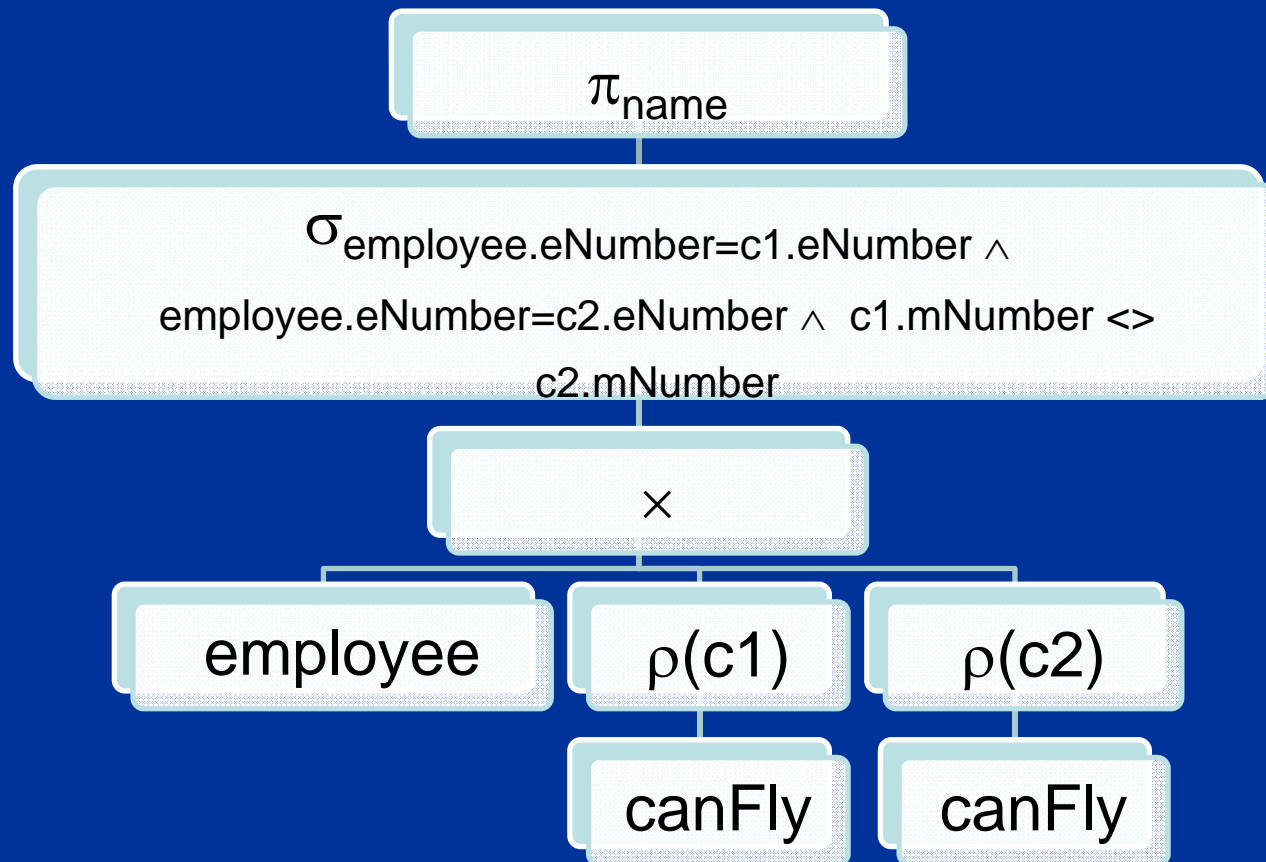
Find the names of employees who can fly two planes or more





## Example

Find the names of employees who can fly two planes or more



## Example (SQL)

Find the names of employees who can fly two planes or more

```
SELECT employee.name  
FROM canFly c1, canFly c2, employee  
WHERE c1.eNumber=employee.eNumber  
AND c2.eNumber=employee.eNumber  
AND c1.mNumber <> c2.mNumber
```

## Example (SQL with aggregates)

Find the names of employees who can fly two planes or more (no aggregates in Algebra)

```
SELECT employee.name  
FROM canFly, employee  
WHERE employee.eNumber=canFly.eNumber  
GROUP BY employee.eNumber, employee.name  
HAVING COUNT(*) >= 2
```



## Division

- Compute all possible combinations of the first column of  $R_1$  and  $R_2$ .

$$(\pi_A(R_1) \times R_2)$$

- Then remove those rows that exist in  $R_1$

$$(\pi_A(R_1) \times R_2) - R_1$$

- Keep only the first column of the result. These are the ***disqualified*** values

$$\pi_A((\pi_A(R_1) \times R_2) - R_1)$$

- $A/B$  is the first column of  $R_1$  **except** the disqualified values

$$R_1 / R_2 = \pi_A(R_1) - \pi_A((\pi_A(R_1) \times R_2) - R_1)$$

## Example

Find the employment numbers of employees who can fly **all** MD planes

eNumber	mNumber
1001	B727
1001	B747
1001	DC10
1002	A320
1002	A340
1002	B757
1002	DC9
1003	A310
1003	DC9
1003	DC10

maker	mNumber
Airbus	A310
Airbus	A320
Airbus	A330
Airbus	A340
Boeing	B727
Boeing	B747
Boeing	B757
MD	DC10
MD	DC9

## Division (Example)

Find the Employment numbers of the pilots who can fly **all** MD planes

$\text{canFly} / \pi_{\text{mNumber}}(\sigma_{\text{maker}='MD'}(\text{plane}))$

$\pi_{\text{eNumber}}(\text{canFly}) -$   
 $\pi_{\text{eNumber}}((\pi_{\text{eNumber}}(\text{canFly})$   
 $\times \pi_{\text{mNumber}}(\sigma_{\text{Maker}='MD'}(\text{plane}))) - \text{canFly})$

## Example: step-by-step

$\pi_{\text{eNumber}}(\text{canFly}) \times \pi_{\text{mNumber}}(\sigma_{\text{Maker}='MD'}(\text{plane}))$

eNumber	mNumber
1001	DC9
1001	DC10
1002	DC9
1002	DC10
1003	DC9
1003	DC10

## Example: step-by-step

$(\pi_{\text{eNumber}}(\text{canFly}) \times$   
 $\pi_{\text{mNumber}}(\sigma_{\text{Maker}='MD'}(\text{plane}))) - \text{canFly}$

eNumber	mNumber
1001	DC9
1002	DC10

## Example: step-by-step

$\pi_{\text{eNumber}}((\pi_{\text{eNumber}}(\text{canFly})$   
 $\times \pi_{\text{mNumber}}(\sigma_{\text{Maker}='MD'}(\text{plane}))) -$   
 $\text{canFly})$

eNumber
1001
1002

## Example: step-by-step

$\pi_{\text{eNumber}}(\text{canFly}) -$   
 $\pi_{\text{eNumber}}((\pi_{\text{eNumber}}(\text{canFly})$   
 $\times \pi_{\text{mNumber}}(\sigma_{\text{Maker}='MD'}(\text{plane}))) -$   
 $\text{canFly})$

eNumber
1003

## Division (SQL)

```
SELECT DISTINCT c1.eNumber
FROM canFly c1
WHERE NOT EXISTS
    (SELECT c.eNumber, p.mNumber
     FROM canFly c, plane p
     WHERE p.maker='MD'
           AND c1.eNumber=c.eNumber
     EXCEPT
     SELECT enumber, mNumber
     FROM canFly)
```



## Division (SQL)

```
SELECT DISTINCT c1.eNumber
FROM canFly c1
WHERE NOT EXISTS(
    SELECT *
    FROM plane p
    WHERE p.maker='MD' AND NOT EXISTS(
        SELECT *
        FROM canFly c2
        WHERE c1.eNumber=c2.eNumber
            AND p.mNumber=c2.mNumber))
```

## Credits

The content of this lecture is based  
on chapter 4 of the book  
“Introduction to database  
Systems”

By  
S. Bressan and B. Catania,  
McGraw Hill publisher

Animated characters are animated  
using VocaliseTTS under  
license from Digital Curiosity

Clipart and media are licensed from  
Microsoft Office Online Clipart  
and Media

Copyright © 2012 by Stéphane Bressan

