
Programming Language Concepts, CS2104

Lecture 11

Declarative Concurrency

Overview

- Declarative concurrency
- Mechanisms of concurrent program
- Streams
- Demand-driven execution
 - execute computation, if variable needed
 - needs suspension by a thread
 - requested computation is running in new thread
- By-Need triggers
- Lazy functions

The World is Concurrent!

- Concurrent programs
 - several activities execute simultaneously (concurrently)
- Most of the software used are concurrent
 - operating system: IO, user interaction, many processes, ...
 - web browser, Email client, Email server, ...
 - telephony switches handling many calls
 - ...

Why Should We Care?

- Software must be concurrent...
 - ... for many application areas
- Concurrency can be helpful for constructing programs
 - organize programs into independent parts
 - concurrency allows to make them independent with respect to how to execute
 - essential: how do concurrent programs interact?
- Concurrent programs can run faster on parallel machines (including clusters and cores)

Concurrency and Parallelism

- **Concurrency** is *logically simultaneous processing* which can also run on sequential machine.
- **Parallelism** is *physically simultaneous processing* and it involves multiple processing elements and/or independent device operations.
- A **computer cluster** is a group of connected computers that work together as a unit. One popular implementation is a cluster with nodes running Linux with support library (for parallelism).

Concurrent Programming is Difficult...

- This is the traditional belief
- The truth is: concurrency is *very* difficult...
 - ... if used with inappropriate tools and programming languages
- Particularly troublesome : *state* and *concurrency*

Concurrent Programming is Easy...

- Oz (as well as Erlang) has been designed to be very good at concurrency...
- Essential for concurrent programming here
 - data-flow variables
 - very simple interaction between concurrent programs, mostly automatic
 - light-weight threads

Declarative Concurrent Programming

- What stays the same
 - the result of your program
 - concurrency does not change the result
- What changes
 - programs can compute incrementally
 - incremental input... (such as reading from a network connection) ... and incremental processing

Our First Concurrent Program

```
declare x0 x1 x2 x3  
thread x1 = 1 + x0 end  
thread x3 = x1 + x2 end  
{Browse [x0 x1 x2 x3]}
```

- Browser will show [x0 x1 x2 x3]
 - variables are not yet assigned

Our First Program

```
declare x0 x1 x2 x3  
thread x1 = 1 + x0 end  
thread x3 = x1 + x2 end  
{Browse [x0 x1 x2 x3]}
```

- Both threads are suspended
 - $x1 = 1 + x0$ suspended; $x0$ unassigned
 - $x3 = x1 + x2$ suspended; $x1, x2$ unassigned

Our First Program

```
declare X0 X1 X2 X3  
thread X1 = 1 + X0 end  
thread X3 = X1 + X2 end  
{Browse [X0 X1 X2 X3]}
```

■ Feeding $X0 = 4$

Our First Program

```
declare x0 x1 x2 x3  
thread x1 = 1 + x0 end  
thread x3 = x1 + x2 end  
{Browse [x0 x1 x2 x3]}
```

- Feeding $x0 = 4$
 - First thread can execute, binds $x1$ to 5

Our First Program

```
declare x0 x1 x2 x3  
thread x1 = 1 + x0 end  
thread x3 = x1 + x2 end  
{Browse [x0 x1 x2 x3]}
```

- **Feeding** $x0 = 4$
 - First thread can execute, binds $x1$ to 5
 - Browser shows $[4 \ 5 \ x2 \ x3]$

Our First Program

```
declare x0 x1 x2 x3  
thread x1 = 1 + x0 end  
thread x3 = x1 + x2 end  
{Browse [x0 x1 x2 x3]}
```

- Second thread is still suspended
 - Variable x2 is still not assigned

Our First Program

```
declare x0 x1 x2 x3  
thread x1 = 1 + x0 end  
thread x3 = x1 + x2 end  
{Browse [x0 x1 x2 x3]}
```

- **Feeding** $x2 = 2$
 - Second thread can execute, binds $x3$ to 7
 - Browser shows [4 5 2 7]

Threads

- A **thread** is simply an executing program.
- A program can have more than one thread.
- A thread is created by :

`thread <s> end`

- Threads compute
 - independently
 - as soon as their statements can be executed
 - interact by binding variables in store

The Browser

- Browser is implemented in Oz as a thread.
- It also runs whenever browsed variables are bound
- It uses some extra functionality to look at unbound variables

Basic Concepts

- Model allows multiple statements to execute "*simultaneously*" ?
- Can imagine that these threads really execute in parallel, each has its own processor, but share the same memory
- Reading and writing different variables can be done simultaneously by different threads
- Reading the same variable can be done *concurrently*.
- Writing to the same variable to be done *sequentially*.

Causal Order

- In a sequential program, all execution states are *totally ordered*
- In a concurrent program, all execution states *of a given thread* are totally ordered
- But, ... the execution state of the concurrent program as a whole is **partially ordered**


Total Order

- In a sequential program all execution states are *totally ordered*
- Computation step: transition between two consecutive execution states




Causal Order = Partial Order

- In a concurrent program all execution states of a given thread are totally ordered
- The execution state of the concurrent program is *partially ordered*




Causal Order = Partial Order



Nondeterminism

- An execution is *nondeterministic* if there is a computation step in which there is a **choice** what to do next
- Nondeterminism appears naturally when there are multiple concurrent states

Example of Nondeterminism



- The thread that binds x first will continue, the other thread will raise an exception

Nondeterminism

- If there is only one binder for each dataflow variable, nondeterminism is not *observable* on the store.
- That is the store has the same final results.
- Hence, for correctness we can ignore the concurrency
- This concept is known as "Declarative Concurrency".

Declarative concurrency

- *Declarative programming (Reminder):*
 - the output of a declarative program should be a mathematical function of its input.
- *Functional programming (Reminder):*
 - the program executes with some input values and when it terminates, it has returned some output values.
- *Data-driven concurrent model:* a **concurrent** program is **declarative** if all executions with a given set of inputs have one of two results:
 - (1) they all do not terminate or
 - (2) they all eventually reach partial termination and give results that are logically equivalent.

Partial Termination. Example

```
fun {Double Xs}  
case Xs of  
  nil then nil  
  [ ] X|Xr then 2*X|{Double Xr} end  
end  
Ys={Double Xs}
```

- As long as input stream x_s grows, then output stream y_s grows too. The program never terminates.
- However, if the input stream stops growing, then the program will eventually stop executing too.
- The program does *a partial termination*.

Partial Termination. Examples

- If the inputs are bound to some partial values, then the program will eventually end up in partial termination. Also, the outputs will be bound to some partial values.
- What is the relation of outputs in terms of inputs when we consider partial values?
- Example:

$X_S = 1 \mid 2 \mid 3 \mid \text{Xr} \rightarrow Y_S$ will be bound to $2 \mid 4 \mid 6 \mid _$

- Having $\text{Xr} = 4 \mid 5 \mid \text{Xr1}$, we get Y_S bound to $2 \mid 4 \mid 6 \mid 8 \mid 10 \mid _$
- Making $\text{Xr1} = \text{nil}$, we get Y_S bound to $[2 \ 4 \ 6 \ 8 \ 10]$

Scheduling

- The choice of which thread to execute next and for how long is done by the *scheduler*
- A thread is *runnable* if its next statement to execute is not blocked on a dataflow variable, otherwise the thread is *suspended*

Scheduling

- A scheduler is *fair* if it does not starve each runnable thread
 - All runnable threads execute eventually
- Fair scheduling makes it easier to reason about programs
- Otherwise some runnable programs will never get its turn for execution.

Example of Runnable Threads

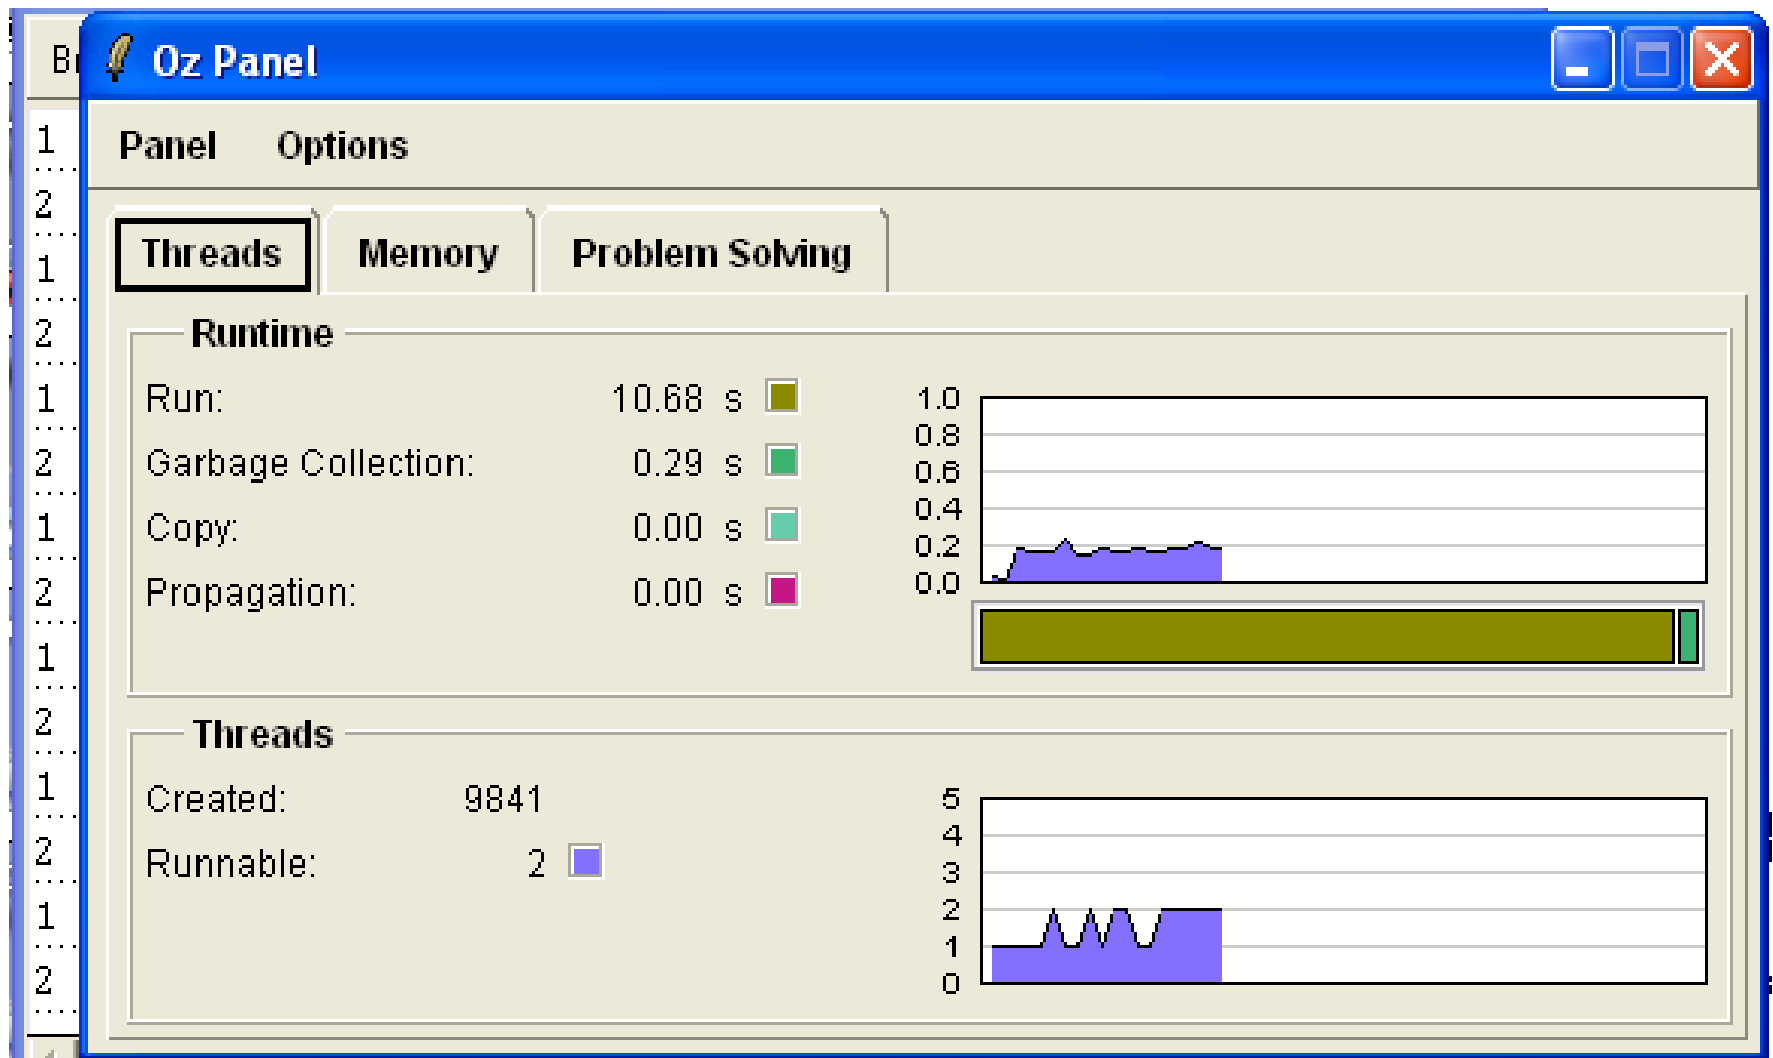
```
thread
  for I in 1..10000 do {Browse 1} end
end
thread
  for I in 1..10000 do {Browse 2} end
end
```

Example of Runnable Threads

```
thread
  for I in 1..10000 do {Browse 1} end
end
thread
  for I in 1..10000 do {Browse 2} end
end
```

- This program will interleave the execution of two threads, one printing 1, and the other printing 2
- fair scheduler

Example of Runnable Threads



Dataflow Computation

- Threads suspend when dataflow variables needed are not yet bound
- {Delay X} primitive makes the thread suspends for x milliseconds, after that the thread is runnable

```
declare X
{Browse X}
local Y in
  thread {Delay 1000} Y = 10*10 end
  X = Y + 100*100
end
```

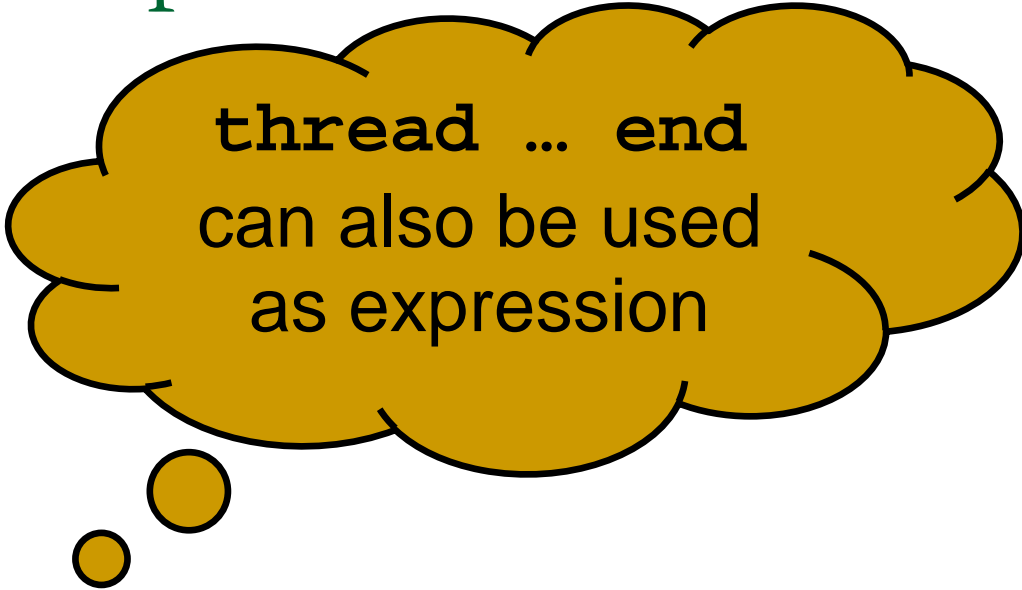
Concurrency is Transparent

Example : a concurrent map operation

```
fun {CMap Xs F}  
  case Xs  
  of nil then nil  
  [] X|Xr then  
    thread {F X} end | {CMap Xr F}  
  end  
end
```

Concurrency is Transparent

```
fun {CMap Xs F}  
  case Xs  
  of nil then nil  
  [] X|Xr then  
    thread {F X} end | {CMap Xr F}  
  end  
end
```



thread ... end
can also be used
as expression

Concurrency is Transparent

- What happens:

declare F

`{Browse {CMap [1 2 3 4] F} }`

- Browser shows `[_ _ _ _]`

- `CMap` computes the list skeleton

- newly created threads suspend until F becomes bound

Concurrency is Transparent

- What happens:

$F = \mathbf{fun} \ \{\$ \ X\} \ X+1 \ \mathbf{end}$

- Browser shows [2 3 4 5]

Cheap Concurrency and Dataflow

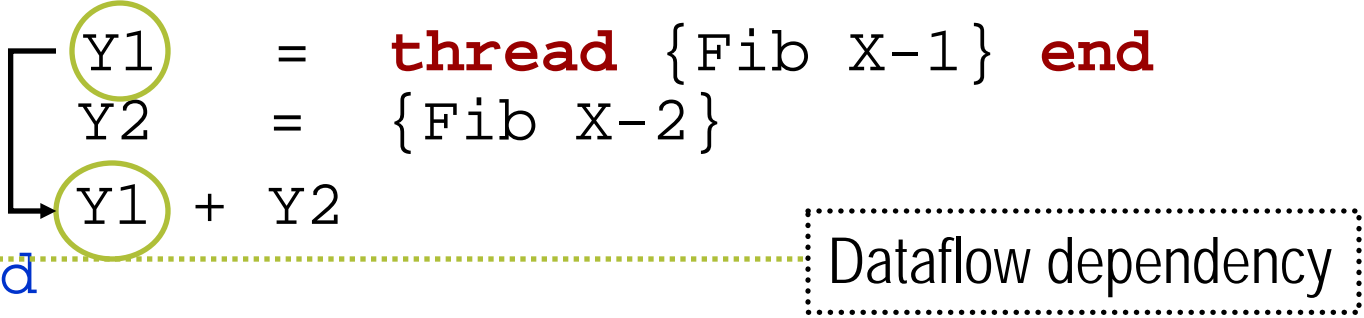
- Declarative programs can be easily made concurrent
- Just use the `thread` statement where concurrency is needed

Cheap Concurrency and Dataflow

```
fun {Fib X}
  if X==0 then 0
  elseif X==1 then 1
  else
    thread {Fib X-1} end + {Fib X-2}
  end
end
```


Understanding why


```
fun {Fib X}
  if X==0 then 0 elseif X==1 then 1
  else Y1 Y2 in
    [Y1 = thread {Fib X-1} end
     Y2 = {Fib X-2}
     Y1 + Y2]
  end
end
```



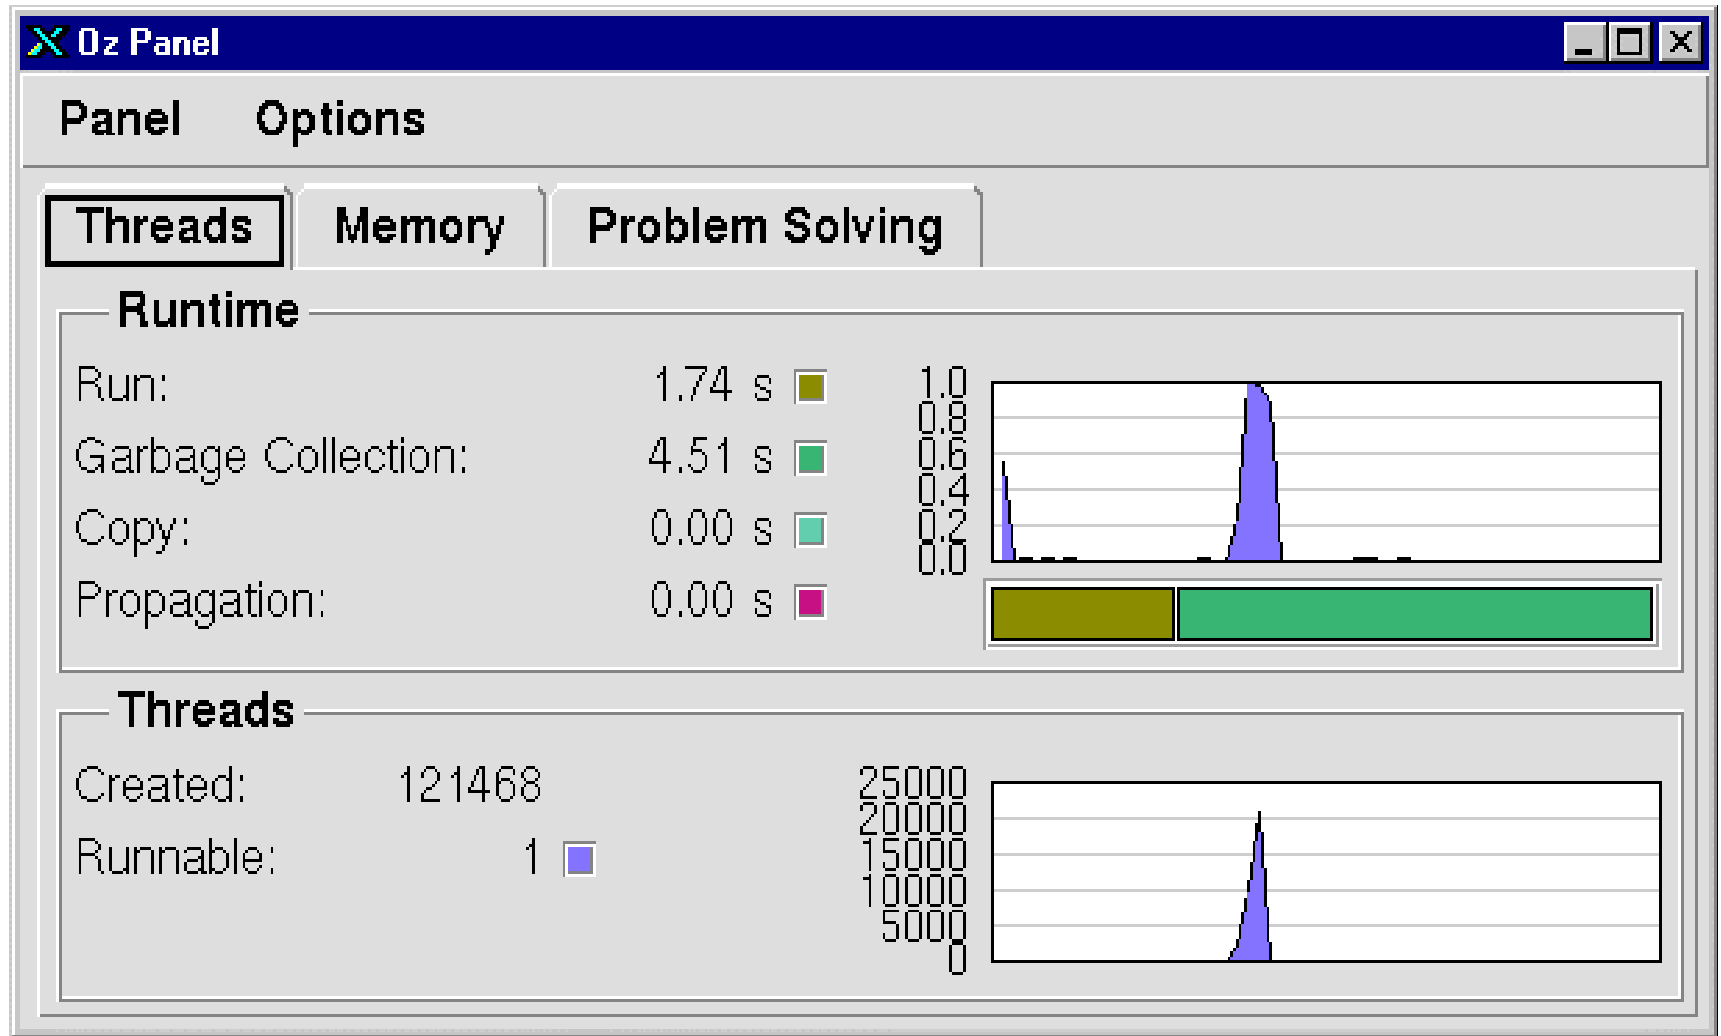
Dataflow dependency

Execution of {Fib 6}

{Fib 6} is denoted as F6,...



Fib



Streams

- A most useful technique for declarative concurrent programming to use **streams** to communicate between threads.
- A **stream** is a potentially unbounded list of messages, i.e., it is a list whose tail is an unbound dataflow variable.
- A thread communicating through streams is a kind of “active object”, also called **stream object**.
- A sequence of stream objects each of which feeds the next is called a **pipeline**.
- **Deterministic stream programming**: each stream object always knows for each input where the next message will come from.


Producer \Leftrightarrow Consumer

```
thread X={Produce} end
```

```
thread Result={Consume X} end
```

- Typically, what is produced will be put on a list that never ends (without `nil`), called **stream**
- **Consumer** (also called **sink**) consumes as soon as **producer** (also called **source**) produces

Producer/Consumer Stream



$Xs = \{\text{Produce } 0 \text{ Limit}\}$

$S = \{\text{Consume } Xs \ 0\}$

Example: Producer \Leftrightarrow Consumer


```
fun {Produce N Limit}
  if N<Limit then
    N|{Produce N+1 Limit}
  else nil end
end
fun {Consume Xs Acc}
  case Xs of X|Xr then
    {Consume Xr Acc+X}
  [] nil then Acc
  end
end
```

Stream Transducer. Example

```
thread Stream={Produce 0 1000} end
thread FilterResult={Filter Stream IsOdd} end
thread Result={Consume FilterResult 0} end
```

- **Transducer:** a stream which reads the producer's output and computes a filtered stream for the consumer.
- Can be: filtering, mapping, ...
- Advantages of pipeline:
 - there is no need to wait the final value of the producer
 - producer, transducer, and consumer are executed concurrently

Simple Pipeline



$Ys = \{\text{Filter } Xs \dots\}$

Client \Leftrightarrow Server

- Similar to producer \Leftrightarrow consumer
- Typical scenario:
 - more clients than servers
 - server has a fixed identity
 - clients send messages to server
 - server replies
- See Next Lecture: message sending

Fairness

- Essential that even though producer can always produce, consumer also gets a chance to run
- Threads are scheduled with **fairness**
 - if a thread is runnable, it will eventually run

Thread Scheduling

- More guarantees than just fairness
- Threads are given a time slice to run
 - approximately 10ms
 - when time slice is over: thread is **preempted**
 - next runnable thread is **scheduled**
- Can be influenced by priorities
 - high, medium, low
 - controls relative size of time slice (Sections 4.2.4-4.2.6)

How to Control Producers?

- *Eager model*: the producer decides when enough data has been sent
- *Possible problem*: producer should not produce more than needed
- *One attempt*: make consumer the driver
 - consumer produces stream skeleton
producer fills skeleton

Make Consumer be the Driver

```
fun {DConsume ?Xs A Limit}
  if Limit>0 then
    local X Xr in
      Xs=X|Xr {DConsume Xr A+X Limit-1}
    else A end
  end
proc {DProduce N Xs}
  case Xs of X|Xr then
    X=N
    {DProduce N+1 Xr}
  end
end
```

Overall program :

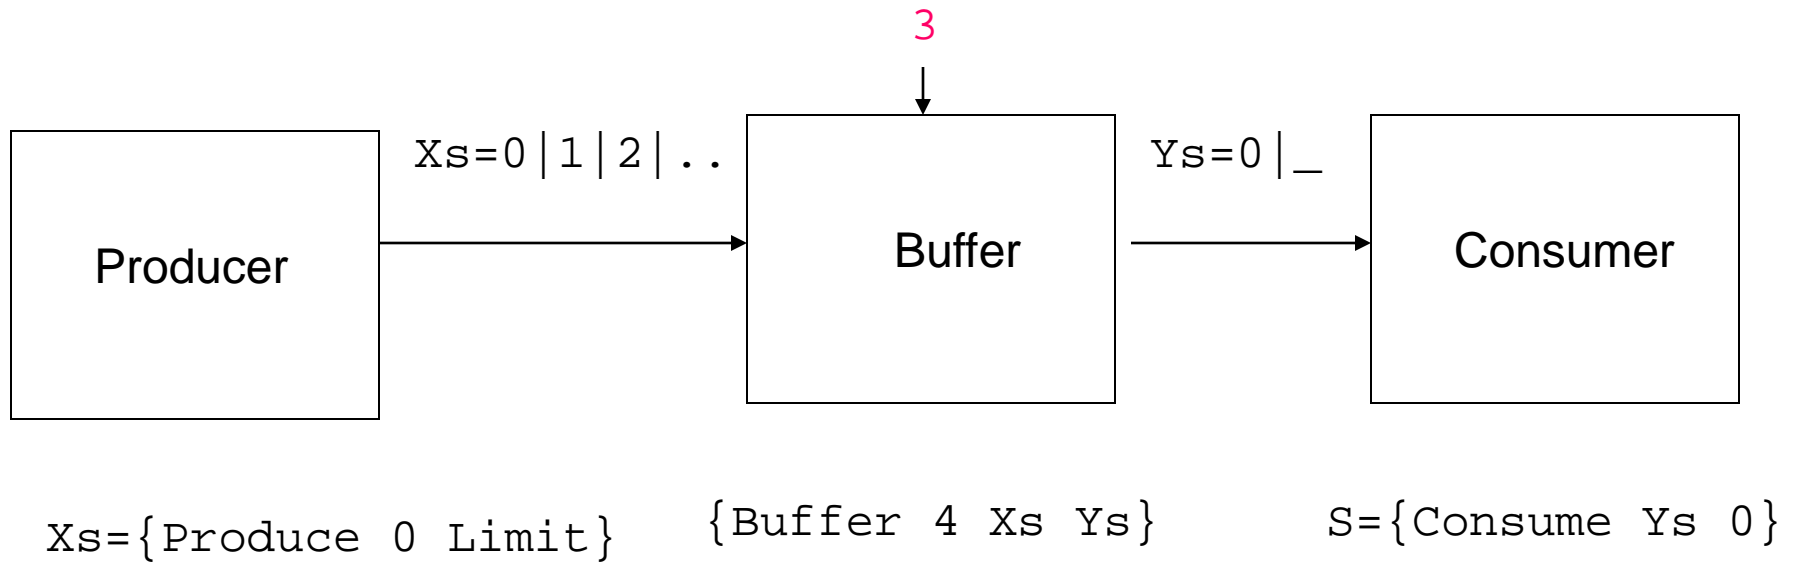
```
local Xs S in
  thread {DProduce 0 Xs} end
  thread S={DConsume Xs 0 150000} end
  {Browse S}
end
```

Note that consumer controls how many elements are needed.

Bounded Buffer

- *Eager – producer may run ahead*
- *Demand-driven – consumer in control but more complex execution.*
- *Compromise : Bounded Buffer*

Bounded Buffer



Bounded Buffer Code

input

output

```
proc {Buffer N Xs Ys}
  fun {Startup N ?Xs}
    if N==0 then Xs
    else Xr in Xs=_|Xr {Startup N-1 Xr} end
  end
  proc {AskLoop Ys ?Xs ?End}
    case Ys of Y|Yr then Xr End2 in
      Xs=Y|Xr % get element from buffer
      End=_|End2 % replenish the buffer
      {AskLoop Yr Xr End2}
    [] nil then End=nil
    end
  end
end
End={Startup N Xs}
in
  {AskLoop Ys Xs End}
end
```

Lazy Streams

- Better solution for demand-driven concurrency
Use Lazy Streams

That is consumer decides, so producer runs on request.

Needed Variables

- Idea: start execution, when value for variable needed
short: **variable needed**
- Value for variable needed...
...a thread suspends on variable!

Lazy Execution (Reminder)

- Up to now the execution order of each thread follows textual order.
That is each statement is executed in order, whether or not its results are needed later.
- This execution scheme is called *eager execution*, or *supply-driven* execution
- Another execution order is that a statement is executed only if its results **are needed** somewhere in the program
- This scheme is called **lazy evaluation**, or **demand-driven evaluation**

Lazy Execution. Reminder

declare

fun lazy {F1 X} 2*X end

fun {F2 Y} Y*Y end

B = {F1 3}

{Browse B} → nothing (simply unbound B)

C = {F2 4}

{Browse C} → display 16

A = B+C → display 6 for B

- F1 is a lazy function
- B = {F1 3} is executed only if its result is needed in A = B+C

Example

declare

```
fun lazy {F1 X} 2*X end
```

```
fun lazy {F2 Y} Y*Y end
```

```
B = {F1 3}
```

```
{Browse B} % → nothing (simply unbound B)
```

```
C = {F2 4}
```

```
{Browse C} % → nothing (simply unbound C)
```

- F1 and F2 are now lazy functions
- B = {F1 3} and C = {F2 4} are executed only if their results are needed in an expression, like: A = B+C


Example

```
declare
fun lazy {F1 X} 2*X end
fun lazy {F2 Y} Y*Y end
B = {F1 3}
{Browse B} % → display 6
C = {F2 4}
{Browse C} % → display 16
A = B+C
```

- F1 and F2 are now lazy functions
- B = {F1 3} and C = {F2 4} are executed because their results are needed in A = B+C


Example

- In **lazy execution**, an operation suspends until its result is needed
- Each suspended operation is triggered when another operation needs the value for its arguments
- In general, multiple suspended operations can start concurrently



Example II

- In **data-driven execution**, an operation suspends until the values of its arguments results are available
- In general, the suspended computation can start concurrently



Lazy Production

```
fun lazy {Produce N}  
    N | {Produce N+1}  
end
```

- Intuitive understanding: function executes only, if its output is needed

Example: Lazy Production

```
fun lazy {Produce N}  
    N | {Produce N+1}  
end  
declare Ns={Produce 0}  
{Browse Ns}
```

- Shows again Ns
 - Remember: `Browse` does not need the values of the variables

Example: Lazy Production

```
fun lazy {Produce N}  
    N | {Produce N+1}  
end  
declare Ns={Produce 0}
```

- **Execute** `_ = Ns.1`
 - needs the variable `Ns`
 - Browser now shows `0 | _` or `0 | <Future>`

Example: Lazy Production

```
fun lazy {Produce N}  
    N | {Produce N+1}  
end  
declare Ns = {Produce 0}
```

- Execute `_ = Ns . 2 . 2 . 1`
 - needs the variable `Ns . 2 . 2`
 - Browser now shows `0 | 1 | 2 | _`

Everything can be Lazy!

- Not only producers, but also transducers can be made lazy
- Sketch
 - consumer needs variable
 - transducer is triggered, needs variable
 - producer is triggered

Lazy Transducer. Example

```
fun lazy {Inc Xs}  
  case Xs  
  of X|Xr then X+1 | {Inc Xr}  
  end  
end  
  
declare Xs={Inc {Inc {Produce N}}}
```


Stream Object

Diagram illustrating the Stream Object structure and its recursive definition:

```
fun {StreamObject S1 X1 ?T1}
  case S1 of M|S2 then N X2 T2 in
    {NextState M X1 N X2}
    T1 = N|T2 {StreamObject S2 X2 T2}
  [] nil then T1=nil end
end

declare S0 X0 T0
thread {StreamObject S0 X0 T0} end
```

Making the Driver into a Stream

Diagram illustrating the transformation of a driver into a stream, showing the flow of data and state variables.

Labels and Arrows:

- input**: Points to `?S1`.
- accumulator**: Points to `X1`.
- output**: Points to `?T1`.

```
fun {StreamObject ?S1 X1 ?T1}
  S1=M|S2
  local N X2 T2 in
    {NextState M X1 N X2}
    T1 = N|T2
    {StreamObject S2 X2 T2}
  end
end

declare S0 X0 T0
thread {StreamObject S0 X0 T0} end
```

The diagram shows the flow of data and state variables in a stream processing context. The `fun` block defines a stream object `StreamObject` with parameters `?S1`, `X1`, and `?T1`. The `input` label points to `?S1`, the `accumulator` label points to `X1`, and the `output` label points to `?T1`. The code inside the `fun` block shows the state transitions and the creation of a new stream object. The `local` block contains the state variables `N`, `X2`, and `T2`, and the `NextState` function. The `StreamObject` is created with parameters `S2`, `X2`, and `T2`. The `end` block shows the `declare` and `thread` statements.

Fork-Join for Threads

```
local  $X_1$   $X_2$   $\dots$   $X_{n-1}$   $X_n$  in  
  thread <stmt1>  $X_1 = \text{unit}$  end  
  thread <stmt2>  $X_2 = X_1$  end  
    :  
  thread <stmt $n$ >  $X_n = X_{n-1}$  end  
  {Wait  $X_n$ }  
end
```



wait for all threads to complete through variable binding

Barrier Synchronization

list of threads

```
proc {Barrier Ps}
  fun {Loop Ps L}
    case Ps of P|Pr then M in
      thread {P} M=L end
      {Loop Pr M}
    [] nil then L
    end
  end
  S={Loop Ps unit}
in
  {Wait S}
end
```

wait for all threads to complete