

School of Computing
National University of Singapore
CS4243 Computer Vision and Pattern Recognition
Semester 1, AY 2013/14

Lab 5 : Principal Component Analysis

Objectives:

- Learn to use NumPy for numerical computations.
- Learn to write a Python program to do Principal Component Analysis.

Preparation:

- Create a folder in the PC with your name, e.g., `d:/myname`. This folder will be used as your working directory.
- Download the following files into your working directory: `pcaAY1314_v3.doc` & `md-data.txt`.

Part 0. Initialisation

Set the working directory, e.g., `d:/myname`

```
>>> import os  
>>> os.chdir("d:/myname")
```

Open `md-data.txt` using your favourite editor. You should see 30 rows of values. Each row contains 11 values corresponding to the 11 components of a data point. That is, these data points are 11-dimensional.

Part 1. PCA

Write a program to perform PCA as follows. After writing the codes for each step, run the program to verify the results.

1. Import NumPy (`linalg` stands for “linear algebra”):

```
import numpy as np  
import numpy.linalg as la
```

2. Read data in `md-data.txt` into a $d \times n$ matrix called `data`:

```
infile = open("md-data.txt")
data = np.genfromtxt(infile, delimiter=",")
data = np.matrix(data).T
infile.close()
```

Print `data` to examine its content (see on your screen only, no need to print out hardcopy). It should contain 30 column vectors of 11 components each.

3. Get the number of dimensions d and number of data points n :

```
d = data.shape[0]
n = data.shape[1]
```

Print the values of d and n . Are they correct?

4. Compute the mean of the data points:

```
mean = np.mean(data, 1)
```

The function `np.mean` computes the mean along axis 1, i.e., horizontally. Print `mean` to verify that it is a column vector.

5. Subtract mean from data points:

```
sdata = data - mean
```

This subtraction shifts the data points so that their centroid is at the origin. You can verify this by printing the mean of `sdata`:

```
print np.mean(sdata, 1)
```

The values printed should be very close to 0.

6. Compute the covariance matrix of the shifted data points:

```
cov = np.cov(sdata, None, 1, 1)
```

To call this function correctly, you need to indicate that the vectors in `sdata` are arranged in columns. Check up NumPy reference guide for the meanings of arguments.

7. Compute eigen-decomposition:

```
eigvalues, eigvectors = la.eig(cov)
index = np.argsort(eigvalues)
```

The function `la.eig` performs eigen-decomposition and returns two arrays. The first array contains the eigenvalues and the second contains the eigenvectors arranged in columns. The function `la.eig` does not necessarily return eigenvalues in sorted order.

The function `np.argsort` sorts the eigenvalues and returns an array of indices ordered in **increasing order** of eigenvalues.

Part 2. Dimensionality Reduction

Add the following codes to the program to perform dimensionality reduction.

1. Compute the ratio of **unaccounted variance** for each possible reduced dimension l from 1 to d :

$$R = \frac{\sum_{j=l+1}^d \lambda_j}{\sum_{j=1}^d \lambda_j}$$

Note that in this formula, the eigenvalues are assumed to be sorted in **decreasing order**. For example, if $l = 3$, then only the first three largest eigenvectors are included, and the variances from λ_4 to λ_d are not accounted for. So, the above formula measures the ratio of unaccounted variance over the total variance.

Be careful of the following:

- a. The indices of python arrays and matrices **begin with 0** instead of 1. For example `eigvalues[0]` is the first eigenvalue in `eigvalues`. On the other hand, the indices in the above formula start with 1.
- b. The function `np.argsort` sorts the eigenvalues in **increasing order** instead of decreasing order. So, `eigvalues[index[0]]` is the smallest eigenvalue λ_d . The following figure illustrates the relationship between the eigenvalues stored in `eigvalues` and those represented by λ_j in the above formula:

	i	0				d-j					d-1
eigvalues[index[i]]		λ_d				λ_j					λ_1

2. Plot a graph of R vs. number of dimensions l by doing the following:

Add the following statement to your program:

```
import matplotlib.pyplot as plt
```

Then you can use `plt.plot` to plot the graph.

3. Based on the plotted graph, select a smaller number of dimensions l that you think is enough for representing the data points.
4. Compose the matrix Q of eigenvectors from `eigvectors`. Arrange the eigenvectors in Q in **decreasing order** of eigenvalues. That is, $Q[:, 0]$ contains the eigenvector with the largest eigenvalue. Remember that Q must be a matrix instead of a 2D array.
5. Compute the y_i vectors of the data points in the eigenspace:

$$y_i = Q^T (x_i - m)$$

Write this python statement without using explicit `for` loop to compute the y_i vector for each data point x_i . This will let your program run faster.

6. Keep the first l components of the y_i vectors, where l is the number of dimensions you selected in Step 3. Call these reduced dimensional vectors \hat{y}_i . Print \hat{y}_i to make sure that their dimensions are correct.
7. Create a reduced dimensional Q matrix, call it \hat{Q} , by copying the first l eigenvectors from the original Q to \hat{Q} . Print \hat{Q} to make sure that it has the correct l eigenvectors.
8. Compute estimates of the data points from \hat{Q} and \hat{y}_i

$$\hat{x}_i = \hat{Q} \hat{y}_i + m$$

Again, write this python statement without explicit `for` loop.

9. Use `la.norm` to compute the (square-root of the squared) difference between the original data points x_i and the estimated data points \hat{x}_i (summed over all data points). This difference should be small if l dimensions are enough to represent the data points sufficiently.
10. Repeat steps 4 to 10 for l ranging from 1 to d . Plot a graph with the `la.norm` values in step 10 along the vertical axis, and l along the horizontal axis. Comment on the curve w.r.t. the curve plotted in step 3.

Lab Report

1. Your lab report should contain the program that you write and the following results:

- a. The values of d and n in Part 1 Step 3.
- b. The mean vector of the input data points computed in Part 1 Step 4.
- c. The mean vector of the shifted data points computed in Part 1 Step 5.
- d. The eigenvalues and sorted indices computed in Part 1 Step 7.
- e. The graph of R vs. l plotted in Part 2 Step 3
- f. The l value you selected in Part 2 Step 4
- g. The difference you computed in Part 2 Step 10 for the selected l value.
- h. The graph of $la.norm$ vs l plotted in Part 2 Step 11

2. Submit the softcopy of your Python program to IVLE

Please put your python program in a folder and submit the folder. Use the following convention to name your folder:

MatriculationNumber_yourName_Lab#_Session. For example, if your matriculation number is A1234567B, and your name is Chow Yuen Fatt, for this lab, your file name should be A1234567B_ChowYuenFatt_Lab5_Wed630pm.